
REWARDING THE GRAPH BEHIND THE CHAIN: TOPOLOGY-AWARE PROCESS SUPERVISION FOR RL AND REASONING DISTILLATION

Anonymous authors

Paper under double-blind review

ABSTRACT

Long chain-of-thought reasoning is not purely linear: intermediate results are reused across non-adjacent steps, while unsupported conclusions may refer backward or form circular support. Outcome rewards and step-local process rewards do not expose these failures. We propose TopoPRM, a topology-aware implicit process reward model that first recovers an order-agnostic directed support graph from a free-form trace and then projects it onto a dependency DAG. Crucially, the backward and cyclic edges removed by this projection are retained as direction-consistency and acyclicity signals. We combine outcome, format, topology, continuity, and length through a topology-conditioned Choquet reward whose interaction terms explicitly value agreement between structural and non-structural sources. The same projected graph guides distillation toward a compact dependency-preserving trace, without requiring graph-format generation from the policy. We evaluate whether TopoPRM advances the fixed-scale accuracy–length Pareto frontier over outcome-only GRPO and additive or multiplicative aggregation, together with controlled edge-recovery, topology-perturbation, and reward-hacking tests.

1 INTRODUCTION

Reinforcement learning with verifiable rewards (RLVR) has substantially advanced long chain-of-thought (CoT) reasoning in large language models and now serves as a core mechanism for training mathematical and symbolic reasoning ability (Wei et al., 2022; Kojima et al., 2022; DeepSeek-AI, 2025; Yeo et al., 2025; Shao et al., 2024; Ke et al., 2026). From sparse final-answer rewards (Cobbe et al., 2021; DeepSeek-AI, 2025) to widely used process reward models (PRMs) that attach supervision to intermediate positions or prefixes (Lightman et al., 2023; Wang et al., 2024; Uesato et al., 2022), the dominant training interface still represents a trace as a *linear sequence of steps*. In practice, long reasoning is rarely linear: in mathematical derivations, a later conclusion often depends on several non-adjacent premises, an intermediate equation is reused by multiple downstream derivations, and subgoals branch before merging into the final answer (Besta et al., 2024; Zhong et al., 2026). Reconciling this structural mismatch has emerged as a core challenge for advancing reasoning supervision beyond linear credit assignment.

Recent reasoning models have made the macroscopic structure of inference clearer and more reliable, yet they remain confined to the CoT-as-result view. DeepSeek-R1 (DeepSeek-AI, 2025), DAPO (Yu et al., 2025), and Tulu 3 (Lambert et al., 2024) demonstrated that outcome-only rewards scale, but they could not distinguish a well-supported derivation from a trace that reached the same answer through unsupported intermediate claims, and they remained vulnerable to reward hacking (Gao et al., 2023; Skalse et al., 2022). On the PRM side, denser signals have been obtained through Monte Carlo rollouts (Wang et al., 2024; Luo et al., 2024), trained verifiers (Zhang et al., 2025; Khalifa et al., 2025; Zhao et al., 2026), and implicit outcome-derived scoring (Yuan et al., 2025; She et al., 2025). These methods can condition on a full prefix, but they do not explicitly recover the **multi-parent support relation** in which one intermediate result supports, or is supported by, several non-adjacent steps (Wang et al., 2025b). Likewise, token-level distillation (Agarwal et al., 2024) does not require a compact target to preserve those cross-step dependencies.




Figure 1: TopoPRM overview in the original reasoning-graph design. A free-form mathematical trace is mapped to a layered dependency graph; direction and acyclicity enter the topology-conditioned multi-source reward, while the same graph diagnostics localize dependency-preserving teacher revisions for distillation. The policy itself never emits graph tags.

We propose TOPOPRM, which supervises reasoning through its *implicit topological dependency structure* rather than assuming that surface order is already valid (Figure 1). A frozen, order-agnostic edge encoder first recovers a raw directed support graph $\mathcal{H}_y$ from an unchanged free-form trace. Projecting $\mathcal{H}_y$ onto a text-order DAG $\mathcal{G}_y$ makes the discarded structure informative: backward-edge mass measures direction inconsistency, while cycle-edge mass measures circular support. These diagnostics then enter a topology-conditioned, non-additive aggregation of outcome, format, topology, continuity, and length. The interaction terms reward agreement between topology and the other sources instead of reducing topology to a small additive bonus. Finally, the same projected graph guides distillation toward a compact dependency-preserving trace. The policy never has to decode graph tags or a special DAG format. Our contributions are threefold:

1. **Raw support graph before acyclic projection.** We recover directed dependencies without imposing textual order, and turn the backward and cyclic structure removed by projection into explicit process supervision.
2. **Topology-conditioned multi-source optimization.** We introduce a monotone 2-additive Choquet reward whose topology–outcome, topology–continuity, and topology–length interactions cannot be represented by an ordinary weighted sum, and reuse the projected graph for Topology-Guided Distillation.
3. **Accuracy–efficiency evaluation.** We compare against outcome-only and explicit length-shaped GRPO, projected-graph controls, additive and multiplicative aggregation, and distillation baselines under a fixed-scale, fixed-budget protocol, reporting both accuracy and the accuracy–length Pareto frontier.

2 RELATED WORK

Reinforcement Learning with Verifiable Rewards. RLVR replaces preference models with programmatic correctness checks on tasks whose answers are objectively verifiable (Shao et al., 2024; DeepSeek-AI, 2025; Lambert et al., 2024; Zeng et al., 2023). Recent work refines GRPO-style optimization through dynamic sampling and clipping (Yu et al., 2025), extended training schedules (Liu et al., 2025), finer credit assignment (Kazemnejad et al., 2025), asymmetric handling of negative samples (Zhu et al., 2025), entropy regularization (Cui et al., 2025), and simplified baselines (Ahmadian et al., 2024; Hu, 2025). Outcome-only signals nevertheless remain structurally blind: traces with very different dependency support can be reward-equivalent (Chen et al., 2026), and reward hacking persists even when the visible chain appears benign (Gao et al., 2023; Skalse et al., 2022; Wang et al., 2026a). *TopoPRM injects frozen topology and continuity diagnostics from latent reasoning graphs, targeting structural credit assignment directly rather than through a correctness surrogate.*

Concise-reasoning objectives introduce a separate confound: penalizing length inside a group-relative reward can assign negative advantage to a correct but long rollout. DRPO addresses this by decoupling the length-driven target distribution from correctness discrimination (Li et al., 2026a). LASER-D instead assigns step length rewards using targets that adapt to training-time behavior and per-question difficulty (Liu et al., 2026b). For multi-source reward integration, HERO stratifies dense reward-model scores by binary verifier outcomes and reweights prompts by within-group variance (Tao et al., 2026); this preserves correctness ordering but makes dense feedback explicitly correctness-first. Our hypothesis is different: raw direction and acyclicity should remain identifiable reward sources rather than only tie-breakers inside correctness strata. We therefore include fixed and adaptive length shaping plus a verifier-stratified control under the same completion budget, and treat accuracy and length as co-primary outcomes.

Process Reward Models for Reasoning. PRMs score intermediate steps to address the sparsity of outcome rewards, originally relying on dense human annotation (Uesato et al., 2022; Lightman et al., 2023). Subsequent work reduces labeling cost via Monte Carlo or tree-search rollouts (Wang et al., 2024; Luo et al., 2024; Zhang et al., 2024; Lu et al., 2024; Zhang et al., 2025; Guan et al., 2025), generative PRMs make verifier reasoning explicit (Khalifa et al., 2025; Zhao et al., 2026; She et al., 2025), while recent process rewards use intrinsic semantic cues or explicitly tagged meta-reasoning actions (He et al., 2026; Zhang et al., 2026b). Process advantage verifiers tie step values to rollout-induced returns (Setlur et al., 2025; Lai et al., 2024); implicit PRMs instead derive them from outcomes alone (Yuan et al., 2025; Snell et al., 2024). Recent benchmarks expose persistent weaknesses on fine-grained error categories and inter-step consistency (Song et al., 2025; Zheng et al., 2025). Although such scorers may read the full prefix, their output is attached to a position in one sequence rather than to an explicit many-to-many support relation among steps. *TopoPRM requires no tagged policy format and trains no scalar neural process verifier; it derives frozen global-structure supervision from a recovered dependency graph together with local continuity.*

Structured and Graph-Based Reasoning. PARC links a completed chain to preceding premises for error identification (Mukherjee et al., 2025), while GRiD generates knowledge/reasoning nodes and verifies each node against declared parents (Wen et al., 2025). Graph of Verification adaptively decomposes a trace into a DAG and verifies node blocks in topological order without training (Fang et al., 2026). Other methods require graph-formatted reasoning: DAG-Math evaluates graph-level logical closeness (Zhang et al., 2026a), Structured Reasoning rewards within-graph flow and cross-sample overlap (Dong et al., 2026a), Structural CoT trains executable computation graphs (Anonymous, 2026b), the Graph Reasoning Paradigm emits labeled symbolic graphs (Liu et al., 2026a), and GoT-R1 reinforces parsability, sparsity, and non-linearity of serialized reasoning graphs with an additive reward (Li et al., 2026b). AtomGraph instead learns a graph-attention verifier over reasoning chains that have already been parsed into DAGs (Karmore, 2026). Contract-checked graph editing enforces acyclicity as a hard gate for composing candidate DAGs during inference-time search (Li & Cao, 2026). RewardFlow propagates terminal reward across an agent state graph (Feng et al., 2026), while FlowTracer derives token credit from an attention-induced answer-flow DAG (Dong et al., 2026b); Graph-GRPO propagates credit over a predefined e-commerce relevance graph (Che et al., 2026). TopoPRM instead leaves policy traces in free text and recovers an unconstrained semantic-support graph before projection. Declared-parent, attention-DAG, learned-DAG, valid-DAG, and hard-gated representations rule out or do not score backward and cyclic semantic support; we retain these residuals as direction and circular-justification signals.

Post-hoc graph probes support non-sequential reasoning (Zhong et al., 2026; Matsutani et al., 2026); CoT2Graph uses recovered graphs for failure analysis, outcome prediction, and Best-of- N selection (Anonymous, 2026a), and Reasoning Consensus merges several recovered DAGs for inference-time structural ensembling (Parulekar et al., 2026). SARL clusters reasoning steps into latent functions and rewards clustering and short paths in their *undirected* transition graph, explicitly abstracting away directional control flow (Wang et al., 2026b). This small-world prior is complementary to our semantic-support relation: our direction score detects backward dependency mass, while acyclicity detects global self-support. Hidden-state cycles can reflect useful recurrent visits (Minegishi et al., 2025), while LoopBench studies self-reinforcing textual repetition during generation (Duan et al., 2026). Our cycles instead denote circular semantic support in the committed derivation: we do not penalize exploration, revision, or repeated wording, only a final support relation in which a




Figure 2: Overview of TopoPRM and Topology-Guided Distillation. An order-agnostic encoder recovers a raw directed support graph; its acyclic projection provides the training DAG while discarded backward and cyclic edges remain reward signals. A topology-conditioned Choquet reward aggregates multiple sources, and the projected graph guides dependency-preserving distillation.

claim depends on itself. The novelty is therefore not “reasoning as a graph,” but raw direction/cycle diagnostics as online, auditable reward sources without graph-formatted generation.

Distillation and Reasoning Compression. Knowledge distillation transfers teacher capability (Hinton et al., 2015); student-generated variants reduce train–test mismatch (Gu et al., 2024; Agarwal et al., 2024; Ko et al., 2024), including ExOPD, BOND, KDRL, and black-box GAD (Yang et al., 2026; Sessa et al., 2024; Xu et al., 2025; Ye et al., 2025). Long-CoT compression uses pruning, length penalties, or budget control (Xia et al., 2025; Luo et al., 2025; Arora & Zanette, 2026; Aggarwal & Welleck, 2025; Muennighoff et al., 2025); graph-based pruning removes weak branches from projected DAGs (Yuan et al., 2026), while Reasoning Scaffolding distills semantic operations (Wen et al., 2026). TopoPRM instead uses raw direction/cycle violations to elicit and filter teacher revisions, preserving dependencies that token budgets or post-projection pruning cannot identify. *Topology-Guided Distillation is evaluated by the accuracy–length frontier and independently audited topology preservation.*

3 METHOD

TopoPRM (Figure 2) is a topology-aware post-training framework whose stages share a graph encoder $\mathcal{E}_\phi$ and an acyclic projection Π . Stage I produces parsable reasoning traces. Stage II applies GRPO to a topology-conditioned multi-source reward. Stage III performs Topology-Guided Distillation, in which the Stage-II teacher revises student rollouts using defects localized on the projected graph.

3.1 PROBLEM SETUP

Let $x = (q, c)$ pair a mathematical problem q with optional context c , and let π_θ generate a trace $y = (s_1, \dots, s_T)$. The frozen encoder produces a raw directed support graph $\mathcal{H}_y = \mathcal{E}_\phi(y) = (\mathcal{V}, \mathcal{A}_H)$, where $i \rightarrow j$ means that step s_i supplies a result, claim, sub-goal, or assumption used by s_j . Unlike a DAG constructed only from pairs $i < j$, $\mathcal{H}_y$ may contain backward edges and cycles. We retain these violations and separately form $\mathcal{G}_y = \Pi(\mathcal{H}_y)$ by keeping support edges consistent with the trace order. Thus $\mathcal{G}_y$ supports DAG operations, while the projection cost of $\mathcal{H}_y$ exposes structural failures that would otherwise be removed by construction.

Scope of the acyclicity assumption. Acyclicity is normative only for the *semantic support relation in the committed trace*. It does not forbid exploratory computation, checking an earlier result, or revising a failed branch. Such behavior may revisit a latent state or repeat a concept while still

yielding an acyclic final justification. A cycle in $\mathcal{H}_y$ has the narrower meaning that a written claim ultimately supports itself; q_{acyc} targets this circular justification rather than recurrent reasoning in general.

3.2 TOPOLOGY-AWARE DATA PROCESSING

Raw-Graph Extraction and Projection. Given a rollout y , $\mathcal{E}_\phi$ proceeds in three stages. **(i) Segmentation and parsing.** It splits y into steps and sub-questions, then extracts normalized expressions, claim keys, variables, and step types. **(ii) Order-agnostic support encoding.** Auditable rules propose high-precision symbolic candidates, while a lightweight pair encoder predicts one of $\{i \rightarrow j, j \rightarrow i, \emptyset\}$ for each unordered intra-component pair. The encoder is trained offline from semantic-teacher edges and frozen during policy optimization; a separately annotated set calibrates edge precision and direction. Retained encoder edges carry their calibrated confidence, while rule edges use unit weight. Semantic predictions override an order-based rule when both judge the same pair, preventing artificial two-cycles caused by conflicting extractors. **(iii) Acyclic projection.** We preserve all predicted edges in $\mathcal{H}_y$, set $\mathcal{G}_y = \Pi(\mathcal{H}_y)$ by removing $i \rightarrow j$ with $i \geq j$, and add an adaptive weak sequential edge only when a node remains unsupported. Full rules and abstention thresholds are in Appendix C.

Topology and Continuity Diagnostics. Let $\mathcal{A}_H^{\text{dep}}$ denote the dependency edges of $\mathcal{H}_y$, let $w_{ij} \geq 0$ be the confidence weight of edge (i, j) , and set $W_H = \sum_{(i,j) \in \mathcal{A}_H^{\text{dep}}} w_{ij}$. Let $\mathcal{A}_{\text{cyc}}$ contain the dependency edges whose endpoints lie in the same cyclic strongly connected component (size greater than one, or a self-loop), and let c_H be the fraction of nodes incident to at least one dependency edge. We first compute d_H , the forward-edge mass, and u_H , the cycle-edge mass, then shrink both diagnostics toward 0.5 when edge coverage is low:

$$q_{\text{dir}}(\mathcal{H}_y) = c_H d_H + (1 - c_H)/2, \quad d_H = \frac{\sum_{(i,j) \in \mathcal{A}_H^{\text{dep}}} w_{ij} \mathbb{I}[i < j]}{W_H}, \quad (1)$$

$$q_{\text{acyc}}(\mathcal{H}_y) = c_H(1 - u_H) + (1 - c_H)/2, \quad u_H = \frac{\sum_{(i,j) \in \mathcal{A}_{\text{cyc}}} w_{ij}}{W_H}. \quad (2)$$

When $W_H = 0$, we use $d_H = 1$ and $u_H = 0$; because $c_H = 0$ for a multi-step empty graph, both coverage-adjusted scores are nevertheless 0.5, not perfect. (A singleton trace has no possible cross-step violation and uses $c_H = 1$.) The topology score q_{topo} is a normalized combination of q_{dir} , q_{acyc} , non-orphan support on $\mathcal{G}_y$, and optional reference-edge agreement. The continuity score q_{cont} measures whether each step is traceable to the question or a supported prefix. Direction and acyclicity are therefore properties of the raw graph, not tautological checks on an already projected DAG. They are related but not redundant: every directed cycle contains a backward edge under the textual order, whereas a backward edge need not close a cycle, and raw graphs with identical backward-edge mass can have different cycle-edge mass. Thus q_{dir} detects local order inconsistency while q_{acyc} detects global circular support.

Topology-Conditioned Multi-Source Aggregation. Let $\mathbf{z} = (r_{\text{out}}, r_{\text{fmt}}, q_{\text{topo}}, q_{\text{cont}}, r_{\text{len}}) \in [0, 1]^5$. A weighted sum assumes that each source has the same marginal value regardless of the others. Instead, we use a monotone 2-additive Choquet model with Möbius coefficients

$$r_{\text{total}}(x, y) = \sum_i m_i z_i + \sum_{(i,j) \in \mathcal{I}_{\text{topo}}} m_{ij} \min(z_i, z_j), \quad \mathcal{I}_{\text{topo}} = \{(o, t), (t, c), (t, l)\}. \quad (3)$$

All $m_i, m_{ij} \geq 0$ and sum to one, so the reward is bounded and monotone. The three interaction terms assign additional value only when topology agrees with outcome, continuity, or compactness. This makes topology a constituent of cross-source utility rather than a generic bonus applied after correctness. We ablate the capacity against equal-weight, tuned additive, and multiplicative alternatives.

3.3 TOPOLOGY-GUIDED POST-TRAINING

Stage I: Supervised Warm-Start. We perform Supervised Fine-Tuning (SFT) on structured mathematical reasoning data to teach the policy to produce stepwise solutions parsable by $\mathcal{E}_\phi$. This

reduces invalid rollouts during subsequent reinforcement learning, since topology diagnostics are meaningful only when the trace contains recoverable reasoning steps.

Stage II: GRPO with TopoPRM Rewards. Starting from the warm-start checkpoint, we optimise the policy with GRPO. For each completion, $\mathcal{E}_\phi$ constructs $\mathcal{H}_y$ and $\mathcal{G}_y$; Eq. 3 then scores the completion. Standard group normalization acts on the complete Choquet reward, so topology can change the ordering of rollouts rather than only break ties inside an outcome stratum. Reward-hacking tests separately measure whether structurally plausible but answer-incorrect traces are over-ranked.

Stage III: Topology-Guided Distillation. The optimized model after Stage II serves as the teacher π_T , while a smaller pretrained checkpoint initializes the student π_S . Topology-Guided Distillation conditions target construction on topology diagnostics from student-generated traces. For each prompt x , the student first rolls out $y_{\text{init}} \sim \pi_S(\cdot | x)$; $\mathcal{E}_\phi$ produces $\mathcal{H}_{y_{\text{init}}}$ and $\mathcal{G}_{y_{\text{init}}}$; a dispatcher ψ then issues a revision instruction targeting a backward edge, cyclic component, orphan conclusion, or continuity break. The teacher generates a revised target under this instruction. Thus topology determines where supervision is requested and which revisions are retained, rather than serving only as a post-hoc score.

Distillation objective. Let $A(y_{\text{init}}, y_{\text{rev}})$ indicate that the revised trace has a correct boxed answer and valid format, satisfies $q_{\text{topo}}^{\text{rev}} \geq \tau_g$, does not reduce either q_{dir} or q_{acyc} relative to the initial trace, and fits a fixed token budget B . The retained dataset is $\mathcal{D}_{\text{TGD}} = \{(x, y_{\text{rev}}) : A(y_{\text{init}}, y_{\text{rev}}) = 1\}$. The student minimises standard teacher-forced distillation loss

$$\mathcal{L}_{\text{distill}}(\theta) = -\mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{TGD}}} \sum_{t=1}^{|y|} \log \pi_S(y_t | x, y_{<t}). \quad (4)$$

Unlike static trace distillation, the target distribution is concentrated on revisions elicited by the current student’s structural failures. The explicit direction and acyclicity constraints prevent apparent compression from being obtained by deleting necessary support or introducing circular justification. The full pipeline is summarised in Appendix Algorithm 1.

4 EXPERIMENTS

Draft audit status. Red cells are not paper claims. Each method uses one canonical seed-0 training run. Parenthesized uncertainty will be a paired per-example bootstrap interval over the same evaluation items; it will not be simulated or described as training variance.

4.1 CANONICAL PROTOCOL

Fixed model scale and budget. The primary comparison starts from the same Qwen3.5-9B Stage-I checkpoint and changes only the Stage-II reward. The public frontier is restricted to dense released checkpoints with 7B–9B parameters; sparse models with similar total but much smaller active parameter counts are outside this matched scale. All variants use rank-64 LoRA, 300 GRPO steps, four generations per prompt, learning rate 5×10^{-6} , $\beta=0.04$, and a 4,096-token completion budget. Stage I uses 18,487 structured math solutions for training and 983 source-disjoint solutions for validation, three epochs, learning rate 5×10^{-5} , and global batch 32. Stage II uses 14,893 unique public train-split questions after exact normalized-question exclusion against every primary evaluation set. We remove 76 of 14,969 leakage-free candidates whose exact Qwen3.5 chat prompts exceed 512 tokens, rather than truncating their problem statements. Reference-solution DAGs are withheld because their step indices are not aligned with generated traces; q_{topo} is computed only from each rollout’s $\mathcal{H}_y$ and $\mathcal{G}_y$. The pair encoder is frozen in every topology-aware run; the 9B semantic teacher is used only to create its offline supervision and never appears in the GRPO reward loop.

Benchmarks and metrics. The primary fixed-scale frontier uses GSM8K, MATH-500, the official 675-item English text-only mathematics subset of OlympiadBench, and AIME 2024, spanning elementary through competition mathematics. Pass@1 exact-answer accuracy and mean visible response tokens are co-primary: responses are retokenized with the evaluated model’s tokenizer, with-

Table 1: Fixed-budget accuracy–length comparison at dense 7B–9B scale. Each cell reports pass@1 accuracy (%) and mean visible response tokens, retokenized without special tokens, from one seed-0 sampled response per item under the same sampling parameters, 4,096-token budget, and task content. Released weights, chat templates, and required thinking controls are retained; each adapted method uses one canonical seed-0 training run. **Pending intervals will come from per-example bootstrap resampling, not training variance.**

Method	GSM8K		MATH-500		OlympiadBench		AIME 2024		Macro avg.	
	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.
<i>Public checkpoints, reproduced with the same evaluator</i>										
Qwen3-8B (Qwen Team, 2025)	—	—	—	—	—	—	23.3	4,096	—	—
GoT-R1-8B (Li et al., 2026b)	—	—	—	—	—	—	—	—	—	—
DeepSeek-R1-0528-Qwen3-8B (DeepSeek-AI, 2025)	—	—	—	—	—	—	—	—	—	—
MiMo-7B-RL-0530 (Xiaomi LLM Core Team, 2025)	—	—	—	—	—	—	—	—	—	—
Nemotron-Cascade-8B-Thinking (Wang et al., 2025a)	—	—	—	—	—	—	—	—	—	—
Klear-Reasoner-8B (Su et al., 2025)	—	—	—	—	—	—	—	—	—	—
OpenReasoning-Nemotron-7B (NVIDIA, 2025)	—	—	—	—	—	—	0.0	4,096	—	—
AceMath-RL-Nemotron-7B (NVIDIA, 2026)	—	—	—	—	—	—	—	—	—	—
Qwen3.5-9B (Qwen Team, 2026)	—	—	—	—	—	—	0.0	4,096	—	—
<i>Matched Qwen3.5-9B adaptation from the same Stage-I checkpoint</i>										
Stage-I SFT	85.8	137	56.4	283	16.6	604	3.3	1,080	40.5	526
Outcome-only GRPO	—	—	—	—	—	—	—	—	—	—
Outcome + format + length GRPO	—	—	—	—	—	—	—	—	—	—
LASER-D (matched) (Liu et al., 2026b)	—	—	—	—	—	—	—	—	—	—
Matched additive reward	—	—	—	—	—	—	—	—	—	—
Matched multiplicative reward	—	—	—	—	—	—	—	—	—	—
TopoPRM (Choquet)	—	—	—	—	—	—	—	—	—	—

out special tokens, so EOS padding introduced by batched generation is excluded. Method *A* Pareto-dominates *B* only when *A* is no less accurate and no longer, with one strict inequality. Pass@5 and maj@5 are secondary diagnostics. Every reproduced row uses one seed-0 sampled response per item with temperature 0.6, top- $p=0.95$, top- $k=20$, min- $p=0$, repetition penalty 1.0, identical task content and verifier, and a 4,096-token response cap. We retain each released chat template and only its documented role or thinking control: GoT-R1 and the two Nemotron math checkpoints receive a user-only task; MiMo uses an empty system turn; Nemotron-Cascade appends “/think” to the user task; Qwen3-8B, DeepSeek-R1-0528-Qwen3-8B, Klear-Reasoner-8B, and Qwen3.5-9B retain their released thinking templates. This avoids conflating prompt compatibility with a model-specific best-case decoder. Matched rows share tokenizer and message roles.

Statistics and provenance. Every reported row maps to a checkpoint, evaluation manifest, and per-example output file. Accuracy differences and length differences use paired bootstrap resampling over shared items. Because training is run once per method, these intervals measure evaluation-set uncertainty only. Several public-frontier checkpoints were released with substantially longer reasoning budgets and repeated decoding. We therefore re-evaluate their released weights under the same single-response 4,096-token protocol rather than transferring quoted scores into Table 1. The human edge audit reports annotator agreement, edge precision/recall, direction accuracy, and abstention coverage separately from teacher-held-out classifier validation.

4.2 ACCURACY–LENGTH FRONTIER (RQ1)

Table 1 tests the central empirical claim: at a fixed 9B adaptation scale and response budget, does the raw-directed TopoPRM reward move the accuracy–length frontier beyond Stage-I SFT, outcome-only GRPO, explicit outcome–format–length shaping, the dynamic difficulty-aware LASER-D baseline, and matched additive or multiplicative multi-source rewards? Our matched LASER-D row follows the released mechanism: training-time length coverage determines a target for each rollout-difficulty band, and only correct responses within that target receive its step reward (Liu et al., 2026b). The completed Stage-I row has macro accuracy 40.5% (stratified item-bootstrap 95% interval [38.7, 42.8]) and 526 visible tokens ([423, 647]); these intervals quantify evaluation-set uncertainty, not training variance. A SOTA claim will be made only if TopoPRM is non-dominated

Table 2: Direction and acyclicity mechanism tests. **Top:** matched policy ablations use the canonical decoder and training budget; accuracy and mean visible tokens are co-primary. **Bottom:** frozen-encoder falsification tests perturb the same 120 source traces (58 matched pairs for cycle closure). Deltas are perturbed minus source; each row uses one deterministic seed-0 perturbation, with uncertainty obtained by paired source-trace bootstrap. This bottom panel is fully measured and is not a policy result. **Red cells require canonical policy outputs.**

Policy variant	GSM8K		MATH-500		AIME 2024		Macro		Generated topology	
	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.	q_{dir}	q_{acyc}
Full TopoPRM	-	-	-	-	-	-	-	-	-	-
w/o direction source	-	-	-	-	-	-	-	-	-	-
w/o acyclicity source	-	-	-	-	-	-	-	-	-	-
w/o direction and acyclicity	-	-	-	-	-	-	-	-	-	-
Projected DAG only	-	-	-	-	-	-	-	-	-	-
Rules only (no pair encoder)	-	-	-	-	-	-	-	-	-	-
w/o continuity	-	-	-	-	-	-	-	-	-	-
<i>Frozen-encoder falsification tests on identical source traces</i>										
Trace perturbation	n	p_{real}	Δq_{dir}	Δq_{acyc}	$\Delta \text{Cov.}$	$\Delta \text{Back.}$	ΔCycle	Targeted failure mode		
Reverse all steps	120	-	-.748	+.022	+.035	+.851	-.005	Textual order / direction		
Matched cycle closure	58	-	.000	-.342	.000	.000	> 0	Circular support at fixed backward mass		
Random edge deletion ($p = .10$)	120	.095	-.028	-.025	-.057	-	-.005	Missing dependencies		
Random direction flip ($p = .30$)	120	.305	-.216	-.076	.000	-	+.086	Orientation corruption		
Random false-edge addition ($p = .30$)	120	.258	-.072	-.171	+.037	-	+.195	Spurious and cyclic support		

Table 3: Matched aggregation and nearest graph-reward controls. All aggregation rows use the same five normalized sources and policy compute; only the aggregation operator changes. Structural columns are measured on generated traces from each policy, while Pareto wins count benchmarks on which a row is non-dominated in accuracy and visible length. The SARL-style control replaces directed semantic support with an undirected small-world score. **Cells await canonical outputs.**

Aggregation / graph control	Structural interaction	Macro Acc.	Avg. Tok.	Pareto wins $\uparrow$	$q_{\text{dir}} \uparrow$	$q_{\text{acyc}} \uparrow$	Hack% $\downarrow$	Zero-std $\downarrow$	ms/trace $\downarrow$
Equal-weight sum	None	-	-	-	-	-	-	-	-
Tuned additive	None	-	-	-	-	-	-	-	-
Verifier-stratified (HERO-style) (Tao et al., 2026)	Correct/incorrect strata	-	-	-	-	-	-	-	-
Multiplicative	All-source product	-	-	-	-	-	-	-	-
Choquet w/o topology interactions	Non-topology pairs only	-	-	-	-	-	-	-	-
Topology-conditioned Choquet	Direction/acyclicity pairs	-	-	-	-	-	-	-	-
SARL-style graph reward	Undirected transitions	-	-	-	-	-	-	-	-

against all matched rows and improves the reproduced accuracy-length frontier over the listed 7B-9B public checkpoints. Until then, the table remains a preregistered comparison rather than a result narrative. The closest training-time graph baseline, FlowTracer, reports 13.0, 79.7, and 46.7 accuracy on AIME 2024, MATH-500, and OlympiadBench with Qwen3-8B-Base under its separate 1K-output protocol (Dong et al., 2026b). The explicit-output Graph Reasoning Paradigm reports 91.2 on MATH-500 and 46.7 on AIME 2024, with mean reasoning lengths 1,260 and 7,040, respectively (Liu et al., 2026a). Its policy emits labeled graph nodes and the latter length exceeds our 4,096-token cap. Neither work releases a trained checkpoint with all four primary measurements, so their values define external accuracy-length bars but are not mixed into the reproduced Pareto table.

4.3 WHERE THE GAIN COMES FROM (RQ2)

Raw direction and acyclicity. Table 2 separates our central structural hypothesis from generic reward shaping at two levels: causal policy ablations test downstream utility, while measured counterfactuals test whether the frozen encoder detects the intended failure at all. On 120 traces, reversing the same steps reduces q_{dir} by 0.748 (paired bootstrap 95% interval [0.696, 0.798]), while q_{acyc} changes by only -0.022 because reversal need not create circular support. Conversely, in 58 graph pairs that hold edge count, backward mass, and incident coverage fixed, changing only whether the backward edge closes a cycle produces a 0.342 acyclicity gap ([0.299, 0.385]) and exactly zero direction-score difference. Under controlled encoder noise, 9.5% random edge deletion changes either score by only about 0.03, whereas 30.5% direction flips reduce q_{dir} by 0.216 and 25.8% false-edge addition reduces q_{acyc} by 0.171 (Appendix Table 7). Thus direction detects textual-order

Table 4: Independent human edge audit on direction-balanced traces. Every extractor is evaluated against the same strict-majority reference from three blinded copies. Precision, recall, F1, direction accuracy, cycle-detection F1, and non-abstention coverage are overall metrics; the final three columns expose degradation with trace length. Agreement (κ) is computed before reconciliation. **Cells await complete annotation.**

Extractor	Overall edge recovery			Topology validity			Edge F1 by trace length			Agreement
	Prec.	Rec.	F1	Dir.	Cyc.	Cov.	Short	Medium	Long	κ
Adjacent-step chain	–	–	–	–	–	1.00	–	–	–	–
Rules only	–	–	–	–	–	–	–	–	–	–
Pair encoder (no abstention)	–	–	–	–	–	1.00	–	–	–	–
Frozen pair encoder (calibrated)	–	–	–	–	–	–	–	–	–	–
Semantic teacher (audit upper bound)	–	–	–	–	–	–	–	–	–	–
Projected DAG only	–	–	–	–	–	–	–	–	–	–
Hybrid raw-graph encoder	–	–	–	–	–	–	–	–	–	–

inconsistency while acyclicity adds information about global self-support; the policy panel tests whether both signals translate into accuracy–length gains.

Non-additive aggregation. Table 3 holds the five sources and compute fixed within the aggregation block while changing only their combination; the verifier-stratified control follows HERO’s correctness-group normalization and variance-aware prompt weighting but uses our frozen dense diagnostics, avoiding an extra learned reward model (Tao et al., 2026). The final SARL-style control instead replaces directed semantic support with its undirected small-world score under matched policy and optimization budgets. The Choquet hypothesis is supported only if topology interactions improve the Pareto frontier or reduce reward hacking beyond tuned additive, verifier-stratified, and multiplicative controls. The nearest graph-reward control separately tests whether undirected transition structure is sufficient. Reward variance is diagnostic, not a standalone success metric.

4.4 INDEPENDENT STRUCTURAL AUDIT AND COST (RQ3)

Three blinded copies balance order variants and yield pairwise agreement plus a strict-majority reference. The online trace audit (Appendix Figure 3) exposes per-trace graphs, direction and cycle failures, score calculations, and direct links to each annotation copy; Table 4 reports human edge, direction, cycle, and coverage metrics across controls. A policy audit tests whether scores predict human edges or over-rank incorrect plausible traces, keeping the extractor from serving as its own evaluator.

The cost report separates offline semantic labeling and pair-encoder adaptation from online reward computation. For Stage I and Stage II we report wall-clock time, peak per-GPU memory, graph pairs per second, and the fraction of GRPO time spent in graph construction. Source data and large offline artifacts reside on shared object storage. Canonical Stage II checkpoints are saved atomically on the training filesystem, validated as readable safetensors, and only then archived to object storage; logs record the exact $4 \times L20$ allocation.

4.5 TOPOLOGY-GUIDED DISTILLATION (RQ4)

Table 5 asks whether Topology-Guided Distillation moves the 4B accuracy–length frontier beyond static teacher-trace distillation, generic teacher revision, and explicit length-only revision. All methods share the same Stage-II 9B teacher, student initialization, candidate prompts, accepted-target count, response budget, and optimizer budget. Direction and acyclicity retention must survive the independent edge audit rather than only the training extractor. Canonical values will replace the red pending cells after Stage II and the single 4B distillation run complete.

Table 5: Matched 9B-to-4B distillation protocol. All student methods use the same teacher, candidate prompts, accepted-target count, token budget, and optimizer budget. Accuracy is pass@1; Dir. and Acyc. are paired with the independent edge audit. **Cells remain pending until the single canonical run; intervals will be paired per-example bootstrap intervals, not training variance.**

Method	GSM8K	MATH-500	Avg. Tok	Dir.	Acyc.	Accept%
9B teacher	—	—	—	—	—	—
4B initialization	—	—	—	—	—	—
Static teacher traces	—	—	—	—	—	—
Generic revision	—	—	—	—	—	—
Length-only revision	—	—	—	—	—	—
Topology-Guided Distillation	—	—	—	—	—	—

5 CONCLUSION

We presented TopoPRM, a topology-aware process reward model that recovers a raw directed support graph before acyclic projection. The discarded backward and cyclic structure supplies direction and acyclicity supervision; a topology-conditioned Choquet reward couples these diagnostics with outcome, continuity, format, and length; and Topology-Guided Distillation uses the same violations to elicit and filter compact teacher revisions. The policy retains free-form reasoning and requires no graph-format decoding.

LIMITATIONS

TopoPRM provides a structural process signal rather than a proof-level verifier. Consequently, locally plausible yet semantically wrong steps can still receive positive structural credit, especially on competition problems where a single invalid deduction decides the outcome. The raw-graph diagnostics also depend on edge-recovery quality. We therefore report edge precision, direction accuracy, inter-annotator agreement, and abstention coverage rather than treating the encoder as an oracle. Future work can expand the verified-edge set and add uncertainty-aware abstention, couple structural credit with a semantic process verifier, and extend TopoPRM to code, scientific, and multimodal reasoning where executable or simulator-based oracles are available.

REPRODUCIBILITY STATEMENT

This draft is being audited for implementation-to-equation consistency and row-level result provenance. The submission version will point to the exact method definitions, configurations, raw-result manifests, and anonymized code bundle used to reproduce every retained result.

AI USE STATEMENT

Generative AI tools assisted with prior-review analysis, literature search and synthesis, hypothesis and experimental-control refinement, method implementation and code review, dataset cleaning, reproducibility auditing, scientific-figure editing, and manuscript drafting and revision. Before submission, the authors will manually inspect the cited literature, review code and generated artifacts, and verify every reported result against raw outputs. The authors take responsibility for the final claims, code, artifacts, and results; this disclosure will be updated if additional uses occur before submission.

REFERENCES

Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated

mistakes. In *International Conference on Learning Representations*, 2024.

Pranjal Aggarwal and Sean Welleck. L1: Controlling how long a reasoning model thinks with reinforcement learning. *arXiv preprint arXiv:2503.04697*, 2025.

Arash Ahmadian, Chris Cremer, Matthias Gallé, Marzieh Fadaee, Julia Kreutzer, Olivier Pietquin, Ahmet Üstün, and Sara Hooker. Back to basics: Revisiting reinforce-style optimization for learning from human feedback in llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 12248–12267, 2024.

Anonymous. Understanding failures in LLM reasoning by learning structured representations of chain-of-thought. Anonymous ACL submission, 2026a. URL <https://openreview.net/forum?id=0XfuJjhaI5>.

Anonymous. Structural chain-of-thought: Reasoning with executable computational graphs. ACL Rolling Review March 2026 submission, 2026b. URL <https://openreview.net/forum?id=iv6poj9Y4n>.

Daman Arora and Andrea Zanette. Training language models to reason efficiently. *Advances in Neural Information Processing Systems*, 38:60770–60808, 2026.

Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Nyczyk, Marcin Copik, and Torsten Hoeffler. Graph of thoughts: Solving elaborate problems with large language models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(16):17682–17690, 2024.

Jiarui Che, Yifei Chen, Zhixing Tian, Chenyang Wang, and Ziguang Cheng. Graph-GRPO: Dependency-aware credit assignment for generative e-commerce search relevance. *arXiv preprint arXiv:2605.31003*, 2026. URL <https://arxiv.org/abs/2605.31003>.

Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang. Does reinforcement learning really incentivize reasoning capacity in llms beyond the base model? *Advances in Neural Information Processing Systems*, 38:57654–57689, 2026.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, et al. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025.

DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

Yubo Dong, Hehe Fan, Linchao Zhu, and Yi Yang. Structured reasoning for LLMs: A unified framework for efficiency and explainability. In *International Conference on Learning Representations*, pp. 106183–106215, 2026a. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/ad5b3f324b24c17cdc2f3712298c76bd-Paper-Conference.pdf.

Zhichen Dong, Yang Li, Yuhan Sun, Weixun Wang, Yijia Luo, Zinian Peng, Taiheng Ye, Chao Yang, Wenbo Su, Yu Cheng, Bo Zheng, and Junchi Yan. How does reasoning flow? tracing attention-induced information flow for targeted RL in LLMs. In *International Conference on Machine Learning*, 2026b. URL <https://arxiv.org/abs/2606.10646>.

Zenghao Duan, Liang Pang, Zihao Wei, Wenbin Duan, Yuxin Tian, Shicheng Xu, Jingcheng Deng, Zhiyi Yin, and Xueqi Cheng. Circular reasoning: Understanding self-reinforcing loops in large reasoning models. *arXiv preprint arXiv:2601.05693*, 2026. URL <https://arxiv.org/abs/2601.05693>.

Jiwei Fang, Bin Zhang, Changwei Wang, Jin Wan, and Zhiwei Xu. Graph of verification: Structured verification of llm reasoning with directed acyclic graphs. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 30665–30672, 2026. doi: 10.1609/aaai.v40i36.40322.

- Xiao Feng, Bo Han, Zhanke Zhou, Jiaqi Fan, Jiangchao Yao, Ka Ho Li, Dahai Yu, and Michael Kwok-Po Ng. Rewardflow: Topology-aware reward propagation on state graphs for agentic rl with large language models. *arXiv preprint arXiv:2603.18859*, 2026. URL <https://arxiv.org/abs/2603.18859>.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *International Conference on Machine Learning*, pp. 10835–10866. PMLR, 2023.
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *International Conference on Learning Representations*, 2024.
- Xinyu Guan, Li Lyna Zhang, Yifei Liu, Ning Shang, Youran Sun, Yi Zhu, Fan Yang, and Mao Yang. rstar-math: Small llms can master math reasoning with self-evolved deep thinking, 2025. URL <https://arxiv.org/abs/2501.04519>.
- Tao He, Rongchuan Mu, Lizi Liao, Yixin Cao, Li Yang, Yijia Luo, Weixun Wang, Ming Liu, and Bing Qin. Smarter not harder: Generative process evaluation with intrinsic-signal driving and ability-adaptive reward shaping. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/da842df0d1e35e15a521267acb62bbd0-Abstract-Conference.html.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Jian Hu. Reinforce++: A simple and efficient approach for aligning large language models. *arXiv e-prints*, pp. arXiv–2501, 2025.
- Aryan Karmore. Atomgraph: Reasoning isn’t linear, why should verification be? In *ICLR 2026 Workshop on LLM Reasoning*, 2026. URL <https://openreview.net/forum?id=dNsIZln2uv>.
- Amirhossein Kazemnejad, Milad Aghajohari, Eva Portelance, Alessandro Sordoni, Siva Reddy, Aaron Courville, and Nicolas Le Roux. VinePPO: Refining credit assignment in RL training of LLMs. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025.
- Zixuan Ke, Fangkai Jiao, Yifei Ming, Xuan-Phi Nguyen, Austin Xu, Do Xuan Long, Minzhi Li, Chengwei Qin, Peifeng Wang, Silvio Savarese, Caiming Xiong, and Shafiq Joty. A survey of frontiers in llm reasoning: Inference scaling, learning to reason, and agentic systems, 2026. URL <https://arxiv.org/abs/2504.09037>.
- Muhammad Khalifa, Rishabh Agarwal, Lajanugen Logeswaran, Jaekyeom Kim, Hao Peng, Moon-tae Lee, Honglak Lee, and Lu Wang. Process reward models that think. *arXiv preprint arXiv:2504.16828*, 2025.
- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. Distillm: Towards streamlined distillation for large language models. *arXiv preprint arXiv:2402.03898*, 2024.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35: 22199–22213, 2022.
- Xin Lai, Zhuotao Tian, Yukang Chen, Senqiao Yang, Xiangru Peng, and Jiaya Jia. Step-dpo: Step-wise preference optimization for long-chain reasoning of llms. *arXiv preprint arXiv:2406.18629*, 2024.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Xinxin Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm, Luca Soldaini, Noah A Smith, Yizhong Wang, Pradeep Dasigi, and Hannaneh Hajishirzi. Tulu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*, 2024.

- Gang Li, Yan Chen, Ming Lin, and Tianbao Yang. DRPO: Efficient reasoning via decoupled reward policy optimization. In *International Conference on Learning Representations*, 2026a. arXiv:2510.04474.
- Rui Li and Shuang Cao. From trajectories to graphs: Contract-checked editing for verifier-guided LLM reasoning. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pp. 43259–43306, 2026. doi: 10.18653/v1/2026.acl-long.2004. URL <https://aclanthology.org/2026.acl-long.2004/>.
- Zuchao Li, Qiwei Li, Yao Yao, Hai Zhao, Lefei Zhang, and Bo Du. GoT-R1: Internalizing graph-of-thought via structural reinforcement for high-density reasoning. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 7090–7104, 2026b. doi: 10.18653/v1/2026.findings-acl.352. URL <https://aclanthology.org/2026.findings-acl.352/>.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023.
- Mingjie Liu, Shizhe Diao, Ximing Lu, Jian Hu, Xin Dong, Yejin Choi, Jan Kautz, and Yi Dong. ProRL: Prolonged reinforcement learning expands reasoning boundaries in large language models. *Advances in Neural Information Processing Systems*, 2025.
- Runxuan Liu, Xianhao Ou, Xinyan Ma, Jiyuan Wang, Jiafeng Liang, Jiaqi Li, Tao He, Zheng Chu, Rongchuan Mu, Zekun Wang, et al. Graph reasoning paradigm: Structured and symbolic reasoning with topology-aware reinforcement learning for large language models. *arXiv preprint arXiv:2601.12995*, 2026a.
- Wei Liu, Ruochen Zhou, Yiyun Deng, Yuzhen Huang, Junteng Liu, Yuntian Deng, Yizhe Zhang, and Junxian He. Learn to reason efficiently with adaptive length-based reward shaping. In *International Conference on Learning Representations*, 2026b. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/47795c4ae2f7d07ea2fb0d11fa2c3c90-Abstract-Conference.html.
- Jianqiao Lu, Zhiyang Dou, Hongru Wang, Zeyu Cao, Jianbo Dai, Yingjia Wan, Yunlong Feng, and Zhijiang Guo. Autopsv: Automated process-supervised verifier. *Advances in Neural Information Processing Systems*, 37:79935–79962, 2024.
- Haotian Luo, Li Shen, Haiying He, Yibo Wang, Shiwei Liu, Wei Li, Naiqiang Tan, Xiaochun Cao, and Dacheng Tao. O1-pruner: Length-harmonizing fine-tuning for o1-like reasoning pruning. *arXiv preprint arXiv:2501.12570*, 2025.
- Liangchen Luo, Yinxiao Liu, Rosanne Liu, Samrat Phatale, Harsh Lara, Yunxuan Li, Lei Shu, Yun Zhu, Lei Meng, Jiao Sun, and Abhinav Rastogi. Improve mathematical reasoning in language models by automated process supervision. *arXiv preprint arXiv:2406.06592*, 2024.
- Kohsei Matsutani, Shota Takashiro, Gouki Minegishi, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo. RL squeezes, SFT expands: A comparative study of reasoning LLMs. In *International Conference on Learning Representations*, pp. 141482–141526, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/e52554a70e0df57a0bea11d1eca0c9b5-Paper-Conference.pdf.
- Gouki Minegishi, Hiroki Furuta, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo. Topology of reasoning: Understanding large reasoning models through reasoning graph properties. In *Advances in Neural Information Processing Systems*, volume 38, 2025. URL https://papers.neurips.cc/paper_files/paper/2025/hash/147521db4d049c75cf18d9ff96bf68f5-Abstract-Conference.html.
- Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori B Hashimoto. s1: Simple test-time scaling. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 20286–20332, 2025.

- Sagnik Mukherjee, Abhinav Chinta, Takyoung Kim, Tarun Anoop Sharma, and Dilek Hakkani-Tür. Premise-augmented reasoning chains improve error identification in math reasoning with LLMs. In *International Conference on Machine Learning*, 2025. URL <https://arxiv.org/abs/2502.02362>.
- NVIDIA. OpenReasoning-Nemotron-7B model card. Hugging Face model release, 2025. URL <https://huggingface.co/nvidia/OpenReasoning-Nemotron-7B>.
- NVIDIA. AceMath-RL-Nemotron-7B model card. Hugging Face model release, 2026. URL <https://huggingface.co/nvidia/AceMath-RL-Nemotron-7B>.
- Amruta Parulekar, Jinu Lee, Dilek Hakkani-Tür, and Hari Sundaram. Reasoning consensus: Structural ensembling of LLM reasoning via weighted DAG aggregation. *arXiv preprint arXiv:2607.27783*, 2026. URL <https://arxiv.org/abs/2607.27783>.
- Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Qwen Team. Qwen3.5-9B model card. Hugging Face model release, 2026. URL <https://huggingface.co/Qwen/Qwen3.5-9B>.
- Pier Giuseppe Sessa, Robert Dadashi, Léonard Hussenot, Johan Ferret, Nino Vieillard, Alexandre Ramé, Bobak Shahriari, Sarah Shariari, Sertan Perrin, Abe Friesen, Geoffrey Cideron, Sertan Girgin, Piotr Stanczyk, Andrea Michi, Danila Sinopalnikov, Sabela Ramos, Morgane Riviere, Joaquin Lin, Dylan Hadfield-Menell, Rémi Munos, Olivier Pietquin, and Olivier Bachem. BOND: Aligning LLMs with best-of-N distillation. *arXiv preprint arXiv:2407.14622*, 2024.
- Amrith Setlur, Chirag Nagpal, Adam Fisch, Xinyang Geng, Jacob Eisenstein, Rishabh Agarwal, Alekh Agarwal, Jonathan Berant, and Aviral Kumar. Rewarding progress: Scaling automated process verifiers for llm reasoning. In *International Conference on Learning Representations*, volume 2025, pp. 60808–60838, 2025.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y K Li, Y Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Shuaijie She, Junxiao Liu, Yifeng Liu, Jiajun Chen, Xin Huang, and Shujian Huang. R-PRM: Reasoning-driven process reward modeling. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025.
- Joar Skalse, Nikolaus Howe, Dmitrii Krashennnikov, and David Krueger. Defining and characterizing reward gaming. *Advances in Neural Information Processing Systems*, 35:9460–9471, 2022.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling llm test-time compute optimally can be more effective than scaling model parameters, 2024. URL <https://arxiv.org/abs/2408.03314>.
- Mingyang Song, Zhaochen Su, Xiaoye Qu, Jiawei Zhou, and Yu Cheng. PRMBench: A fine-grained and challenging benchmark for process-level reward models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*, 2025.
- Zhenpeng Su, Leiyu Pan, Xue Bai, Dening Liu, Guanting Dong, Jiaming Huang, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui Zhou. Klear-Reasoner: Advancing reasoning capability via gradient-preserving clipping policy optimization. *arXiv preprint arXiv:2508.07629*, 2025.
- Leitian Tao, Ilia Kulikov, Swarnadeep Saha, Tianlu Wang, Jing Xu, Sharon Li, Jason E. Weston, and Ping Yu. Hybrid reinforcement: When reward is sparse, better to be dense. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/dcda966df09d8771a3db592c5fe4c587-Abstract-Conference.html.
- Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. Solving math word problems with process-and outcome-based feedback. *arXiv preprint arXiv:2211.14275*, 2022.

- Boxin Wang, Chankyu Lee, Nayeon Lee, Sheng-Chieh Lin, Wenliang Dai, Yang Chen, Yangyi Chen, Zhuolin Yang, Zihan Liu, Mohammad Shoneybi, Bryan Catanzaro, and Wei Ping. Nemotron-cascade: Scaling cascaded reinforcement learning for general-purpose reasoning models. *arXiv preprint arXiv:2512.13607*, 2025a. URL <https://arxiv.org/abs/2512.13607>.
- Hongru Wang, Deng Cai, Wanjun Zhong, Shijue Huang, Jeff Z. Pan, Zeming Liu, and Kam-Fai Wong. Self-reasoning language models: Unfold hidden reasoning chains with few reasoning catalyst, 2025b. URL <https://arxiv.org/abs/2505.14116>.
- Peiyi Wang, Lei Li, Zhihong Shao, Runxin Xu, Damai Dai, Yifei Li, Deli Chen, Yu Wu, and Zhifang Sui. Math-shepherd: Verify and reinforce llms step-by-step without human annotations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 9426–9439, 2024.
- Xinpeng Wang, Nitish Joshi, Barbara Plank, Rico Angell, and He He. Is it thinking or cheating? detecting implicit reward hacking by measuring reasoning effort. In *International Conference on Learning Representations*, 2026a. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/9d0ee47d732bd526c4a03d4aecc39bc1-Abstract-Conference.html.
- Yifan Wang, Bolian Li, David Cho, Ruqi Zhang, Fanping Sui, and Ananth Grama. SARL: Label-free reinforcement learning by rewarding reasoning topology. *arXiv preprint arXiv:2603.27977*, 2026b. URL <https://arxiv.org/abs/2603.27977>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- Xiangyu Wen, Min Li, Junhua Huang, Jianyuan Zhong, Zhijian Xu, Zeju Li, Yongxiang Huang, Mingxuan Yuan, and Qiang Xu. Dependency matters: Enhancing LLM reasoning with explicit knowledge grounding. In *Advances in Neural Information Processing Systems*, volume 38, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/5401bede0941e568fa42b7b5910fa72c-Abstract-Conference.html.
- Xiangyu Wen, Junhua Huang, Zeju Li, Min Li, Jianyuan Zhong, Zhijian Xu, Mingxuan Yuan, Yongxiang Huang, and Qiang Xu. Reasoning scaffolding: Distilling the flow of thought from LLMs. In *International Conference on Learning Representations*, pp. 80433–80459, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/8248b1ded388fcdbbd121bcdfea3068c-Paper-Conference.pdf.
- Heming Xia, Chak Tou Leong, Wenjie Wang, Yongqi Li, and Wenjie Li. Tokenskip: Controllable chain-of-thought compression in llms. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 3351–3363, 2025.
- Xiaomi LLM Core Team. MiMo: Unlocking the reasoning potential of language model—from pretraining to posttraining. *arXiv preprint arXiv:2505.07608*, 2025. URL <https://arxiv.org/abs/2505.07608>.
- Hongling Xu, Qi Zhu, and Ruifeng Xu. KDRL: Post-training reasoning LLMs via unified knowledge distillation and reinforcement learning. *arXiv preprint arXiv:2506.02208*, 2025.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024a.
- An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, et al. Qwen2.5-math technical report: Toward mathematical expert model via self-improvement. *arXiv preprint arXiv:2409.12122*, 2024b.
- Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning beyond teacher: Generalized on-policy distillation with reward extrapolation. *arXiv preprint arXiv:2602.12125*, 2026.
- Tianzhu Ye, Li Dong, Zewen Chi, Xun Wu, Shaohan Huang, and Furu Wei. Black-box on-policy distillation of large language models. *arXiv preprint arXiv:2511.10643*, 2025.

- Edward Yeo, Yuxuan Tong, Xinyao Niu, Graham Neubig, and Xiang Yue. Demystifying long chain-of-thought reasoning in LLMs. *arXiv preprint arXiv:2502.03373*, 2025.
- Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Weinan Dai, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025.
- Hongyuan Yuan, Xinran He, Run Shao, Bolei He, Xianwei Xue, Mengke Chen, Qitong Pan, Haiwei Wang, and Haifeng Li. Graph-based chain-of-thought pruning for reducing redundant reflections in reasoning LLMs. *Findings of the Association for Computational Linguistics: ACL 2026*, 2026. URL <https://arxiv.org/abs/2604.05643>.
- Lifan Yuan, Wendi Li, Huayu Chen, Ganqu Cui, Ning Ding, Kaiyan Zhang, Bowen Zhou, Zhiyuan Liu, and Hao Peng. Free process rewards without process labels. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025.
- Weihao Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He. Simplerl-zoo: Investigating and taming zero reinforcement learning for open base models in the wild, 2025. URL <https://arxiv.org/abs/2503.18892>, 2023.
- Dan Zhang, Sining Zhoubian, Ziniu Hu, Yisong Yue, Yuxiao Dong, and Jie Tang. Rest-mcts*: Llm self-training via process reward guided tree search. *Advances in Neural Information Processing Systems*, 37:64735–64772, 2024.
- Yuanhe Zhang, Ilja Kuzborskij, Jason Lee, Chenlei Leng, and Fanghui Liu. DAG-Math: Graph-of-thought guided mathematical reasoning in LLMs. In *International Conference on Learning Representations*, pp. 74347–74389, 2026a. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/79179a9c80561608fa85f925cc88aaa9-Paper-Conference.pdf.
- Zhenru Zhang, Chujie Zheng, Yangzhen Wu, Beichen Zhang, Runji Lin, Bowen Yu, Dayiheng Liu, Jingren Zhou, and Junyang Lin. The lessons of developing process reward models in mathematical reasoning. *arXiv preprint arXiv:2501.07301*, 2025.
- Zijing Zhang, Ziyang Chen, Mingxiao Li, Zhaopeng Tu, and Xiaolong Li. RLVMR: Reinforcement learning with verifiable meta-reasoning rewards for robust long-horizon agents. In *International Conference on Learning Representations*, 2026b. URL <https://iclr.cc/virtual/2026/poster/10008499>.
- Jian Zhao, Runze Liu, Kaiyan Zhang, Zhimu Zhou, Junqi Gao, Dong Li, Jiafei Lyu, Zhouyi Qian, Biqing Qi, Xiu Li, and Bowen Zhou. GenPRM: Scaling test-time compute of process reward models via generative reasoning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2026.
- Chujie Zheng, Zhenru Zhang, Beichen Zhang, Runji Lin, Keming Lu, Bowen Yu, Dayiheng Liu, Jingren Zhou, and Junyang Lin. Processbench: Identifying process errors in mathematical reasoning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1009–1024, 2025.
- Tianjun Zhong, Linyang He, Ziyang Li, and Nima Mesgarani. From chains to dags: Probing the graph structure of reasoning in llms. *arXiv preprint arXiv:2601.17593*, 2026. URL <https://arxiv.org/abs/2601.17593>.
- Xinyu Zhu, Mengzhou Xia, Zhepei Wei, Wei-Lin Chen, Danqi Chen, and Yu Meng. The surprising effectiveness of negative reinforcement in LLM reasoning. *Advances in Neural Information Processing Systems*, 2025.

A IMPLEMENTATION DETAILS

A.1 HYPERPARAMETERS

Table 6 lists the key hyperparameters for each training stage.

Table 6: Hyperparameter summary for the three training stages.

Parameter	Value	Description
<i>Stage I: Supervised Fine-Tuning (SFT)</i>		
Base model	Qwen3.5-9B	–
Adapter	LoRA $r=64$, $\alpha=128$	all-linear modules
Learning rate	5×10^{-5}	cosine, 3% warmup
Epochs	3	–
Batch size	1/device, accum 8	global batch 32 on 4 GPUs
Max seq length	4096	–
Precision	bf16	–
<i>Stage II: GRPO with TopoPRM Rewards</i>		
Base model	SFT checkpoint	–
Adapter	LoRA $r=64$, $\alpha=128$	all-linear modules
Learning rate	5×10^{-6}	cosine, 5% warmup
Gradient steps	300	–
Batch size	1/device, accum 4	global prompt batch 16 on 4 GPUs
Num generations	4 per prompt	for GRPO grouping
Temperature	0.9	rollout sampling
β (KL penalty)	0.04	–
Max completion length	4096	prompt capped at 512
Precision	bf16	–
<i>Stage III: Topology-Guided Distillation</i>		
Teacher model	Stage-II 9B	frozen during distillation
Student model	Qwen3.5-4B	canonical cross-scale run
Adapter	LoRA $r=64$, $\alpha=128$	all-linear modules
Learning rate	2×10^{-5}	cosine, 3% warmup
Epochs	1	–
Batch size	2/device, accum 4	global batch 32 on 4 GPUs
Student temperature	0.7	student-sampled trace
Teacher temperature	0.0	greedy revision target
Target budget	1024 tokens	fixed across methods
Target selection	answer, format, direction, acyclicity	threshold calibrated on validation
Max seq length	4096	–
Precision	bf16	–

A.2 REWARD CONFIGURATION

The Choquet reward in Eq. 3 uses singleton coefficients $(m_o, m_f, m_t, m_c, m_l) = (0.40, 0.05, 0.12, 0.08, 0.03)$ and interactions $(m_{ot}, m_{tc}, m_{tl}) = (0.12, 0.13, 0.07)$. These non-negative coefficients sum to one. The additive and multiplicative alternatives use the same five normalized sources and are tuned under the same budget; the main method uses the complete Choquet reward directly in standard group normalization. In canonical Stage II runs, $\lambda_k=0$: no reference-solution graph is supplied because its nodes are not aligned with a generated trace. Reference-edge alignment remains an optional definition only for settings with explicit node correspondence.

A.3 HARDWARE AND INFRASTRUCTURE

Canonical ICLR runs use 4×NVIDIA L20 GPUs. We report measured wall-clock time and peak memory from this configuration rather than extrapolating from the earlier A100 runs.

A.4 EVALUATION PROTOCOL

All public benchmark evaluations use the Transformers generation backend. Pass@1 is one sampled response per item from a single seed-0 evaluation pass; this is neither a multi-seed training result nor an average over repeated decodes. The shared decoder uses temperature 0.6, top- $p=0.95$, top- $k=20$, min- $p=0$, and repetition penalty 1.0. This matches Qwen3 and DeepSeek-R1’s reasoning-time recommendations and keeps the public and adapted rows comparable; it differs from Qwen3.5’s general-task recommendation only in using temperature 0.6 instead of 1.0. The four primary mathematics benchmarks use the same 4,096-token completion cap and the same instruction to provide a step-by-step solution with the final answer in `\boxed{}`. Released chat templates are retained. For OpenReasoning-Nemotron-7B and AceMath-RL-Nemotron-7B, the instruction is folded into one user turn to match their released mathematics prompts; no model receives an uncounted response prefix. Answer extraction takes the last boxed expression, and `math_verify` checks symbolic equivalence. Pass@5 and maj@5, when run as secondary diagnostics, draw five samples from the same decoder and reuse this equivalence relation rather than raw string equality. The manifest stores the message-role protocol, prefix flag, decoder parameters, evaluation seed, batch size, model and adapter paths, dataset fingerprint, and full per-example responses; cached results are reused only when these protocol fields match.

B CHOQUET CAPACITY AND BASELINES

For non-negative Möbius coefficients summing to one, Eq. 3 is monotone in every source, maps 0 to 0 and 1 to 1, and remains in $[0, 1]$. The topology interaction mass $m_{ot} + m_{tc} + m_{tl}$ is separately ablated to distinguish the proposed non-additive mechanism from simply increasing topology’s singleton weight. We also test traces that are structurally plausible but answer-incorrect to detect reward hacking.

C DAG EXTRACTION DETAILS

Extraction pipeline. Given a reasoning trace in `<think>`, we apply five frozen stages:

1. Step segmentation: split by numbered-step markers, mathematical discourse markers (`‘:’, ‘.’`), and line boundaries.
2. Sub-question partitioning: markers such as “Sub-question k ” define connected components; cross-sub-question edges are disallowed.
3. Information extraction: for each step, extract expressions, claims, and variables with rule-based patterns.
4. Raw support encoding: combine high-precision symbolic candidates with a lightweight pair classifier that selects $i \rightarrow j$, $j \rightarrow i$, or abstention without imposing text order.
5. Projection: retain the raw graph for diagnostics, remove backward edges to obtain the training DAG, then add adaptive weak sequential edges only when support is absent.

Information-density filters. To avoid low-information nodes and noisy links: (a) claims are extracted at sentence level, with heading fragments and unfinished prefixes removed; (b) trivial tokens (single symbols, pure numerals, non-relational fragments) are excluded from dependency matching.

Dependency rules. For each unordered pair of steps in the same sub-question, rules propose candidates from informative expression reuse, canonicalized claim overlap, variable roles, and explicit step references. The pair classifier assigns direction or abstains. Only after raw-graph diagnostics are computed do we insert a fallback edge for an unsupported derivation or conclusion in the projected DAG.

Pair-encoder adaptation. The lightweight encoder is Qwen3.5-0.8B with a rank-8 LoRA three-way classification head. Offline supervision contains 540 usable source traces labeled by the frozen 9B semantic teacher. Each source trace is presented in original, reversed, and interleaved step order, while the train-validation split is performed before augmentation at the source-trace level. The encoder is frozen during policy optimization; its confidence and direction-margin abstention thresholds are calibrated on the held-out teacher split and audited separately against human edge labels. A retained encoder edge keeps its calibrated confidence as w_{ij} ; auditable rule edges use $w_{ij} = 1$.

Controlled edge noise. Table 7 measures how the two raw-graph diagnostics respond when frozen predictions are corrupted after extraction. Deletion probes missed edges, false-edge addition probes spurious support, and direction reversal probes orientation errors. The audit uses one deterministic perturbation draw per source trace and paired bootstrap across traces; it is not a repeated-training or multi-seed result.

Table 7: Frozen-encoder sensitivity on 120 source traces. One deterministic seed-0 perturbation is applied per trace; intervals are paired source-trace bootstrap 95% intervals. p_{real} is the realized modified-edge fraction. Deltas are corrupted minus original.

Perturbation	p	p_{real}	Δq_{dir}	Δq_{acyc}	$\Delta \text{Cov.}$	$\Delta \text{Cycle mass}$
Random edge deletion	.10	.095	-.028 [-.045, -.014]	-.025 [-.043, -.009]	-.057	-.005
	.20	.192	-.050 [-.072, -.030]	-.054 [-.076, -.035]	-.107	+.001
	.30	.260	-.049 [-.072, -.029]	-.056 [-.077, -.036]	-.121	-.005
Random false-edge addition	.10	.053	-.016 [-.025, -.007]	-.050 [-.075, -.027]	+.007	+.058
	.20	.134	-.032 [-.048, -.016]	-.098 [-.138, -.061]	+.027	+.116
	.30	.258	-.072 [-.095, -.050]	-.171 [-.220, -.125]	+.037	+.195
Random direction flip	.10	.061	-.047 [-.069, -.026]	-.023 [-.042, -.006]	.000	+.023
	.20	.180	-.128 [-.162, -.096]	-.040 [-.069, -.015]	.000	+.048
	.30	.305	-.216 [-.271, -.165]	-.076 [-.116, -.040]	.000	+.086

Adaptive sequential weak edges. A weak sequential edge (weight 0.3) is added only when the target step lacks strong incoming support and the source step does not already emit strong outgoing dependencies. This prevents chain-dominant artifacts in long traces.

Node typing. Each node is classified into DEFINITION, DERIVATION, COMPUTATION, CONCLUSION, AUXILIARY, SUBSTITUTION, or CASE_ANALYSIS by keyword rules.

Relationship between DAG quality and reasoning compression. We observe empirically that traces with higher DAG scores (acyclicity, low orphan ratio, strong directional consistency) tend to admit shorter compressed forms under Topology-Guided Distillation. Intuitively, a well-structured DAG has fewer redundant branches and shorter critical paths, so the transitive reduction yields a more compact chain. The topology source q_{topo} enters the Choquet reward in Eq. 3, while the projected DAG supplies the compression target.

D REWARD FUNCTION DETAILS

Outcome reward. Extracts the final mathematical answer from the last `\boxed{\}` expression, falling back to an explicit final numeric expression when needed. Symbolic equivalence to the reference answer is checked with `math_verify`; a correct answer receives 1 and an incorrect or unparseable answer receives 0.

Format reward. Scores recognizable mathematical reasoning and conclusion structure: a `<think>` block followed by a boxed answer receives 1.0; multi-step reasoning with a boxed answer receives 0.8; a boxed answer alone receives 0.5; a `<think>` block without a boxed conclusion receives 0.3; and unboxed step-by-step reasoning receives 0.1. Empty or unstructured output receives 0.




Figure 3: Trace-level topology audit interface. For each sampled reasoning trace, the dashboard exposes the recovered support graph, edge direction, coverage-adjusted direction and acyclicity, the auditable topology-score calculation, and links to three blinded human-annotation copies. Original, reversed, and interleaved controls switch the same trace online. This frozen-encoder sample contains 10 retained edges, including two backward and three cycle-participating edges; its values are not an aggregate evaluation result.

Length reward. Linear decay from 1.0 (at $\leq 2,000$ characters) to 0.0 (at $\geq 8,000$ characters). Evaluation tables report tokenizer-token counts independently of this training-time shaping signal.

Topology reward (r_{topo}). The topology source used in GRPO is computed jointly from the raw support graph and projected DAG (Eq. 5). It aggregates graph existence, coverage-adjusted acyclicity, absence of orphan conclusions, coverage-adjusted direction consistency, and optional reference-edge alignment.

Continuity reward. Uses regex-based expression matching to verify that each step references expressions or claims from prior steps or the problem statement. The continuity score q_{cont} (Eq. 6) measures the fraction of supported steps.

E SUPPLEMENTARY MATHEMATICAL DEFINITIONS

Topology score (full form). Let c_H be the fraction of nodes incident to a dependency edge in the raw graph, and let d_H and u_H be the confidence-weighted forward-edge and cycle-edge mass defined in Eqs. (1)–(2). Coverage-adjusted diagnostics treat missing edge evidence as uncertainty:

$$q_{topo}(\mathcal{H}, \mathcal{G}) = \sigma(\lambda_b \mathbb{1}[|\mathcal{V}| > 0] + \lambda_a [c_H(1 - u_H) + (1 - c_H)/2] + \lambda_o \mathbb{1}[\rho_{orphan}(\mathcal{G}) = 0] + \lambda_d [c_H d_H + (1 - c_H)/2] + \lambda_k \kappa(\mathcal{G}, \mathcal{G}^*)), \quad (5)$$

Algorithm 1 Topology-Guided Distillation Training Pipeline

Require: Base policy π_0 ; data $\mathcal{D}_{\text{sft}}, \mathcal{D}_{\text{rl}}$; Choquet coefficients $\{m_i, m_{ij}\}$

Ensure: Teacher π_T , compact student π_S

```

1:  $\pi_\theta \leftarrow \text{SFT}(\pi_0, \mathcal{D}_{\text{sft}})$  ▷ Stage I
2: for each GRPO step on  $\mathcal{D}_{\text{rl}}$  do ▷ Stage II
3:   Sample  $\{y^{(i)}\}_{i=1}^G \sim \pi_\theta(\cdot | x)$ ; build  $\mathcal{H}^{(i)} = \mathcal{E}_\phi(y^{(i)})$  and  $\mathcal{G}^{(i)} = \Pi(\mathcal{H}^{(i)})$ 
4:   Compute  $r_{\text{out}}^{(i)}, r_{\text{fmt}}^{(i)}, r_{\text{len}}^{(i)}, q_{\text{topo}}^{(i)}, q_{\text{cont}}^{(i)}$ 
5:   Compute  $r_{\text{total}}^{(i)}$  by Eq. 3; normalize within the GRPO group
6:   Update  $\pi'_\theta$  with the GRPO objective
7: end for
8:  $\pi_T \leftarrow \pi'_\theta$ ; init.  $\pi_S$  from a smaller pretrained checkpoint
9: for each prompt  $x \in \mathcal{D}_{\text{rl}}$  do ▷ Stage III
10:   $y_{\text{init}} \sim \pi_S(\cdot | x)$ ; extract defect record  $F(y_{\text{init}})$  from  $\mathcal{H}$  and  $\mathcal{G}$ 
11:   $P_r \leftarrow \psi(F(y_{\text{init}}))$ ; sample  $y_{\text{rev}} \sim \pi_T(\cdot | x, y_{\text{init}}, P_r)$ 
12:  Retain  $(x, y_{\text{rev}})$  only if answer, format, direction, acyclicity, and budget constraints pass
13:  Update  $\pi_S$  with  $\mathcal{L}_{\text{distill}}$  (Eq. 4)
14: end for
15: return  $(\pi_T, \pi_S)$ 

```

where $\sigma(\cdot)$ normalizes by the active weights, ρ_{orphan} is orphan-conclusion ratio on the projected DAG, and κ is optional reference-edge alignment. The terms in Eq. (5) are aggregated with weights $(\lambda_b, \lambda_a, \lambda_o, \lambda_d, \lambda_k) = (0.20, 0.15, 0.15, 0.15, 0.25)$. The normalizer σ divides by the sum of active weights; when no reference graph exists, λ_k is removed before normalization. This guarantees $q_{\text{topo}} \in [0, 1]$ while preventing an empty raw graph from receiving perfect direction and acyclicity credit.

Continuity score (full form). The local continuity quality is

$$q_{\text{cont}}(y) = \begin{cases} 1.0, & \eta(y) = 1, \\ 0.8 \eta(y), & \eta(y) < 1, \end{cases} \quad \eta(y) = \frac{1}{T} \sum_{i=1}^T \mathbb{1}[\text{supported}(s_i, s_{<i}, q)]. \quad (6)$$

Remark. Eq. (6) is the engineering-tuned continuity term used in the implementation. It is a separate source in the Choquet reward rather than being folded into topology before aggregation.

Topology-conditioned revision dispatch. The revision instruction is generated by

$$P_r = \psi(r_{\text{out}}, q_{\text{dir}}, q_{\text{acyc}}, \rho_{\text{orphan}}, q_{\text{cont}}, \ell), \quad (7)$$

where ℓ identifies the offending edge, cyclic component, or unsupported step when one can be localized.

F TOPOLOGY-GUIDED DISTILLATION DETAILS

Topology-Guided Distillation is an optional compression stage whose goal is to retain the teacher’s structural reasoning profile while reducing generation length in compact students.

Why static trace distillation is insufficient. Conventional trace distillation (Hinton et al., 2015; Gu et al., 2024) presents every student with the same teacher solutions. It therefore cannot target a student’s particular missing dependency, and a length-only target can become shorter by deleting support that is necessary for the conclusion. Topology-Guided Distillation instead constructs supervision after observing the compact student’s own trace.

Defect localization and revision. For each initial student attempt, we retain the raw edge list and projected DAG. The defect record contains the endpoints of backward edges, nodes in a non-trivial strongly connected component, orphan conclusions in the projected DAG, and unsupported local transitions. The dispatcher selects the highest-priority localized defect and asks the frozen Stage-II teacher to rewrite the solution within a fixed response budget while preserving a boxed final answer.

Target acceptance. A revised trace is retained only if it has a correct boxed answer and valid format, reaches the calibrated topology threshold, does not reduce coverage-adjusted direction consistency or acyclicity relative to the initial trace, and stays within the shared token budget. We log every rejection reason and report acceptance coverage. This makes direction and acyclicity part of the distillation data contract, not merely an evaluation diagnostic.

Student training and controls. The 4B student is trained on accepted revisions with Eq. 4. Matched controls use the same teacher, prompts, number of accepted targets, and optimization budget but replace topology-conditioned dispatch with either a generic revision instruction or an explicit length-only instruction. External answer accuracy and generated tokens are primary; extractor-based structural retention is reported as internal-consistency evidence and paired with the independent human edge audit.

G RETAINED EMNLP SUBMISSION RESULTS

Archival evidence, not current ICLR claims. The following tables are retained from the EMNLP submission so that prior measured work is not discarded. They mix earlier checkpoints, decoders, response budgets, and in some cases explicitly marked approximate entries. The ICLR main tables use only the canonical single-run protocol in Section 4.1; archival rows will be promoted only after row-level provenance and protocol revalidation.

Table 8: Per-benchmark results and component ablations of TOPOPRM. We report **p@1**, **p@5**, **maj@5**, mean generated tokens / total wall-clock seconds (**Tok/Time**, in ks), and token-efficiency **Acc/kTok** = $p@1 / (\text{tokens}/1000)$. We use the same protocol in A. **Bold** / underline: best / second-best per column.

Dataset / Method	GSM8K (1,319)					MATH-500 (500)					OlympiadBench (134)					Omni-MATH (262)				
	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok
Base (Qwen3.5-9B)	55.2	87.8	55.7	1536/45.0	35.9	19.8	39.2	21.4	1536/18.7	12.9	11.2	25.4	14.2	2560/26.5	4.4	16.4	33.2	17.9	2559/52.0	6.4
Outcome Only	93.3	97.7	93.6	1047/34.2	89.1	<u>50.8</u>	59.0	55.2	1511/18.9	33.6	31.3	46.3	37.3	<u>2043/19.7</u>	<u>15.3</u>	41.6	<u>55.3</u>	48.9	<u>1845/35.2</u>	<u>22.5</u>
w/o Topology	93.0	97.7	93.9	<u>946/25.0</u>	<u>98.3</u>	<u>51.0</u>	<u>59.4</u>	55.6	<u>1506/16.2</u>	<u>33.9</u>	29.8	<u>50.7</u>	<u>38.1</u>	2468/11.8	12.1	<u>42.4</u>	54.2	45.8	2430/17.8	17.4
w/o Continuity	51.0	86.1	52.6	1528/48.3	33.4	21.6	39.6	23.2	1536/19.1	14.1	8.2	25.4	9.0	2560/ 8.7	3.2	14.9	30.5	17.6	2560/18.5	5.8
w/o ACE	<u>93.8</u>	98.0	<u>94.6</u>	1006/28.3	93.2	<u>50.8</u>	<u>59.8</u>	<u>55.2</u>	1515/18.2	33.5	32.8	47.0	35.8	2491/ 9.6	13.2	42.0	53.8	47.3	2463/19.4	17.1
★ TopoPRM (Full)	94.1	<u>97.9</u>	94.8	796/19.7	118.2	51.0	81.9	54.6	1454/16.6	35.1	32.8	50.7	38.8	1768/18.1	18.6	42.8	55.7	<u>47.7</u>	1592/19.0	26.9

Dataset / Method	AIME 2024 (30)					AIME 2025 (90)					CNMO 2024 (30)					MMLU (500)				
	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok
Base (Qwen3.5-9B)	16.7	33.3	33.3	4083/ 8.5	4.1	2.2	4.4	1.1	2560/ 5.9	0.9	10.0	16.7	3.3	2560/ 1.9	3.9	28.6	41.4	26.2	512/55.1	55.9
Outcome Only	16.7	33.3	20.0	<u>2546/ 5.6</u>	6.6	<u>13.3</u>	15.6	8.9	<u>2547/16.7</u>	<u>5.2</u>	16.7	<u>33.3</u>	<u>26.7</u>	<u>2530/ 5.5</u>	6.6	63.0	84.0	61.0	512/146.2	123.0
w/o Topology	<u>20.0</u>	<u>40.0</u>	<u>30.0</u>	2560/ 2.1	<u>7.8</u>	12.2	<u>16.7</u>	10.0	2560/ 6.2	4.8	<u>20.0</u>	40.0	<u>26.7</u>	2560/ 2.0	<u>7.8</u>	<u>65.8</u>	<u>85.0</u>	<u>63.0</u>	512/32.3	<u>128.5</u>
w/o Continuity	6.7	13.3	3.3	2560/ 2.3	2.6	1.1	4.4	0.0	2560/ 7.0	0.4	6.7	20.0	3.3	2560/ 2.4	2.6	20.2	34.8	18.6	512/80.4	39.5
w/o ACE	<u>20.0</u>	30.0	23.3	2560/ 2.6	<u>7.8</u>	<u>13.3</u>	<u>17.8</u>	12.2	2560/ 7.7	<u>5.2</u>	<u>20.0</u>	30.0	<u>26.7</u>	2560/ 2.6	<u>7.8</u>	61.1	83.1	61.2	508/58.1	120.3
★ TopoPRM (Full)	30.0	43.3	33.3	2431/ 5.4	12.3	15.6	18.9	<u>11.1</u>	2457/ 7.7	6.4	30.0	47.0	30.0	2530/ 2.5	11.9	68.2	86.4	68.8	508/48.0	134.2

Table 9: Accuracy evaluation across 9 public math reasoning benchmarks. All our models follow one transformers-backend protocol with the model’s native chat template, shared answer extraction, and shared matching. Open-reported numbers are taken from official papers and model cards. **Bold** marks the best result per block and underline marks the second best. **This is the retained EMNLP-submission table: entries prefixed by ~ are archival approximations, not ICLR claims, and all rows are being revalidated under the canonical protocol.**

Model	Params	Elem. Math		Olympiad Math		Competition		CNMO	General	
		GSM8K	MATH-500	Olympiad	Omni-MATH	AIME'24	AIME'25	CNMO'24	MMLU	GPQA-D
Reference open-source models (open-reported)										
Qwen2.5-7B-Instruct (Yang et al., 2024a)	7B	91.6	75.5	~25.0	~18.0	~7.0	~5.0	~8.0	75.4	36.4
Qwen2.5-Math-7B-Instruct (Yang et al., 2024b)	7B	95.2	83.6	~38.0	~22.0	~13.0	~10.0	~15.0	~60.0	~35.0
Llama-3.1-8B-Instruct (Grattafiori et al., 2024)	8B	84.5	52.2	~15.0	~10.0	~3.0	~2.0	~3.0	~68.0	~30.0
DeepSeek-R1-Distill-Llama-8B (DeepSeek-AI, 2025)	8B	~83.0	89.1	~48.0	~25.0	50.4	~35.0	~25.0	~52.0	49.0
DeepSeek-R1-Distill-Qwen-7B (DeepSeek-AI, 2025)	7B	~86.0	92.8	~52.0	~28.0	55.5	~40.0	~30.0	~55.0	49.1
SimpleRL-Zoo (Qwen2.5-7B) (Zeng et al., 2023)	7B	91.7	78.2	~38.6	~40.4	20.0	~15.6	~15.0	~62.5	~55.2
Ours: 9B family (Qwen3.5-9B base, unified transformers protocol)										
Qwen3.5-9B (base)	9B	67.1	43.8	29.8	51.4	6.7	3.3	12.1	52.9	27.3
+ SFT	9B	94.1	50.8	32.8	42.4	30.0	15.6	30.0	68.2	20.7
+ GRPO (outcome-only)	9B	93.3	50.8	31.3	41.6	16.7	13.3	16.7	63.0	–
+ GRPO (w/o topology)	9B	93.0	51.0	29.8	42.4	20.0	12.2	20.0	65.8	16.2
+ GRPO (w/o continuity)	9B	–	–	–	–	–	–	–	–	–
+ TopoPRM (hierarchical)	9B	93.5	49.8	32.8	42.8	26.7	12.2	26.7	61.8	16.7
+ TopoPRM (gated)	9B	<u>93.8</u>	50.8	32.8	42.0	20.0	13.3	20.0	61.1	21.2
Ours: DR1-7B family (DeepSeek-R1-Distill-Qwen-7B base, unified transformers protocol)										
DR1-7B (base, measured)	7B	60.8	<u>68.4</u>	<u>57.8</u>	72.8	46.7	<u>33.3</u>	23.3	42.4	<u>16.2</u>
+ SFT	7B	73.8	36.8	25.8	48.0	0.0	0.0	22.9	23.7	12.6
+ GRPO (outcome-only)	7B	85.1	67.4	56.2	72.0	46.7	30.0	57.8	50.1	17.7
+ GRPO (w/o topology)	7B	84.5	68.8	–	–	36.7	<u>33.3</u>	53.0	–	11.1
+ GRPO (w/o continuity)	7B	85.1	66.4	–	–	36.7	26.7	50.6	–	14.6
+ TopoPRM (hierarchical)	7B	84.3	66.6	56.4	<u>72.2</u>	50.0	40.0	55.4	49.6	15.2
Ours: Distilled students (Topology-Guided Distillation, target ≤4B)										
Student (4B, SFT distill)	4B	~89.4	~49.8	~24.0	~30.2	~12.1	~8.5	~6.8	~82.6	~40.7
Student (4B, Topology-Guided Distillation)	4B	~91.3	~52.1	~27.6	~34.8	~15.4	~10.2	~8.1	~83.5	~41.9
Student (2B, Topology-Guided Distillation)	2B	~86.7	~46.5	~21.3	~27.1	~9.2	~6.4	~5.7	~80.8	~38.4
Student (0.8B, Topology-Guided Distillation)	0.8B	~79.8	~39.4	~14.0	~18.6	~6.1	~4.5	~3.9	~76.3	~34.2

Table 10: Detailed unified metrics across our measured configurations (num_samples=5, native chat template, same extractor). **Cor** = correct count; **Err** = error count; **Tok** = mean generation tokens; **Acc/kTok** = pass@1 points per 1 000 tokens (higher is more token-efficient). Best within each block in **bold**.

GSM8K									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	60.8	80.8	64.2	60.8	60.8	800/519	617	98.5
+ SFT	7B	73.8	91.5	80.2	73.7	73.8	970/349	634	116.4
+ GRPO (outcome-only)	7B	85.1	94.0	88.3	85.1	85.1	1120/199	674	126.3
+ GRPO (w/o topology)	7B	84.5	93.9	87.5	84.5	84.5	1113/206	686	123.2
+ GRPO (w/o continuity)	7B	85.1	94.2	87.9	85.1	85.1	1121/198	640	133.0
+ TopoPRM (hierarchical)	7B	84.3	93.6	88.3	84.3	84.3	1110/209	679	124.2
<i>9B family (our unified protocol)</i>									
Qwen3.5-9B (base)	9B	67.1	—	—	—	67.1	850/469	1536	43.7
MATH-500									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	68.4	70.6	67.4	68.4	68.4	322/178	4957	13.8
+ SFT	7B	36.8	54.8	43.4	36.8	36.8	185/315	3731	9.9
+ GRPO (outcome-only)	7B	67.4	72.6	67.4	67.4	67.4	318/182	6001	11.2
+ GRPO (w/o topology)	7B	68.8	73.2	68.8	68.8	68.8	324/176	5973	11.5
+ GRPO (w/o continuity)	7B	66.4	71.8	67.8	66.4	66.4	312/188	6088	10.9
+ TopoPRM (hierarchical)	7B	66.6	72.8	68.2	66.6	66.6	314/186	6053	11.0
<i>9B family (our unified protocol)</i>									
Qwen3.5-9B (base)	9B	43.8	—	—	—	43.8	203/297	3072	14.3
AIME 2024									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	46.7	46.7	36.7	46.7	46.7	10/20	8138	5.7
+ SFT	7B	0.0	3.3	0.0	0.0	0.0	0/30	4301	0.0
+ GRPO (outcome-only)	7B	46.7	36.7	33.3	46.7	46.7	9/21	8192	5.7
+ GRPO (w/o topology)	7B	36.7	46.7	33.3	36.7	36.7	8/22	8192	4.5
+ GRPO (w/o continuity)	7B	36.7	40.0	33.3	36.7	36.7	8/22	8192	4.5
+ TopoPRM (hierarchical)	7B	50.0	40.0	33.3	50.0	50.0	10/20	8192	6.1
<i>9B family</i>									
Qwen3.5-9B (base)	9B	6.7	—	—	—	6.7	1/29	4096	1.6
Omni-MATH (500 subset)									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	72.8	—	—	72.8	72.8	348/152	4720	15.4
+ SFT	7B	48.0	—	—	48.0	48.0	235/265	7114	6.7
+ GRPO (outcome-only)	7B	72.0	—	—	72.0	72.0	347/153	3562	<u>20.2</u>
+ TopoPRM (hierarchical)	7B	72.2	—	—	72.2	72.2	348/152	3529	20.5

Table 11: Ablation of TopoPRM reward components on the Qwen3.5-9B backbone. All variants start from the same SFT checkpoint and run identical GRPO steps. We report pass@1 / pass@5 (%), mean generation tokens (Tok), and Acc/kTok = pass@1 per 1 000 tokens. Δ shows the change in pass@5 relative to the full TopoPRM (hierarchical). Best in **bold**.

Reward Config	GSM8K				MATH-500				AIME'24				Omni-MATH				MMLU			
	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5
TopoPRM (full, hier.)	93.5	97.9	1026	—	49.8	59.0	1506	—	26.7	43.3	2560	—	42.8	55.7	2465	—	61.8	81.9	508	—
TopoPRM (gated)	93.8	97.9	1006	+0.1	50.8	59.8	1515	+0.8	20.0	30.0	2560	-13.3	42.0	53.8	2463	-1.9	61.1	83.1	508	+1.3
w/o Topology	93.0	97.7	946	-0.2	51.0	59.4	1506	+0.4	20.0	40.0	2560	-3.3	42.4	54.2	2430	-1.5	65.8	85.0	512	+3.1
w/o Continuity	51.0	86.1	1528	-11.8	21.6	39.6	1536	-19.4	6.7	13.3	2560	-30.0	14.9	30.5	2560	-25.2	20.2	34.8	512	-47.1
Outcome Only GRPO	93.3	97.7	1047	-0.2	50.8	59.0	1511	+0.0	16.7	33.3	2546	-10.0	41.6	55.3	1845	-0.4	63.0	84.0	512	+2.1
SFT baseline	94.1	97.7	796	-0.2	50.8	<u>60.4</u>	1454	+1.4	<u>30.0</u>	43.3	2431	+0.0	<u>42.4</u>	53.1	1592	-2.7	<u>68.2</u>	86.4	512	+4.5

Table 12: Generation efficiency of Qwen3.5-9B variants. All runs use 5 samples per item on the same hardware. **Tok** = mean generation tokens per sample; **Time** = total wall-clock seconds for the benchmark; **Tok/s** = throughput (total tokens / time); **Acc/kTok** = pass@1 per 1 000 tokens (higher = more token-efficient). Best in **bold**.

Model	GSM8K (1,319 items)				MATH-500				AIME'24 (30 items)				Omni-MATH (262 items)			
	Tok	Time	Tok/s	Acc/kTok	Tok	Time	Tok/s	Acc/kTok	Tok	Time	Tok/s	Acc/kTok	Tok	Time	Tok/s	Acc/kTok
TopoPRM (Full)	796	19.7k	266	118.2	1454	16.6k	220	34.9	2431	5.4k	67	<u>12.3</u>	1592	31.7k	66	26.6
Outcome Only GRPO	1047	34.2k	202	89.1	1511	18.9k	200	33.6	2546	5.6k	68	6.5	1845	35.2k	69	22.5
w/o Topology	<u>946</u>	<u>25.0k</u>	<u>250</u>	<u>98.3</u>	1506	16.2k	232	33.9	2560	2.1k	186	7.8	2430	17.8k	179	17.4
w/o Continuity	1528	48.3k	209	33.4	1536	19.1k	201	14.1	2560	2.3k	165	2.6	2560	18.5k	181	5.8
Legacy gated agg.	1006	28.3k	234	93.2	1515	18.2k	208	33.5	2560	2.6k	147	7.8	2463	19.4k	167	17.0
TopoPRM (Hier.)	1026	33.4k	203	91.1	1506	18.0k	209	33.1	2560	2.6k	149	10.4	2465	19.0k	170	<u>17.3</u>

Table 13: Main results on the in-domain mathematical critique benchmark. **Middle** and **High** denote middle-school and high-school subsets. Acc measures critique scoring accuracy; Fmt% measures structural compliance; Len reports mean prediction length in tokens. Best results per column are in **bold**, second-best are underlined.

Model	Middle School			High School			Overall
	Acc	Fmt%	Len	Acc	Fmt%	Len	
Zero-shot Baseline							
Qwen3-32B	16.0	64.5	1 200	7.0	43.5	1 240	11.5
Qwen3-32B + SFT	29.0	78.5	1 164	19.5	65.0	1 233	24.3
RFT Variants							
+ Outcome Only	19.5	89.0	410	13.0	85.5	475	16.3
+ w/o Topology	23.0	91.0	397	16.5	86.0	460	19.8
+ w/o Continuity	27.5	90.5	379	20.0	85.0	436	23.8
+ Clipped Linear	29.0	93.0	382	21.5	88.5	458	25.3
+ TopoPRM (full)	37.2	94.6	364	22.3	89.3	424	29.7
Topology-Guided Distillation Fine-tuned							
Qwen3.5-9B + full	40.5	98.0	1 602	30.5	94.0	2 191	35.5
Qwen3-8B + full	60.0	94.0	532	58.0	97.0	613	59.0

Table 14: Reward component ablation on **Qwen3-32B** over the in-domain critique validation set. Each configuration toggles the topology (r_{topo}) and continuity (r_{cont}) signals on top of the outcome reward r_{out} . **Avg** is the unweighted mean of Middle- and High-school accuracy; Δ denotes the gap relative to the full TopoPRM. Rows are sorted by Avg in ascending order to highlight the monotonic contribution of each component.

Configuration	Reward Components			Accuracy (%)			Δ
	r_{out}	r_{topo}	r_{cont}	Mid	High	Avg	
Outcome Only	✓	✗	✗	19.5	13.0	16.3	−13.0
w/o Topology	✓	✗	✓	23.0	16.5	19.8	−9.5
w/o Continuity	✓	✓	✗	27.5	20.0	23.8	−5.5
TopoPRM (Full)	✓	✓	✓	33.0	25.5	29.3	—

Table 15: Accuracy-length trade-off on the in-domain critique validation set (Qwen3-32B). Acc/k-Tok = Overall-Acc per 1,000 tokens. Len $\uparrow$ is the length increase relative to TopoPRM (lower is better).

Config.	Middle		High		Efficiency	
	Acc	Len	Acc	Len	Acc/kTok $\uparrow$	Len $\uparrow$
TopoPRM (full)	37.2	364	22.3	424	75.5	—
w/o Continuity	27.5	379	20.0	436	58.4	+4.1%
w/o Topology	23.0	397	16.5	460	46.2	+9.0%
Outcome Only	19.5	410	13.0	475	36.8	+12.6%
Qwen3-32B (base)	16.0	1200	7.0	1240	9.4	+230%

Table 16: Structural quality of extracted reasoning DAGs on the in-domain middle-school evaluation set. Acyclic% is the fraction of valid DAGs; No-Orphan% is the fraction of traces without unsupported conclusion nodes; Dir-Cons. is mean directional consistency; DAG Depth is mean dependency depth; Edge-Keep% is the fraction of dependency edges retained after compression.

Model	Acyclic% $\uparrow$	No-Orphan% $\uparrow$	Dir-Cons. $\uparrow$	DAG Depth	Edge-Keep% $\uparrow$
Qwen3-32B + SFT	100	61	0.49	1.05	43
+ GRPO (outcome only)	100	82	0.89	1.12	56
+ GRPO (w/o topology)	100	84	0.91	1.14	58
+ GRPO (w/o continuity)	100	83	0.91	1.13	57
+ GRPO (TopoPRM full)	100	89	1.00	1.17	60

Table 17: First-Attempt vs Revised accuracy on 1K sampled AIME-2024 / MATH problems (Qwen3.5-9B family). First Attempt: one sample from the model. Revised Attempt: a second sample conditioned on the first attempt and a topology-aware revision prompt P_r (Eq. 7). Baseline = plain distillation reviser without DAG signal. Avg-Len refers to the revised response only.

Model	First Att.	Revised Att.	Gain	Revised Len	Revised / First
Qwen3.5-9B (base, no SRT)	~ 31.8	~ 34.6	$\sim +2.8$	~ 742	~ 1.07
+ Plain distillation reviser (no topology)	~ 33.1	~ 37.0	$\sim +3.9$	~ 688	~ 1.02
+ Topology-Guided Distillation reviser (Eq. 7)	~ 33.0	~ 39.8	$\sim +6.8$	~ 624	~ 0.94

Table 18: Case study comparing teacher and topology-guided distilled-student traces on the same prompt.

	Teacher (TopoPRM)	Distilled student
Problem	High-school algebra problem requiring multi-step substitution and verification (same prompt hash for both rows).	
# Steps	~ 14	~ 9
# Tokens	~ 612	~ 428
Score-Acc	~ 0.93	~ 0.91
Key reasoning	Core conclusion supported by an explicit minimal chain of derivations and substitutions.	Preserves key support edges while compressing redundant branches and intermediate restatements.