

REWARDING THE GRAPH BEHIND THE CHAIN: TOPOLOGY-AWARE PROCESS SUPERVISION FOR RL AND REASONING DISTILLATION

Anonymous authors

Paper under double-blind review

ABSTRACT

Long chain-of-thought reasoning is not purely linear: intermediate results are reused across non-adjacent steps, while unsupported conclusions may refer backward or form circular support. Outcome rewards and step-local process rewards do not expose these failures. We propose TopoPRM, a topology-aware implicit process reward model that first recovers an order-agnostic directed support graph from a free-form trace and then projects it onto a dependency DAG. Crucially, the backward and cyclic edges removed by this projection remain separate direction-consistency and acyclicity signals. We aggregate outcome, format, direction, acyclicity, and continuity with a topology-conditioned Choquet reward whose interactions explicitly encode structural agreement rather than averaging the two graph diagnostics into a generic topology score. The same projected graph guides distillation toward a compact dependency-preserving trace, without requiring graph-format generation from the policy. Matched counterfactuals verify that the signals are identifiable: reversing unchanged steps changes direction by 0.715, matched cycle closure changes acyclicity by 0.328 with zero direction shift, and premature averaging retains only $0.213\times$ the direction score’s cross-trace variation. We pair these construct and edge-recovery audits with a fixed-scale, fixed-budget test of whether TopoPRM advances the accuracy–length Pareto frontier over outcome-only GRPO and additive or multiplicative aggregation.

1 INTRODUCTION

Reinforcement learning with verifiable rewards (RLVR) has substantially advanced long chain-of-thought (CoT) reasoning in large language models and now serves as a core mechanism for training mathematical and symbolic reasoning ability (Wei et al., 2022; Kojima et al., 2022; DeepSeek-AI, 2025; Yeo et al., 2025; Shao et al., 2024; Ke et al., 2026). From sparse final-answer rewards (Cobbe et al., 2021; DeepSeek-AI, 2025) to widely used process reward models (PRMs) that attach supervision to intermediate positions or prefixes (Lightman et al., 2023; Wang et al., 2024; Uesato et al., 2022), the dominant training interface still represents a trace as a *linear sequence of steps*. In practice, long reasoning is rarely linear: in mathematical derivations, a later conclusion often depends on several non-adjacent premises, an intermediate equation is reused by multiple downstream derivations, and subgoals branch before merging into the final answer (Besta et al., 2024; Zhong et al., 2026). Reconciling this structural mismatch has emerged as a core challenge for advancing reasoning supervision beyond linear credit assignment.

Recent reasoning models have made the macroscopic structure of inference clearer and more reliable, yet they remain confined to the CoT-as-result view. DeepSeek-R1 (DeepSeek-AI, 2025), DAPO (Yu et al., 2025), and Tulu 3 (Lambert et al., 2024) demonstrated that outcome-only rewards scale, but they could not distinguish a well-supported derivation from a trace that reached the same answer through unsupported intermediate claims, and they remained vulnerable to reward hacking (Gao et al., 2023; Skalse et al., 2022). On the PRM side, denser signals have been obtained through Monte Carlo rollouts (Wang et al., 2024; Luo et al., 2024), trained verifiers (Zhang et al., 2025; Khalifa et al., 2025; Zhao et al., 2026), and implicit outcome-derived scoring (Yuan et al., 2025; She et al., 2025). These methods can condition on a full prefix, but they do not explicitly recover the **multi-parent support relation** in which one intermediate result supports, or is supported by,




Figure 1: TopoPRM overview. A free-form mathematical trace is mapped to a layered dependency graph; direction and acyclicity enter the topology-conditioned multi-source reward, while the same graph diagnostics localize dependency-preserving teacher revisions for distillation. The policy itself never emits graph tags.

several non-adjacent steps (Wang et al., 2025b). Likewise, token-level distillation (Agarwal et al., 2024) does not require a compact target to preserve those cross-step dependencies.

We propose TOPOPRM, which supervises reasoning through its *implicit topological dependency structure* rather than assuming that surface order is already valid (Figure 1). A frozen, order-agnostic edge encoder first recovers a raw directed support graph $\mathcal{H}_y$ from an unchanged free-form trace. Projecting $\mathcal{H}_y$ onto a text-order DAG $\mathcal{G}_y$ makes the discarded structure informative: backward-edge mass measures direction inconsistency, while cycle-edge mass measures circular support. We retain these as separate reward sources and aggregate them non-additively with outcome, format, and continuity. In particular, direction–acyclicity synergy rewards traces that satisfy both local support order and global non-circularity, while their outcome and continuity interactions make topology change rollout ordering rather than act as a small generic bonus. Finally, the same projected graph guides distillation toward a compact dependency-preserving trace. The policy never has to decode graph tags or a special DAG format. Our contributions are threefold:

1. **Raw support graph before acyclic projection.** We recover directed dependencies without imposing textual order, and turn the backward and cyclic structure removed by projection into explicit process supervision.
2. **Topology-conditioned multi-source optimization.** We introduce a monotone 2-additive Choquet reward that keeps direction and acyclicity independent and explicitly models their mutual, outcome, and continuity interactions, then reuse the graph diagnostics for Topology-Guided Distillation.
3. **Falsifiable structural and Pareto evaluation.** Matched reversal, cycle-closure, edge-noise, and premature-aggregation controls test whether direction and acyclicity are identifiable; an independent three-copy edge audit and fixed-scale, fixed-budget comparisons test recovery quality and the downstream accuracy–length frontier.

2 RELATED WORK

Reinforcement Learning with Verifiable Rewards. RLVR replaces preference models with programmatic correctness checks on tasks whose answers are objectively verifiable (Shao et al., 2024; DeepSeek-AI, 2025; Lambert et al., 2024; Zeng et al., 2023). Recent work refines GRPO-style optimization through dynamic sampling and clipping (Yu et al., 2025), extended training schedules (Liu et al., 2025), finer credit assignment (Kazemnejad et al., 2025), asymmetric handling of negative samples (Zhu et al., 2025), entropy regularization (Cui et al., 2025), and simplified baselines (Ahmadian et al., 2024; Hu, 2025). Outcome-only signals nevertheless remain structurally blind: traces with very different dependency support can be reward-equivalent (Chen et al., 2026), and reward

hacking persists even when the visible chain appears benign (Gao et al., 2023; Skalse et al., 2022; Wang et al., 2026b). *TopoPRM injects frozen topology and continuity diagnostics from latent reasoning graphs, targeting structural credit assignment directly rather than through a correctness surrogate.*

Concise-reasoning objectives introduce a separate confound: penalizing length inside a group-relative reward can assign negative advantage to a correct but long rollout. DRPO addresses this by decoupling the length-driven target distribution from correctness discrimination (Li et al., 2026a). DeCS localizes the necessary reasoning prefix with a separate detector, applies token-level reward only to the resulting chunks, and schedules prompt difficulty (Jiang et al., 2026); we reproduce its released 7B checkpoint under our fixed evaluator. LASER-D instead assigns step length rewards using targets that adapt to training-time behavior and per-question difficulty (Liu et al., 2026b); ARLCP coordinates difficulty-calibrated length and detected reflection penalties (Yu et al., 2026); and Self-Aligned Reward favors query-specific responses through a relative-perplexity signal rather than explicit trace structure (Han et al., 2026). Step-Tagging monitors semantic step types with lightweight classifiers and stops generation from their running counts (Belkhit et al., 2025); it controls inference behavior but does not recover directed support among steps. These methods target redundancy or response efficiency, whereas our direction and acyclicity scores distinguish support relations that can occur at equal length and with unchanged wording. For multi-source reward integration, fair and stable reward composition alternates between mirror-descent weight updates and policy updates but retains a dynamic weighted sum (Li et al., 2024); HERO stratifies dense reward-model scores by binary verifier outcomes and reweights prompts by within-group variance (Tao et al., 2026). The former adapts source weights without representing source complementarity, while the latter preserves correctness ordering and makes dense feedback explicitly correctness-first. Our hypothesis is different: raw direction and acyclicity should remain identifiable, interacting reward sources rather than only additive terms or tie-breakers inside correctness strata. We therefore train a Shapley-matched additive control, compare additive, multiplicative, and Choquet rankings on one frozen candidate pool, and treat accuracy and length as co-primary policy outcomes.

Process Reward Models for Reasoning. PRMs score intermediate steps to address the sparsity of outcome rewards, originally relying on dense human annotation (Uesato et al., 2022; Lightman et al., 2023). Subsequent work reduces labeling cost via Monte Carlo or tree-search rollouts (Wang et al., 2024; Luo et al., 2024; Zhang et al., 2024; Lu et al., 2024; Zhang et al., 2025; Guan et al., 2025), generative PRMs make verifier reasoning explicit (Khalifa et al., 2025; Zhao et al., 2026; She et al., 2025), while recent process rewards use intrinsic semantic cues or explicitly tagged meta-reasoning actions (He et al., 2026; Zhang et al., 2026c). Process advantage verifiers tie step values to rollout-induced returns (Setlur et al., 2025; Lai et al., 2024); implicit PRMs derive them from outcomes alone (Yuan et al., 2025; Snell et al., 2024), and CRM conditions each step on the full preceding sequence while linking it to the final outcome (Zhang et al., 2026b). Recent benchmarks expose persistent weaknesses on fine-grained error categories and inter-step consistency (Song et al., 2025; Zheng et al., 2025). Such scorers can model a prefix but do not recover an explicit many-to-many support relation among steps. *TopoPRM requires no tagged policy format and trains no scalar neural process verifier; it derives frozen global-structure supervision from a recovered dependency graph together with local continuity.*

Structured and Graph-Based Reasoning. PARC links a completed chain to preceding premises for error identification (Mukherjee et al., 2025), while GRiD generates knowledge/reasoning nodes and verifies each node against declared parents (Wen et al., 2025). Graph of Verification adaptively decomposes a trace into a DAG and verifies node blocks in topological order without training (Fang et al., 2026). Other methods require graph-formatted reasoning: DAG-Math evaluates graph-level logical closeness (Zhang et al., 2026a), Structured Reasoning rewards within-graph flow and cross-sample overlap (Dong et al., 2026a), Structural CoT trains executable computation graphs (Anonymous, 2026b), VeriGraph makes a data-analytic agent emit a heterogeneous evidence DAG and rewards its integrity and coherence (Jin et al., 2026), the Graph Reasoning Paradigm emits labeled symbolic graphs (Liu et al., 2026a), and GoT-R1 reinforces parsability, sparsity, and non-linearity of serialized reasoning graphs with an additive reward (Li et al., 2026b). AtomGraph instead learns a graph-attention verifier over reasoning chains that have already been parsed into DAGs (Karmore, 2026). Contract-checked graph editing enforces acyclicity as a hard gate for composing candidate DAGs during inference-time search (Li & Cao, 2026). TUR-DPO sanitizes extracted claim-support

graphs and linearly combines valid-path, cycle, dangling-node, and contradiction terms to shape offline preference optimization (Abdullah et al., 2026); it neither retains edge-direction residuals nor separates direction from cyclic support. RewardFlow merges observed agent trajectories into an environment-state transition graph and propagates terminal reward along action edges (Feng et al., 2026); it does not recover support direction within one free-text derivation. FlowTracer instead derives token credit from an attention-induced answer-flow DAG (Dong et al., 2026b). Two methods named Graph-GRPO propagate credit over a predefined e-commerce relevance graph or optimize sampled multi-agent communication topologies, rather than dependencies within one free-form trace (Che et al., 2026; Cang et al., 2026). TopoPRM instead leaves policy traces in free text and recovers an unconstrained semantic-support graph before projection. Declared-parent, attention-DAG, learned-DAG, valid-DAG, and hard-gated representations rule out or do not score backward and cyclic semantic support; we retain these residuals as direction and circular-justification signals.

Post-hoc graph probes support non-sequential reasoning (Zhong et al., 2026; Matsutani et al., 2026); CoT2Graph uses recovered graphs for failure analysis, outcome prediction, and Best-of- N selection (Anonymous, 2026a), and Reasoning Consensus merges several recovered DAGs for inference-time structural ensembling (Parulekar et al., 2026). SARL clusters reasoning steps into latent functions and rewards clustering and short paths in their *undirected* transition graph, explicitly abstracting away directional control flow (Wang et al., 2026c). This small-world prior is complementary to our semantic-support relation: our direction score detects backward dependency mass, while acyclicity detects global self-support. Hidden-state cycles can reflect useful recurrent visits (Minegishi et al., 2025), while LoopBench studies self-reinforcing textual repetition during generation (Duan et al., 2026). Our cycles instead denote circular semantic support in the committed derivation: we do not penalize exploration, revision, or repeated wording, only a final support relation in which a claim depends on itself. The novelty is therefore not “reasoning as a graph,” but raw direction/cycle diagnostics as online, auditable reward sources without graph-formatted generation; Appendix G makes this design boundary explicit against the closest alternatives.

Distillation and Reasoning Compression. Knowledge distillation transfers teacher capability (Hinton et al., 2015); student-generated variants reduce train–test mismatch (Gu et al., 2024; Agarwal et al., 2024; Ko et al., 2024), including ExOPD, BOND, KDRL, and black-box GAD (Yang et al., 2026; Sessa et al., 2024; Xu et al., 2025; Ye et al., 2025). Long-CoT compression uses pruning, length penalties, or budget control (Xia et al., 2025; Luo et al., 2025; Arora & Zanette, 2026; Aggarwal & Welleck, 2025; Muennighoff et al., 2025); graph-based pruning removes weak branches from projected DAGs (Yuan et al., 2026), while Reasoning Scaffolding distills semantic operations (Wen et al., 2026). TopoPRM instead uses raw direction/cycle violations to elicit and filter teacher revisions, preserving dependencies that token budgets or post-projection pruning cannot identify. *Topology-Guided Distillation is evaluated by the accuracy–length frontier and independently audited topology preservation.*

3 METHOD

TopoPRM (Figure 2) is a topology-aware post-training framework whose stages share a graph encoder $\mathcal{E}_\phi$ and an acyclic projection Π . Stage I produces parsable reasoning traces. Stage II applies GRPO to a topology-conditioned multi-source reward. Stage III performs Topology-Guided Distillation, in which the Stage-II teacher revises student rollouts using defects localized on the projected graph.

3.1 PROBLEM SETUP

Let $x = (q, c)$ pair a mathematical problem q with optional context c , and let π_θ generate a trace $y = (s_1, \dots, s_T)$. The frozen encoder produces a raw directed support graph $\mathcal{H}_y = \mathcal{E}_\phi(y) = (\mathcal{V}, \mathcal{A}_H)$, where $i \rightarrow j$ means that step s_i supplies a result, claim, sub-goal, or assumption used by s_j . Unlike a DAG constructed only from pairs $i < j$, $\mathcal{H}_y$ may contain backward edges and cycles. We retain these violations and separately form $\mathcal{G}_y = \Pi(\mathcal{H}_y)$ by keeping support edges consistent with the trace order. Thus $\mathcal{G}_y$ supports DAG operations, while the projection cost of $\mathcal{H}_y$ exposes structural failures that would otherwise be removed by construction.

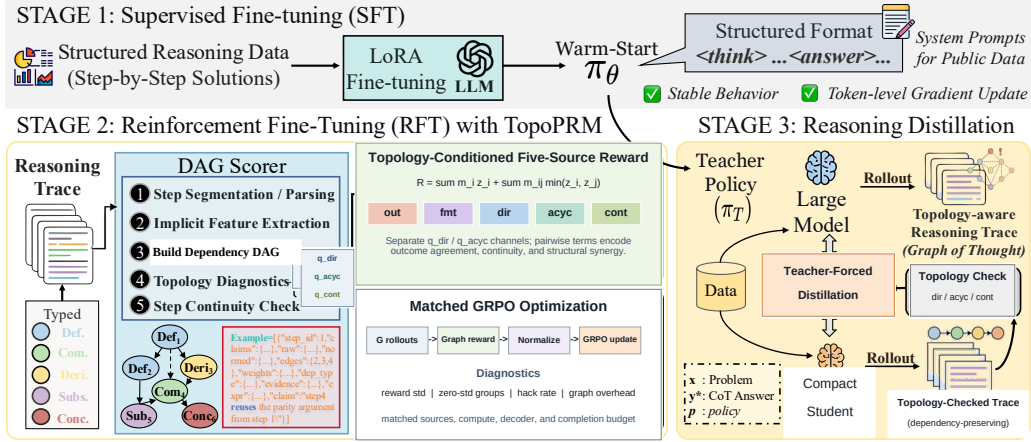


Figure 2: Overview of TopoPRM and Topology-Guided Distillation. An order-agnostic encoder recovers a raw directed support graph; its acyclic projection provides the training DAG while discarded backward and cyclic edges remain reward signals. A topology-conditioned Choquet reward aggregates multiple sources, and the projected graph guides dependency-preserving distillation.

Scope of the acyclicity assumption. Acyclicity is normative only for the *semantic support relation in the committed trace*. It does not forbid exploratory computation, checking an earlier result, or revising a failed branch. Such behavior may revisit a latent state or repeat a concept while still yielding an acyclic final justification. A cycle in $\mathcal{H}_y$ has the narrower meaning that a written claim ultimately supports itself; q_{acyc} targets this circular justification rather than recurrent reasoning in general.

3.2 TOPOLOGY-AWARE DATA PROCESSING

Raw-Graph Extraction and Projection. Given a rollout y , $\mathcal{E}_\phi$ proceeds in three stages. **(i) Segmentation and parsing.** It splits y into steps and sub-questions, then extracts normalized expressions, claim keys, variables, and step types. **(ii) Order-agnostic support encoding.** Auditable rules propose high-precision symbolic candidates, while a lightweight pair encoder predicts one of $\{i \rightarrow j, j \rightarrow i, \emptyset\}$ for each unordered intra-component pair. The encoder is trained offline from semantic-teacher edges and frozen during policy optimization; a separately annotated set calibrates edge precision and direction. Retained encoder edges carry their calibrated confidence, while rule edges use unit weight. Semantic predictions override an order-based rule when both judge the same pair, preventing artificial two-cycles caused by conflicting extractors. **(iii) Acyclic projection.** We preserve all predicted edges in $\mathcal{H}_y$, set $\mathcal{G}_y = \Pi(\mathcal{H}_y)$ by removing $i \rightarrow j$ with $i \geq j$, and add an adaptive weak sequential edge only when a node remains unsupported. Full rules and abstention thresholds are in Appendix C.

Topology and Continuity Diagnostics. Let $\mathcal{A}_H^{\text{dep}}$ denote the dependency edges of $\mathcal{H}_y$, let $w_{ij} \geq 0$ be the confidence weight of edge (i, j) , and set $W_H = \sum_{(i,j) \in \mathcal{A}_H^{\text{dep}}} w_{ij}$. Let $\mathcal{A}_{\text{cyc}}$ contain the dependency edges whose endpoints lie in the same cyclic strongly connected component (size greater than one, or a self-loop), and let c_H be the fraction of nodes incident to at least one dependency edge. We first compute d_H , the forward-edge mass, and u_H , the cycle-edge mass, then shrink both diagnostics toward 0.5 when edge coverage is low:

$$q_{\text{dir}}(\mathcal{H}_y) = c_H d_H + (1 - c_H)/2, \quad d_H = \frac{\sum_{(i,j) \in \mathcal{A}_H^{\text{dep}}} w_{ij} \mathbb{1}[i < j]}{W_H}, \quad (1)$$

$$q_{\text{acyc}}(\mathcal{H}_y) = c_H(1 - u_H) + (1 - c_H)/2, \quad u_H = \frac{\sum_{(i,j) \in \mathcal{A}_{\text{cyc}}} w_{ij}}{W_H}. \quad (2)$$

When $W_H = 0$, we use $d_H = 1$ and $u_H = 0$; because $c_H = 0$ for a multi-step empty graph, both coverage-adjusted scores are nevertheless 0.5, not perfect. (A singleton trace has no possible cross-step violation and uses $c_H = 1$.) The continuity score q_{cont} measures whether each step is traceable

to the question or a supported prefix of $\mathcal{G}_y$. Direction and acyclicity remain independent properties of the raw graph rather than being averaged into one topology score or checked tautologically after projection. They are related but not redundant: every directed cycle contains a backward edge under the textual order, whereas a backward edge need not close a cycle, and raw graphs with identical backward-edge mass can have different cycle-edge mass. Thus q_{dir} detects local order inconsistency while q_{acyc} detects global circular support.

Topology-Conditioned Multi-Source Aggregation. Let $\mathbf{z} = (r_{\text{out}}, r_{\text{fint}}, q_{\text{dir}}, q_{\text{acyc}}, q_{\text{cont}}) \in [0, 1]^5$. A weighted sum assumes that each source has the same marginal value regardless of the others. Instead, we use a monotone 2-additive Choquet model with Möbius coefficients (Choquet, 1954; Grabisch, 1997)

$$r_{\text{total}}(x, y) = \sum_i m_i z_i + \sum_{(i,j) \in \mathcal{I}_{\text{topo}}} m_{ij} \min(z_i, z_j), \quad \mathcal{I}_{\text{topo}} = \{(o, d), (o, a), (d, c), (a, c), (d, a)\}. \quad (3)$$

Here o, d, a, c denote outcome, direction, acyclicity, and continuity. All $m_i, m_{ij} \geq 0$ and sum to one, so the reward is bounded and monotone. The five interactions assign additional value when each raw-graph diagnostic agrees with outcome or continuity and when direction agrees with acyclicity. This makes topology a constituent of cross-source utility rather than a correctness-first bonus; no length source enters the main reward. We ablate the capacity against equal-weight, Shapley-matched additive, and multiplicative alternatives, and remove the direction–acyclicity interaction in isolation.

3.3 TOPOLOGY-GUIDED POST-TRAINING

Stage I: Supervised Warm-Start. We perform Supervised Fine-Tuning (SFT) on structured mathematical reasoning data to teach the policy to produce stepwise solutions parsable by $\mathcal{E}_\phi$. This reduces invalid rollouts during subsequent reinforcement learning, since topology diagnostics are meaningful only when the trace contains recoverable reasoning steps.

Stage II: GRPO with TopoPRM Rewards. Starting from the warm-start checkpoint, we optimise the policy with GRPO. For each completion, $\mathcal{E}_\phi$ constructs $\mathcal{H}_y$ and $\mathcal{G}_y$; Eq. 3 then scores the completion. Standard group normalization acts on the complete Choquet reward, so topology can change the ordering of rollouts rather than only break ties inside an outcome stratum. Reward-hacking tests separately measure whether structurally plausible but answer-incorrect traces are over-ranked.

Stage III: Topology-Guided Distillation. The optimized model after Stage II serves as the teacher π_T , while a smaller pretrained checkpoint initializes the student π_S . Topology-Guided Distillation conditions target construction on topology diagnostics from student-generated traces. For each prompt x , the student first rolls out $y_{\text{init}} \sim \pi_S(\cdot | x)$; $\mathcal{E}_\phi$ produces $\mathcal{H}_{y_{\text{init}}}$ and $\mathcal{G}_{y_{\text{init}}}$; a dispatcher ψ then issues a revision instruction targeting a backward edge, cyclic component, orphan conclusion, or continuity break. The teacher generates a revised target under this instruction. Thus topology determines where supervision is requested and which revisions are retained, rather than serving only as a post-hoc score.

Distillation objective. Let $A(y_{\text{init}}, y_{\text{rev}})$ indicate that the revised trace has a correct boxed answer and valid format, satisfies $q_{\text{dir}}^{\text{rev}} \geq \tau_d$, $q_{\text{acyc}}^{\text{rev}} \geq \tau_a$, and $q_{\text{cont}}^{\text{rev}} \geq \tau_c$, does not reduce either raw-graph diagnostic relative to the initial trace, and fits a fixed token budget B . The retained dataset is $\mathcal{D}_{\text{distill}} = \{(x, y_{\text{rev}}) : A(y_{\text{init}}, y_{\text{rev}}) = 1\}$. The student minimises standard teacher-forced distillation loss

$$\mathcal{L}_{\text{distill}}(\theta) = -\mathbb{E}_{(x,y) \sim \mathcal{D}_{\text{distill}}} \sum_{t=1}^{|y|} \log \pi_S(y_t | x, y_{<t}). \quad (4)$$

Unlike static trace distillation, the target distribution is concentrated on revisions elicited by the current student’s structural failures. The explicit direction and acyclicity constraints prevent apparent compression from being obtained by deleting necessary support or introducing circular justification. The full pipeline is summarised in Appendix Algorithm 1.

4 EXPERIMENTS

Draft audit status. Red cells are not claims. Each method uses one canonical seed-0 training run; reported intervals quantify evaluation-set uncertainty, not training variance. The original EMNLP tables remain archival in Appendix H until they satisfy the canonical protocol below.

4.1 CANONICAL PROTOCOL

Fixed model scale and budget. The primary comparison starts from the same Qwen3.5-9B Stage-I checkpoint and changes only the Stage-II reward. The public frontier is restricted to dense released checkpoints with 7B–9B parameters; sparse models with similar total but much smaller active parameter counts are outside this matched scale. All variants use rank-64 LoRA, 300 GRPO steps, four generations per prompt, learning rate 5×10^{-6} , $\beta=0.04$, and a 4,096-token completion budget. Stage I uses 18,487 structured math solutions for training and 983 source-disjoint solutions for validation, three epochs, learning rate 5×10^{-5} , and global batch 32. Stage II uses 14,893 unique public train-split questions after exact normalized-question exclusion against every primary evaluation set. We remove 76 of 14,969 leakage-free candidates whose exact Qwen3.5 chat prompts exceed 512 tokens, rather than truncating their problem statements. Reference-solution DAGs are withheld because their step indices are not aligned with generated traces; q_{dir} , q_{acyc} , and q_{cont} are computed only from each rollout’s $\mathcal{H}_y$ and $\mathcal{G}_y$. The pair encoder is frozen in every topology-aware run; the 9B semantic teacher is used only to create its offline supervision and never appears in the GRPO reward loop.

Benchmarks and metrics. The primary fixed-scale frontier uses GSM8K, MATH-500, the official 675-item English text-only mathematics subset of OlympiadBench, and AIME 2024, spanning elementary through competition mathematics. Pass@1 exact-answer accuracy and mean visible response tokens are co-primary; correctness is read only from the last non-empty box or a terminal explicitly marked answer, so intermediate expressions in truncated traces receive no credit. EOS padding is excluded by retokenizing each decoded response without special tokens. Method *A* Pareto-dominates *B* only when it is no less accurate and no longer, with one strict inequality; pass@5 and maj@5 are secondary diagnostics. Every reproduced row uses one seed-0 sampled response per item with temperature 0.6, top- $p=0.95$, top- $k=20$, min- $p=0$, repetition penalty 1.0, identical task content and verifier, and a 4,096-token cap. Released chat templates and documented role/thinking controls are retained; matched rows share tokenizer and message roles (Appendix A.4).

Statistics and provenance. Every row maps to a checkpoint, manifest, and per-example output. Marginal benchmark accuracy uses Wilson 95% intervals; token means and stratified macros use item bootstrap, while matched accuracy and length differences use paired bootstrap over shared items. With one training run, these quantify evaluation-set uncertainty only. Public checkpoints with longer budgets or repeated decoding are re-evaluated under our single-response 4,096-token protocol rather than quoted in Table 1. The human audit reports agreement, edge precision/recall, direction accuracy, and abstention coverage separately from teacher-held-out validation.

4.2 ACCURACY–LENGTH FRONTIER (RQ1)

Table 1 tests the central empirical claim: at a fixed 9B adaptation scale and response budget, does the raw-directed TopoPRM reward move the accuracy–length frontier beyond Stage-I SFT, outcome-only GRPO, the dynamic difficulty-aware LASER-D baseline, and a Shapley-matched additive multi-source reward? Our matched LASER-D adaptation preserves the released mechanism: training-time length coverage determines a target for each rollout-difficulty band, and only correct responses within that target receive its step reward (Liu et al., 2026b); Appendix B records the budget-scaled thresholds and monitor protocol. The completed Stage-I row has macro accuracy 42.1% (stratified item-bootstrap 95% interval [40.2, 44.4]) and 526 visible tokens ([423, 647]); these intervals quantify evaluation-set uncertainty, not training variance. A SOTA claim will be made only if TopoPRM is non-dominated against all matched rows and improves the reproduced accuracy–length frontier over the listed 7B–9B public checkpoints. Until then, the table remains a preregistered comparison rather than a result narrative. *External bars.* FlowTracer uses an attention-induced answer-flow DAG under a 1K-output protocol (Dong et al., 2026b); Structured Reasoning

Table 1: Fixed-budget accuracy–length comparison at dense 7B–9B scale. Cells report pass@1 (%) and mean visible response tokens from one seed-0 sample per item under matched sampling, task content, and a 4,096-token cap; released chat templates and required thinking controls are retained. Adapted methods use one canonical seed-0 training run. **Accuracy intervals are Wilson; token and paired-difference intervals use item bootstrap. Neither measures training variance.**

Method	GSM8K		MATH-500		OlympiadBench		AIME 2024		Macro avg.	
	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.	Acc.	Tok.
<i>Public checkpoints, reproduced with the same evaluator</i>										
Qwen3-8B (Qwen Team, 2025)	91.4	1,948	64.6	3,185	–	–	0.0	4,096	–	–
SpyRL-Qwen3-8B-Math (Wang et al., 2026a)	–	–	–	–	–	–	–	–	–	–
GoT-R1-8B (Li et al., 2026b)	95.2	357	–	–	–	–	16.7	2,018	–	–
DeepSeek-R1-0528-Qwen3-8B (DeepSeek-AI, 2025)	–	–	–	–	–	–	–	–	–	–
DeCS-7B (Jiang et al., 2026)	–	–	–	–	–	–	–	–	–	–
MiMo-7B-RL-0530 (Xiaomi LLM Core Team, 2025)	–	–	–	–	–	–	–	–	–	–
Nemotron-Cascade-8B-Thinking (Wang et al., 2025a)	–	–	–	–	–	–	–	–	–	–
Klear-Reasoner-8B (Su et al., 2025)	–	–	–	–	–	–	–	–	–	–
OpenReasoning-Nemotron-7B (NVIDIA, 2025)	83.0	2,429	51.2	3,260	–	–	0.0	4,096	–	–
AceMath-RL-Nemotron-7B (NVIDIA, 2026)	–	–	–	–	–	–	–	–	–	–
Qwen3.5-9B (Qwen Team, 2026)	76.4	2,710	40.4	3,627	–	–	0.0	4,096	–	–
<i>Matched Qwen3.5-9B adaptation from the same Stage-I checkpoint</i>										
Stage-I SFT	85.8	137	59.6	283	19.6	604	3.3	1,080	42.1	526
Outcome-only GRPO	–	–	–	–	–	–	–	–	–	–
LASER-D (matched) (Liu et al., 2026b)	–	–	–	–	–	–	–	–	–	–
Matched additive reward	–	–	–	–	–	–	–	–	–	–
TopoPRM (Choquet)	–	–	–	–	–	–	–	–	–	–

Table 2: Direction and acyclicity mechanism tests. **Top:** matched policy ablations use one decoder and budget. **Bottom:** targeted frozen-encoder counterfactuals; Appendix Table 7 gives the complete ledger and paired intervals. **Red cells await canonical policy outputs.**

Policy variant	Accuracy (%)				Aggregate		Generated topology	
	GSM8K	MATH	Olymp.	AIME	Macro	Tok.	q_{dir}	q_{acyc}
Full TopoPRM	–	–	–	–	–	–	–	–
w/o direction source	–	–	–	–	–	–	–	–
w/o acyclicity source	–	–	–	–	–	–	–	–
w/o raw support	–	–	–	–	–	–	–	–
<i>Frozen-encoder counterfactuals on identical source traces</i>								
Control	n	Direction response	Acyclicity response	Isolated effect				
Reverse all steps	120	$\Delta q_{\text{dir}} = -.715$	$\Delta q_{\text{acyc}} = +.064$	Textual order / direction				
Matched cycle closure	57	$\Delta q_{\text{dir}} = .000$	$\Delta q_{\text{acyc}} = -.328$	Cycle at fixed backward mass				
Premature averaging	120	$\sigma_{\text{dir}} = .110$	$\sigma_{\text{acyc}} = .107$	Avg. $\sigma = .023$ (.213 $\times$); opposite: 39.2%				

applies attention-MaxFlow or cross-response LCS rewards to explicitly tagged traces (Dong et al., 2026a); and the Graph Reasoning Paradigm makes the policy emit labeled graph nodes, reporting a 7,040-token AIME mean that exceeds our cap (Liu et al., 2026a). ARLCP and Self-Aligned Reward instead target efficient reasoning without explicit support direction (Yu et al., 2026; Han et al., 2026). Because their models, training, and decoders differ, these are external bars rather than rows mixed into the reproduced Pareto comparison.

4.3 WHERE THE GAIN COMES FROM (RQ2)

Raw direction and acyclicity. Table 2 pairs policy ablations with matched counterfactuals. Reversal changes ($q_{\text{dir}}, q_{\text{acyc}}$) by $(-0.715, +0.064)$. Their paired 95% intervals are $[0.657, 0.771]$ and $[0.040, 0.091]$, consistent with changed order but no added cycle. Across 57 edge-count, backward-mass, and coverage-matched pairs, cycle closure changes only q_{acyc} by -0.328 $[0.292, 0.366]$. Direction-flip and false-edge controls selectively reduce q_{dir} by 0.180 and q_{acyc} by 0.190 (Appendix Table 7; Figure 3).

Table 3: Frozen four-candidate aggregation diagnostic. Rows rank identical Stage-I samples; benchmark cells are selection, not trained-policy, accuracy. AUC measures correct–incorrect pair ranking and StructWrong counts wrong selections with $q_{\text{dir}}, q_{\text{acyc}} \geq .8$. **Cells await the canonical pool; intervals are item bootstrap, not training variance.**

Aggregator	GSM8K	MATH	Olymp.	AIME	Macro	AUC↑	StructWrong↓	Tok.↓	q_{dir} ↑	q_{acyc} ↑
Outcome only	–	–	–	–	–	–	–	–	–	–
Equal-weight additive	–	–	–	–	–	–	–	–	–	–
Shapley-matched additive	–	–	–	–	–	–	–	–	–	–
Shapley-matched multiplicative	–	–	–	–	–	–	–	–	–	–
Choquet w/o dir.–acyc. term	–	–	–	–	–	–	–	–	–	–
Topology-conditional Choquet	–	–	–	–	–	–	–	–	–	–

Table 4: Independent three-copy blind edge audit. **Cells await annotation.**

Extractor	n	Overall edge recovery			Topology validity			Edge F1 by trace length			Agreement
		Prec.	Rec.	F1	Dir.	Cyc.	Cov.	Short	Medium	Long	κ
Adjacent-step chain	120	–	–	–	–	–	1.00	–	–	–	–
Rules only	120	–	–	–	–	–	–	–	–	–	–
Pair encoder (no abstention)	120	–	–	–	–	–	1.00	–	–	–	–
Frozen pair encoder (calibrated)	120	–	–	–	–	–	–	–	–	–	–
Semantic teacher (audit upper bound)	120	–	–	–	–	–	–	–	–	–	–
Projected DAG only	120	–	–	–	–	–	–	–	–	–	–
Hybrid raw-graph encoder	120	–	–	–	–	–	–	–	–	–	–

On reversed controls the scores are anticorrelated ($r = -0.907$, $[-0.976, -0.779]$): averaging shrinks cross-trace standard deviation from 0.110 to 0.023, a ratio of 0.213 ($[0.110, 0.320]$), and 39.2% of traces move oppositely. Keeping q_{dir} and q_{acyc} separate therefore preserves textual-order and global self-support information that premature averaging cancels; the policy panel tests its downstream value.

Non-additive aggregation. Table 3 isolates scalar aggregation by ranking the same four Stage-I candidates. It cannot populate trained-policy cells: the formal comparison trains only full and Shapley-matched additive rewards. We retain topology interactions only if frozen selection and the trained frontier both improve over matched additive weighting.

4.4 INDEPENDENT STRUCTURAL AUDIT AND COST (RQ3)

The 120 traces balance original, reversed, and interleaved order (40 each). Three blinded copies share each trace’s condition but independently shuffle presentation, supporting pairwise agreement and a strict-majority reference without revealing model edges, scores, or IDs. The dashboard exposes every graph and score calculation (Appendix Figure 4); Table 4 audits edge recovery, direction, cycles, coverage, and incorrect-plausible traces. Cost separates offline labeling/adaptation from online latency, memory, graph pairs/s, and GRPO share under the logged $4 \times \text{L20}$ allocation.

5 CONCLUSION

We presented TopoPRM, which converts backward and cyclic support into direction and acyclicity supervision before DAG projection and reuses both signals for reward aggregation and distillation without graph-format decoding. Matched counterfactuals show that the signals respond differently and that premature averaging suppresses cross-trace variation. They depend on edge recovery and complement, rather than replace, semantic verification.

Scope and limitations. We supervise a *committed* derivation, not its hidden search. A coherent trace may still be false, so topology complements correctness and inherits edge-encoder errors; teacher metrics cannot replace the human audit. Claims are restricted to public mathematics, dense 7B–9B checkpoints, and one budget; Choquet and distillation claims remain conditional on matched non-dominance.

ETHICS STATEMENT

The canonical study uses public mathematics benchmarks and collects no demographic attributes. Human annotators label directed support relations in reasoning traces; their identities are not part of the released annotations. An archival appendix table retains a prior private teacher-grading benchmark for continuity with the earlier submission, but it does not support the ICLR claims. **Before submission, that table requires documented authorization, privacy review, and release constraints; otherwise it will be removed from the submitted appendix.**

REPRODUCIBILITY STATEMENT

This draft is being audited for implementation-to-equation consistency and row-level result provenance. The submission version will point to the exact method definitions, configurations, raw-result manifests, and anonymized code bundle used to reproduce every retained result.

AI USE STATEMENT

Generative AI tools assisted with prior-review analysis, literature search and synthesis, hypothesis and experimental-control refinement, method implementation and code review, dataset cleaning, reproducibility auditing, scientific-figure editing, and manuscript drafting and revision. Before submission, the authors will manually inspect the cited literature, review code and generated artifacts, and verify every reported result against raw outputs. The authors take responsibility for the final claims, code, artifacts, and results; this disclosure will be updated if additional uses occur before submission.

REFERENCES

- Abdulahdy Abas Abdullah, Fatemeh Daneshfar, Seyedali Mirjalili, and Mourad Oussalah. TUR-DPO: Topology- and uncertainty-aware direct preference optimization. In *Proceedings of the 43rd International Conference on Machine Learning*, 2026. URL <https://arxiv.org/abs/2605.00224>.
- Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *International Conference on Learning Representations*, 2024.
- Pranjal Aggarwal and Sean Welleck. L1: Controlling how long a reasoning model thinks with reinforcement learning. *arXiv preprint arXiv:2503.04697*, 2025.
- Arash Ahmadian, Chris Cremer, Matthias Gallé, Marzieh Fadaee, Julia Kreutzer, Olivier Pietquin, Ahmet Üstün, and Sara Hooker. Back to basics: Revisiting reinforce-style optimization for learning from human feedback in llms. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 12248–12267, 2024.
- Anonymous. Understanding failures in LLM reasoning by learning structured representations of chain-of-thought. Anonymous ACL submission, 2026a. URL <https://openreview.net/forum?id=0XfuJjhaI5>.
- Anonymous. Structural chain-of-thought: Reasoning with executable computational graphs. ACL Rolling Review March 2026 submission, 2026b. URL <https://openreview.net/forum?id=iv6poj9Y4n>.
- Daman Arora and Andrea Zanette. Training language models to reason efficiently. *Advances in Neural Information Processing Systems*, 38:60770–60808, 2026.
- Yannis Belkhir, Seshu Tirupathi, Giulio Zizzo, and John D. Kelleher. Step-tagging: Toward controlling the generation of language reasoning models through step monitoring. *arXiv preprint arXiv:2512.14332*, 2025. URL <https://arxiv.org/abs/2512.14332>.

- Maciej Besta, Nils Blach, Ales Kubicek, Robert Gerstenberger, Michal Podstawski, Lukas Gianinazzi, Joanna Gajda, Tomasz Lehmann, Hubert Nyczyk, Marcin Copik, and Torsten Hoeffler. Graph of thoughts: Solving elaborate problems with large language models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(16):17682–17690, 2024.
- Yueyang Cang, Xiaoteng Zhang, Erlu Zhao, Zehua Ji, Yuhang Liu, Yuchen He, Zhiyuan Ning, Chen Yijun, Wenge Que, and Li Shi. Graph-GRPO: Stabilizing multi-agent topology learning via group relative policy optimization. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 20222–20231, 2026. doi: 10.18653/v1/2026.findings-acl.1010. URL <https://aclanthology.org/2026.findings-acl.1010/>.
- Jiarui Che, Yifei Chen, Zhixing Tian, Chenyang Wang, and Ziguang Cheng. Graph-GRPO: Dependency-aware credit assignment for generative e-commerce search relevance. *arXiv preprint arXiv:2605.31003*, 2026. URL <https://arxiv.org/abs/2605.31003>.
- Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, and Gao Huang. Does reinforcement learning really incentivize reasoning capacity in llms beyond the base model? *Advances in Neural Information Processing Systems*, 38:57654–57689, 2026.
- Gustave Choquet. Theory of capacities. *Annales de l’Institut Fourier*, 5:131–295, 1954. doi: 10.5802/aif.53. URL <https://www.numdam.org/articles/10.5802/aif.53/>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, et al. The entropy mechanism of reinforcement learning for reasoning language models. *arXiv preprint arXiv:2505.22617*, 2025.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Yubo Dong, Hehe Fan, Linchao Zhu, and Yi Yang. Structured reasoning for LLMs: A unified framework for efficiency and explainability. In *International Conference on Learning Representations*, pp. 106183–106215, 2026a. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/ad5b3f324b24c17cdc2f3712298c76bd-Paper-Conference.pdf.
- Zhichen Dong, Yang Li, Yuhan Sun, Weixun Wang, Yijia Luo, Zinian Peng, Taiheng Ye, Chao Yang, Wenbo Su, Yu Cheng, Bo Zheng, and Junchi Yan. How does reasoning flow? tracing attention-induced information flow for targeted RL in LLMs. In *International Conference on Machine Learning*, 2026b. URL <https://arxiv.org/abs/2606.10646>.
- Zenghao Duan, Liang Pang, Zihao Wei, Wenbin Duan, Yuxin Tian, Shicheng Xu, Jingcheng Deng, Zhiyi Yin, and Xueqi Cheng. Circular reasoning: Understanding self-reinforcing loops in large reasoning models. *arXiv preprint arXiv:2601.05693*, 2026. URL <https://arxiv.org/abs/2601.05693>.
- Jiwei Fang, Bin Zhang, Changwei Wang, Jin Wan, and Zhiwei Xu. Graph of verification: Structured verification of llm reasoning with directed acyclic graphs. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 40, pp. 30665–30672, 2026. doi: 10.1609/aaai.v40i36.40322.
- Xiao Feng, Bo Han, Zhanke Zhou, Jiaqi Fan, Jiangchao Yao, Ka Ho Li, Dahai Yu, and Michael Kwok-Po Ng. Rewardflow: Topology-aware reward propagation on state graphs for agentic rl with large language models. *arXiv preprint arXiv:2603.18859*, 2026. URL <https://arxiv.org/abs/2603.18859>.
- Leo Gao, John Schulman, and Jacob Hilton. Scaling laws for reward model overoptimization. In *International Conference on Machine Learning*, pp. 10835–10866. PMLR, 2023.
- Michel Grabisch. k -order additive discrete fuzzy measures and their representation. *Fuzzy Sets and Systems*, 92(2):167–189, 1997. doi: 10.1016/S0165-0114(97)00168-1.

-
- Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *International Conference on Learning Representations*, 2024.
- Xinyu Guan, Li Lyna Zhang, Yifei Liu, Ning Shang, Youran Sun, Yi Zhu, Fan Yang, and Mao Yang. rstar-math: Small llms can master math reasoning with self-evolved deep thinking, 2025. URL <https://arxiv.org/abs/2501.04519>.
- Peixuan Han, Adit Krishnan, Gerald Friedland, Jiaxuan You, and Luyang Kong. Self-aligned reward: Towards effective and efficient reasoners. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/fc30caeb45721bab13507c50199e6403-Abstract-Conference.html.
- Tao He, Rongchuan Mu, Lizi Liao, Yixin Cao, Li Yang, Yijia Luo, Weixun Wang, Ming Liu, and Bing Qin. Smarter not harder: Generative process evaluation with intrinsic-signal driving and ability-adaptive reward shaping. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/da842df0d1e35e15a521267acb62bbd0-Abstract-Conference.html.
- Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.
- Jian Hu. Reinforce++: A simple and efficient approach for aligning large language models. *arXiv e-prints*, pp. arXiv-2501, 2025.
- Jingcheng Hu, Yinmin Zhang, Qi Han, Daxin Jiang, Xiangyu Zhang, and Heung-Yeung Shum. Open-reasoner-zero: An open source approach to scaling up reinforcement learning on the base model. *Advances in Neural Information Processing Systems*, 38:162239–162262, 2026.
- Shuyang Jiang, Yusheng Liao, Ya Zhang, Yanfeng Wang, and Yu Wang. Overthinking reduction with decoupled rewards and curriculum data scheduling. In *International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=kdeiRledV6>. Oral.
- Jiajie Jin, Zhao Yang, Wenle Liao, Yuyang Hu, Guanting Dong, Xiaoxi Li, Yutao Zhu, and Zhicheng Dou. Verigraph: Towards verifiable data-analytic agents. *arXiv preprint arXiv:2606.16603*, 2026. URL <https://arxiv.org/abs/2606.16603>.
- Aryan Karmore. Atomgraph: Reasoning isn’t linear, why should verification be? In *ICLR 2026 Workshop on LLM Reasoning*, 2026. URL <https://openreview.net/forum?id=dNsIZln2uv>.
- Amirhossein Kazemnejad, Milad Aghajohari, Eva Portelance, Alessandro Sordoni, Siva Reddy, Aaron Courville, and Nicolas Le Roux. VinePPO: Refining credit assignment in RL training of LLMs. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025.
- Zixuan Ke, Fangkai Jiao, Yifei Ming, Xuan-Phi Nguyen, Austin Xu, Do Xuan Long, Minzhi Li, Chengwei Qin, Peifeng Wang, Silvio Savarese, Caiming Xiong, and Shafiq Joty. A survey of frontiers in llm reasoning: Inference scaling, learning to reason, and agentic systems, 2026. URL <https://arxiv.org/abs/2504.09037>.
- Muhammad Khalifa, Rishabh Agarwal, Lajanugen Logeswaran, Jaekyeom Kim, Hao Peng, Moon-tae Lee, Honglak Lee, and Lu Wang. Process reward models that think. *arXiv preprint arXiv:2504.16828*, 2025.
- Jongwoo Ko, Sungnyun Kim, Tianyi Chen, and Se-Young Yun. Distillm: Towards streamlined distillation for large language models. *arXiv preprint arXiv:2402.03898*, 2024.
- Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. Large language models are zero-shot reasoners. *Advances in neural information processing systems*, 35: 22199–22213, 2022.

- Xin Lai, Zhuotao Tian, Yukang Chen, Senqiao Yang, Xiangru Peng, and Jiaya Jia. Step-dpo: Step-wise preference optimization for long-chain reasoning of llms. *arXiv preprint arXiv:2406.18629*, 2024.
- Nathan Lambert, Jacob Morrison, Valentina Pyatkin, Shengyi Huang, Hamish Ivison, Faeze Brahman, Lester James V Miranda, Alisa Liu, Nouha Dziri, Xinxu Lyu, Yuling Gu, Saumya Malik, Victoria Graf, Jena D Hwang, Jiangjiang Yang, Ronan Le Bras, Oyvind Tafjord, Chris Wilhelm, Luca Soldaini, Noah A Smith, Yizhong Wang, Pradeep Dasigi, and Hannaneh Hajishirzi. Tülu 3: Pushing frontiers in open language model post-training. *arXiv preprint arXiv:2411.15124*, 2024.
- Gang Li, Yan Chen, Ming Lin, and Tianbao Yang. DRPO: Efficient reasoning via decoupled reward policy optimization. In *International Conference on Learning Representations*, 2026a. arXiv:2510.04474.
- Jiahui Li, Hanlin Zhang, Fengda Zhang, Tai-Wei Chang, Kun Kuang, Long Chen, and Jun Zhou. Optimizing language models with fair and stable reward composition in reinforcement learning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pp. 10122–10140. Association for Computational Linguistics, 2024. doi: 10.18653/v1/2024.emnlp-main.565. URL <https://aclanthology.org/2024.emnlp-main.565/>.
- Rui Li and Shuang Cao. From trajectories to graphs: Contract-checked editing for verifier-guided LLM reasoning. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics*, pp. 43259–43306, 2026. doi: 10.18653/v1/2026.acl-long.2004. URL <https://aclanthology.org/2026.acl-long.2004/>.
- Zuchao Li, Qiwei Li, Yao Yao, Hai Zhao, Lefei Zhang, and Bo Du. GoT-R1: Internalizing graph-of-thought via structural reinforcement for high-density reasoning. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 7090–7104, 2026b. doi: 10.18653/v1/2026.findings-acl.352. URL <https://aclanthology.org/2026.findings-acl.352/>.
- Hunter Lightman, Vineet Kosaraju, Yura Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. *arXiv preprint arXiv:2305.20050*, 2023.
- Mingjie Liu, Shizhe Diao, Ximing Lu, Jian Hu, Xin Dong, Yejin Choi, Jan Kautz, and Yi Dong. ProRL: Prolonged reinforcement learning expands reasoning boundaries in large language models. *Advances in Neural Information Processing Systems*, 2025.
- Runxuan Liu, Xianhao Ou, Xinyan Ma, Jiyuan Wang, Jiafeng Liang, Jiaqi Li, Tao He, Zheng Chu, Rongchuan Mu, Zekun Wang, et al. Graph reasoning paradigm: Structured and symbolic reasoning with topology-aware reinforcement learning for large language models. *arXiv preprint arXiv:2601.12995*, 2026a.
- Wei Liu, Ruochen Zhou, Yiyun Deng, Yuzhen Huang, Junteng Liu, Yuntian Deng, Yizhe Zhang, and Junxian He. Learn to reason efficiently with adaptive length-based reward shaping. In *International Conference on Learning Representations*, 2026b. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/47795c4ae2f7d07ea2fb0d11fa2c3c90-Abstract-Conference.html.
- Jianqiao Lu, Zhiyang Dou, Hongru Wang, Zeyu Cao, Jianbo Dai, Yingjia Wan, Yunlong Feng, and Zhijiang Guo. Autopsv: Automated process-supervised verifier. *Advances in Neural Information Processing Systems*, 37:79935–79962, 2024.
- Haotian Luo, Li Shen, Haiying He, Yibo Wang, Shiwei Liu, Wei Li, Naiqiang Tan, Xiaochun Cao, and Dacheng Tao. O1-pruner: Length-harmonizing fine-tuning for o1-like reasoning pruning. *arXiv preprint arXiv:2501.12570*, 2025.
- Liangchen Luo, Yinxiao Liu, Rosanne Liu, Samrat Phatale, Harsh Lara, Yunxuan Li, Lei Shu, Yun Zhu, Lei Meng, Jiao Sun, and Abhinav Rastogi. Improve mathematical reasoning in language models by automated process supervision. *arXiv preprint arXiv:2406.06592*, 2024.

- Kohsei Matsutani, Shota Takashiro, Gouki Minegishi, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo. RL squeezes, SFT expands: A comparative study of reasoning LLMs. In *International Conference on Learning Representations*, pp. 141482–141526, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/e52554a70e0df57a0bea11d1eca0c9b5-Paper-Conference.pdf.
- Gouki Minegishi, Hiroki Furuta, Takeshi Kojima, Yusuke Iwasawa, and Yutaka Matsuo. Topology of reasoning: Understanding large reasoning models through reasoning graph properties. In *Advances in Neural Information Processing Systems*, volume 38, 2025. URL https://papers.neurips.cc/paper_files/paper/2025/hash/147521db4d049c75cf18d9ff96bf68f5-Abstract-Conference.html.
- Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori B Hashimoto. s1: Simple test-time scaling. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 20286–20332, 2025.
- Sagnik Mukherjee, Abhinav Chinta, Takyoun Kim, Tarun Anoop Sharma, and Dilek Hakkani-Tür. Premise-augmented reasoning chains improve error identification in math reasoning with LLMs. In *International Conference on Machine Learning*, 2025. URL <https://arxiv.org/abs/2502.02362>.
- NVIDIA. OpenReasoning-Nemotron-7B model card. Hugging Face model release, 2025. URL <https://huggingface.co/nvidia/OpenReasoning-Nemotron-7B>.
- NVIDIA. AceMath-RL-Nemotron-7B model card. Hugging Face model release, 2026. URL <https://huggingface.co/nvidia/AceMath-RL-Nemotron-7B>.
- Amruta Parulekar, Jinu Lee, Dilek Hakkani-Tür, and Hari Sundaram. Reasoning consensus: Structural ensembling of LLM reasoning via weighted DAG aggregation. *arXiv preprint arXiv:2607.27783*, 2026. URL <https://arxiv.org/abs/2607.27783>.
- Qwen Team. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Qwen Team. Qwen3.5-9B model card. Hugging Face model release, 2026. URL <https://huggingface.co/Qwen/Qwen3.5-9B>.
- Pier Giuseppe Sessa, Robert Dadashi, Léonard Hussenot, Johan Ferret, Nino Vieillard, Alexandre Ramé, Bobak Shahriari, Sarah Shariari, Sertan Perrin, Abe Friesen, Geoffrey Cideron, Sertan Girgin, Piotr Stanczyk, Andrea Michi, Danila Sinopalnikov, Sabela Ramos, Morgane Riviere, Joaquin Lin, Dylan Hadfield-Menell, Rémi Munos, Olivier Pietquin, and Olivier Bachem. BOND: Aligning LLMs with best-of-N distillation. *arXiv preprint arXiv:2407.14622*, 2024.
- Amrith Setlur, Chirag Nagpal, Adam Fisch, Xinyang Geng, Jacob Eisenstein, Rishabh Agarwal, Alekh Agarwal, Jonathan Berant, and Aviral Kumar. Rewarding progress: Scaling automated process verifiers for llm reasoning. In *International Conference on Learning Representations*, volume 2025, pp. 60808–60838, 2025.
- Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y K Li, Y Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models. *arXiv preprint arXiv:2402.03300*, 2024.
- Shuaijie She, Junxiao Liu, Yifeng Liu, Jiajun Chen, Xin Huang, and Shujian Huang. R-PRM: Reasoning-driven process reward modeling. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025.
- Joar Skalse, Nikolaus Howe, Dmitrii Krashennnikov, and David Krueger. Defining and characterizing reward gaming. *Advances in Neural Information Processing Systems*, 35:9460–9471, 2022.
- Charlie Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling llm test-time compute optimally can be more effective than scaling model parameters, 2024. URL <https://arxiv.org/abs/2408.03314>.

- Mingyang Song, Zhaochen Su, Xiaoye Qu, Jiawei Zhou, and Yu Cheng. PRMBench: A fine-grained and challenging benchmark for process-level reward models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*, 2025.
- Zhenpeng Su, Leiyu Pan, Xue Bai, Dening Liu, Guanting Dong, Jiaming Huang, Wenping Hu, Fuzheng Zhang, Kun Gai, and Guorui Zhou. Klear-Reasoner: Advancing reasoning capability via gradient-preserving clipping policy optimization. *arXiv preprint arXiv:2508.07629*, 2025.
- Leitian Tao, Ilia Kulikov, Swarnadeep Saha, Tianlu Wang, Jing Xu, Sharon Li, Jason E. Weston, and Ping Yu. Hybrid reinforcement: When reward is sparse, better to be dense. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/dcd966df09d8771a3db592c5fe4c587-Abstract-Conference.html.
- Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. Solving math word problems with process-and outcome-based feedback. *arXiv preprint arXiv:2211.14275*, 2022.
- Boxin Wang, Chankyu Lee, Nayeon Lee, Sheng-Chieh Lin, Wenliang Dai, Yang Chen, Yangyi Chen, Zhuolin Yang, Zihan Liu, Mohammad Shoeybi, Bryan Catanzaro, and Wei Ping. Nemotron-cascade: Scaling cascaded reinforcement learning for general-purpose reasoning models. *arXiv preprint arXiv:2512.13607*, 2025a. URL <https://arxiv.org/abs/2512.13607>.
- Hongru Wang, Deng Cai, Wanjun Zhong, Shijue Huang, Jeff Z. Pan, Zeming Liu, and Kam-Fai Wong. Self-reasoning language models: Unfold hidden reasoning chains with few reasoning catalyst, 2025b. URL <https://arxiv.org/abs/2505.14116>.
- Peiyi Wang, Lei Li, Zhihong Shao, Runxin Xu, Damai Dai, Yifei Li, Deli Chen, Yu Wu, and Zhifang Sui. Math-shepherd: Verify and reinforce llms step-by-step without human annotations. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 9426–9439, 2024.
- Qinsi Wang, Jing Shi, Huazheng Wang, Kun Wan, Yiran Wu, Bo Liu, Qingyun Wu, Hai Helen Li, Yiran Chen, Handong Zhao, and Wentian Zhao. From RLVR to RLSVR: Task transformation induces self-verifiable rewards for open-ended LLM self-improvement. In *Conference on Language Modeling*, 2026a. URL <https://huggingface.co/SpyRL/SpyRL-Qwen3-8B-Math>.
- Xinpeng Wang, Nitish Joshi, Barbara Plank, Rico Angell, and He He. Is it thinking or cheating? detecting implicit reward hacking by measuring reasoning effort. In *International Conference on Learning Representations*, 2026b. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/9d0ee47d732bd526c4a03d4aecc39bc1-Abstract-Conference.html.
- Yifan Wang, Bolian Li, David Cho, Ruqi Zhang, Fanping Sui, and Ananth Grama. SARL: Label-free reinforcement learning by rewarding reasoning topology. *arXiv preprint arXiv:2603.27977*, 2026c. URL <https://arxiv.org/abs/2603.27977>.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.
- Xiangyu Wen, Min Li, Junhua Huang, Jianyuan Zhong, Zhijian Xu, Zeju Li, Yongxiang Huang, Mingxuan Yuan, and Qiang Xu. Dependency matters: Enhancing LLM reasoning with explicit knowledge grounding. In *Advances in Neural Information Processing Systems*, volume 38, 2025. URL https://proceedings.neurips.cc/paper_files/paper/2025/hash/5401bede0941e568fa42b7b5910fa72c-Abstract-Conference.html.
- Xiangyu Wen, Junhua Huang, Zeju Li, Min Li, Jianyuan Zhong, Zhijian Xu, Mingxuan Yuan, Yongxiang Huang, and Qiang Xu. Reasoning scaffolding: Distilling the flow of thought from LLMs. In *International Conference on Learning Representations*, pp. 80433–80459, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/8248b1ded388fcdabd121bcdfea3068c-Paper-Conference.pdf.

- Heming Xia, Chak Tou Leong, Wenjie Wang, Yongqi Li, and Wenjie Li. Tokenskip: Controllable chain-of-thought compression in llms. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pp. 3351–3363, 2025.
- Xiaomi LLM Core Team. MiMo: Unlocking the reasoning potential of language model—from pretraining to posttraining. *arXiv preprint arXiv:2505.07608*, 2025. URL <https://arxiv.org/abs/2505.07608>.
- Hongling Xu, Qi Zhu, and Ruifeng Xu. KDRL: Post-training reasoning LLMs via unified knowledge distillation and reinforcement learning. *arXiv preprint arXiv:2506.02208*, 2025.
- An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024a.
- An Yang, Beichen Zhang, Binyuan Hui, Bofei Gao, Bowen Yu, Chengpeng Li, Dayiheng Liu, Jianhong Tu, Jingren Zhou, Junyang Lin, et al. Qwen2. 5-math technical report: Toward mathematical expert model via self-improvement. *arXiv preprint arXiv:2409.12122*, 2024b.
- Wenkai Yang, Weijie Liu, Ruobing Xie, Kai Yang, Saiyong Yang, and Yankai Lin. Learning beyond teacher: Generalized on-policy distillation with reward extrapolation. *arXiv preprint arXiv:2602.12125*, 2026.
- Tianzhu Ye, Li Dong, Zewen Chi, Xun Wu, Shaohan Huang, and Furu Wei. Black-box on-policy distillation of large language models. *arXiv preprint arXiv:2511.10643*, 2025.
- Edward Yeo, Yuxuan Tong, Xinyao Niu, Graham Neubig, and Xiang Yue. Demystifying long chain-of-thought reasoning in LLMs. *arXiv preprint arXiv:2502.03373*, 2025.
- Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Weinan Dai, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An open-source LLM reinforcement learning system at scale. *arXiv preprint arXiv:2503.14476*, 2025.
- Zewei Yu, Lirong Gao, Yuke Zhu, Bo Zheng, Junbo Zhao, Sheng Guo, and Haobo Wang. Stop unnecessary reflection: Training LRMs for efficient reasoning with adaptive reflection and length coordinated penalty. In *International Conference on Learning Representations*, 2026. URL https://proceedings.iclr.cc/paper_files/paper/2026/hash/08fe50bf209c57eefc0804f9f9ed639f-Abstract-Conference.html.
- Hongyuan Yuan, Xinran He, Run Shao, Bolei He, Xianwei Xue, Mengke Chen, Qitong Pan, Haiwei Wang, and Haifeng Li. Graph-based chain-of-thought pruning for reducing redundant reflections in reasoning LLMs. *Findings of the Association for Computational Linguistics: ACL 2026*, 2026. URL <https://arxiv.org/abs/2604.05643>.
- Lifan Yuan, Wendi Li, Huayu Chen, Ganqu Cui, Ning Ding, Kaiyan Zhang, Bowen Zhou, Zhiyuan Liu, and Hao Peng. Free process rewards without process labels. In *Proceedings of the 42nd International Conference on Machine Learning*, 2025.
- Weihaio Zeng, Yuzhen Huang, Qian Liu, Wei Liu, Keqing He, Zejun Ma, and Junxian He. Simplerl-zoo: Investigating and taming zero reinforcement learning for open base models in the wild, 2025. URL <https://arxiv.org/abs/2503.18892>, 2023.
- Dan Zhang, Sining Zhoubian, Ziniu Hu, Yisong Yue, Yuxiao Dong, and Jie Tang. Rest-mcts*: Llm self-training via process reward guided tree search. *Advances in Neural Information Processing Systems*, 37:64735–64772, 2024.
- Yuanhe Zhang, Ilja Kuzborskij, Jason Lee, Chenlei Leng, and Fanghui Liu. DAG-Math: Graph-of-thought guided mathematical reasoning in LLMs. In *International Conference on Learning Representations*, pp. 74347–74389, 2026a. URL https://proceedings.iclr.cc/paper_files/paper/2026/file/79179a9c80561608fa85f925cc88aaa9-Paper-Conference.pdf.

-
- Zheng Zhang, Ziwei Shan, Kaitao Song, Yexin Li, and Kan Ren. Linking process to outcome: Conditional reward modeling for LLM reasoning. In *International Conference on Learning Representations*, 2026b. URL <https://openreview.net/forum?id=4DJ0B0QNd0>.
- Zhenru Zhang, Chujie Zheng, Yangzhen Wu, Beichen Zhang, Runji Lin, Bowen Yu, Dayiheng Liu, Jingren Zhou, and Junyang Lin. The lessons of developing process reward models in mathematical reasoning. *arXiv preprint arXiv:2501.07301*, 2025.
- Zijing Zhang, Ziyang Chen, Mingxiao Li, Zhaopeng Tu, and Xiaolong Li. RLVMR: Reinforcement learning with verifiable meta-reasoning rewards for robust long-horizon agents. In *International Conference on Learning Representations*, 2026c. URL <https://iclr.cc/virtual/2026/poster/10008499>.
- Jian Zhao, Runze Liu, Kaiyan Zhang, Zhimu Zhou, Junqi Gao, Dong Li, Jiafei Lyu, Zhouyi Qian, Biqing Qi, Xiu Li, and Bowen Zhou. GenPRM: Scaling test-time compute of process reward models via generative reasoning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2026.
- Chujie Zheng, Zhenru Zhang, Beichen Zhang, Runji Lin, Keming Lu, Bowen Yu, Dayiheng Liu, Jingren Zhou, and Junyang Lin. Processbench: Identifying process errors in mathematical reasoning. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pp. 1009–1024, 2025.
- Tianjun Zhong, Linyang He, Ziyang Li, and Nima Mesgarani. From chains to dags: Probing the graph structure of reasoning in llms. *arXiv preprint arXiv:2601.17593*, 2026. URL <https://arxiv.org/abs/2601.17593>.
- Xinyu Zhu, Mengzhou Xia, Zhepei Wei, Wei-Lin Chen, Danqi Chen, and Yu Meng. The surprising effectiveness of negative reinforcement in LLM reasoning. *Advances in Neural Information Processing Systems*, 2025.

A IMPLEMENTATION DETAILS

A.1 HYPERPARAMETERS

Table 5 lists the key hyperparameters for each training stage.

Table 5: Hyperparameter summary for the three training stages.

Parameter	Value	Description
<i>Stage I: Supervised Fine-Tuning (SFT)</i>		
Base model	Qwen3.5-9B	–
Adapter	LoRA $r=64$, $\alpha=128$	all-linear modules
Learning rate	5×10^{-5}	cosine, 3% warmup
Epochs	3	–
Selected checkpoint	step 1,100	minimum held-out loss 0.381609; final step 1,734
Batch size	1/device, accum 8	global batch 32 on 4 GPUs
Max seq length	4096	–
Precision	bf16	–
<i>Stage II: GRPO with TopoPRM Rewards</i>		
Base model	SFT step 1,100	held-out-loss selection
Adapter	LoRA $r=64$, $\alpha=128$	all-linear modules
Learning rate	5×10^{-6}	cosine, 5% warmup
Gradient steps	300	–
Batch size	1/device, accum 4	global prompt batch 16 on 4 GPUs
Num generations	4 per prompt	for GRPO grouping
Temperature	0.9	rollout sampling
β (KL penalty)	0.04	–
Max completion length	4096	prompt capped at 512
Precision	bf16	–
<i>Stage III: Topology-Guided Distillation</i>		
Teacher model	Stage-II 9B	frozen during distillation
Student model	Qwen3.5-4B	canonical cross-scale run
Adapter	LoRA $r=64$, $\alpha=128$	all-linear modules
Learning rate	2×10^{-5}	cosine, 3% warmup
Epochs	1	–
Batch size	2/device, accum 4	global batch 32 on 4 GPUs
Student temperature	0.7	student-sampled trace
Teacher temperature	0.0	greedy revision target
Target budget	1024 tokens	fixed across methods
Target selection	answer, format, direction, acyclicity	threshold calibrated on validation
Max seq length	4096	–
Precision	bf16	–

A.2 REWARD CONFIGURATION

The Choquet reward in Eq. 3 uses the five independent sources (r_{out} , r_{fmt} , q_{dir} , q_{acyc} , q_{cont}) and the interaction set $\{(o, d), (o, a), (d, c), (a, c), (d, a)\}$. **The exact non-negative capacity coefficients are pending author confirmation and are not yet a training or paper claim.** Once locked, the coefficients will sum to one and remain fixed across the canonical run and its matched controls. The additive control uses the capacity’s Shapley source importance; the multiplicative control uses those same exponents. No length or reference-solution graph enters the canonical main reward because generated traces have no node alignment with a reference solution.

A.3 HARDWARE AND INFRASTRUCTURE

Canonical ICLR runs use $4 \times$ NVIDIA L20 GPUs. We report measured wall-clock time and peak memory from this configuration rather than extrapolating from the earlier A100 runs.

A.4 EVALUATION PROTOCOL

All public benchmark evaluations use the Transformers generation backend. Pass@1 is one sampled response per item from a single seed-0 evaluation pass; this is neither a multi-seed training result nor an average over repeated decodes. The shared decoder uses temperature 0.6, top- $p=0.95$, top- $k=20$, min- $p=0$, and repetition penalty 1.0. This matches Qwen3 and DeepSeek-R1’s reasoning-time recommendations and keeps the public and adapted rows comparable; it differs from Qwen3.5’s general-task recommendation only in using temperature 0.6 instead of 1.0. Generation explicitly enables the standard key-value cache even if a released checkpoint retains a training-time `use_cache=false` setting; this is recorded in the result manifest. The four primary mathematics benchmarks use the same 4,096-token completion cap and the same instruction to provide a step-by-step solution with the final answer in `\boxed{\}`. We retain each released chat template and documented control: GoT-R1, OpenReasoning-Nemotron, and AceMath receive a user-only task; MiMo uses an explicit empty system turn; the released Nemotron-Cascade template injects the terminal “/think” control; and the AceMath, Qwen-family, and Klear templates activate their released reasoning format without an evaluator-added prefix. No model receives an uncounted response prefix. For mathematics tasks, the evaluator verifies the last nonempty boxed answer, or a short terminal region after the last explicit final-answer marker, against the dataset target; it never searches an unsubmitted truncated derivation for a matching intermediate expression. The reference is passed first to the asymmetric `math_verify` verifier. Visible length is obtained by retokenizing the decoded response without special tokens. Pass@5 and maj@5, when run as secondary diagnostics, draw five samples from the same decoder and reuse this equivalence relation rather than raw string equality. The manifest stores the message-role protocol, prefix flag, decoder parameters, evaluation seed, batch size, model and adapter paths, dataset fingerprint, and full per-example responses; cached results are reused only when these protocol fields match.

B CHOQUET CAPACITY AND BASELINES

For non-negative Möbius coefficients summing to one, Eq. 3 is a normalized monotone 2-additive capacity (Choquet, 1954; Grabisch, 1997): it maps **0** to 0 and **1** to 1 and remains in $[0, 1]$. We separately replace the direction-acyclicity interaction with equal singleton mass to preserve both normalization and Shapley source importance, distinguishing the proposed structural synergy from merely increasing either source’s weight. We also test traces that are structurally plausible but answer-incorrect to detect reward hacking.

Matched baseline protocols. The LASER-D adaptation uses the true response-token count and four rollouts per prompt. It labels a prompt high, medium, or low by rollout accuracy thresholds $2/3$ and $1/3$, and adds 0.5 only to correct rollouts at or below the band’s target. Every 20 steps, the target is recomputed from a synchronized rolling monitor of 500 prompts. To match our 4,096-token cap, the released 4,096–16,384 candidate range is scaled by four to 1,024–4,096 in increments of 256; no character-length proxy is used. The matched additive policy control uses the Choquet capacity’s Shapley source importance, preserving marginal source weights while removing every interaction. Equal additive, matched multiplicative, and direction-acyclicity-interaction controls are evaluated on the identical frozen four-candidate pool rather than presented as separately trained policies.

C DAG EXTRACTION DETAILS

Extraction pipeline. Given a reasoning trace in `<think>`, we apply five frozen stages:

1. Step segmentation: split by numbered-step markers, mathematical discourse markers (`:`, `;`, `.`), and line boundaries.
2. Sub-question partitioning: markers such as “Sub-question k ” define connected components; cross-sub-question edges are disallowed.
3. Information extraction: for each step, extract expressions, claims, and variables with rule-based patterns.
4. Raw support encoding: combine high-precision symbolic candidates with a lightweight pair classifier that selects $i \rightarrow j$, $j \rightarrow i$, or abstention without imposing text order.

Table 6: Detailed diagnostics for the frozen candidate pool in Table 3. Tie accuracy averages correctness across all top-scoring candidates rather than choosing the first; tie and zero-spread rates expose uninformative ranking groups. **Cells await the canonical pool and report item-bootstrap intervals, not training variance.**

Aggregator	Tie acc.↑	Oracle acc.↑	Tie groups↓	Zero-spread↓	StructWrong↓
Outcome only	—	—	—	—	—
Equal-weight additive	—	—	—	—	—
Shapley-matched additive	—	—	—	—	—
Shapley-matched multiplicative	—	—	—	—	—
Choquet w/o dir.-acyc. interaction	—	—	—	—	—
Topology-conditioned Choquet	—	—	—	—	—

5. Projection: retain the raw graph for diagnostics, remove backward edges to obtain the training DAG, then add adaptive weak sequential edges only when support is absent.

Information-density filters. To avoid low-information nodes and noisy links: (a) claims are extracted at sentence level, with heading fragments and unfinished prefixes removed; (b) trivial tokens (single symbols, pure numerals, non-relational fragments) are excluded from dependency matching.

Dependency rules. For each unordered pair of steps in the same sub-question, rules propose candidates from informative expression reuse, canonicalized claim overlap, variable roles, and explicit step references. The pair classifier assigns direction or abstains. Only after raw-graph diagnostics are computed do we insert a fallback edge for an unsupported derivation or conclusion in the projected DAG.

Pair-encoder adaptation. The lightweight encoder is Qwen3.5-0.8B with a rank-8 LoRA three-way classification head. The frozen 9B semantic teacher processes 595 candidate source traces; 536 contain parsed semantic edges, and the shared strict segmentation filter rejects one dangling clause. The remaining 535 sources are hash-split before order augmentation into 428 training and 107 validation sources. After original, reversed, and interleaved presentations, deterministic negative subsampling yields 9,301 training pairs; the unsampled validation set contains 2,897 pairs. On these teacher-held-out pairs, the selected epoch-3 encoder reaches 0.7784 three-class macro-F1, 0.7711 edge precision, 0.7603 edge recall, and 0.7375 direction accuracy on gold edges. The calibrated edge-confidence and direction-margin thresholds are 0.70 and 0.20. This teacher-held-out diagnostic is not the independent human result in Table 4. The encoder is frozen during policy optimization; a retained encoder edge keeps its calibrated confidence as w_{ij} , while auditable rule edges use $w_{ij} = 1$. The semantic-teacher cache, adapter weights, corrected 120-trace audit pack, and all-order frozen prediction ledger have SHA-256 fingerprints 69131f16...b19c, 9e8e78ac...3212, 95c440a3...b53, and 8c7a644e...9328, respectively; all order-control statistics are paired by source record.

Controlled edge noise. Table 7 measures how the two raw-graph diagnostics respond when frozen predictions are corrupted after extraction. Deletion probes missed edges, false-edge addition probes spurious support, and direction reversal probes orientation errors. The audit uses one deterministic perturbation draw per source trace and paired bootstrap across traces; it is not a repeated-training or multi-seed result.




Figure 3: Controlled response of the frozen raw-graph encoder. Each point is the paired mean change relative to the identical source trace; bars are paired source-trace bootstrap 95% intervals. Reversal primarily changes direction, matched cycle closure changes acyclicity at fixed backward mass, and stochastic edge corruption moves both diagnostics according to its targeted failure. These are deterministic perturbation audits, not downstream policy results or training-seed variation.

Table 8: Diagnostic Best-of-8 selection accuracy (%) on one shared Qwen2.5-Math-7B-Instruct candidate pool ($n=80$ per benchmark; 2,048-token generation cap). PRM uses Qwen2.5-Math-PRM-7B; Hybrid multiplies normalized PRM by $1 + q_{\text{topo}}$. Oracle reports whether any candidate is correct. This diagnostic is separate from the primary 4,096-token policy comparison.

Benchmark	Pass@1	Maj@8	PRM@8	Topo@8	Hybrid@8	Oracle@8
GSM8K	93.75	95.00	95.00	91.25	95.00	96.25
MATH-500	83.75	85.00	83.75	78.75	83.75	92.50

Table 7: Frozen-encoder sensitivity on 120 source traces. One deterministic seed-0 perturbation is applied per trace; intervals are paired source-trace bootstrap 95% intervals. p_{real} is the realized modified-edge fraction. Deltas are corrupted minus original.

Perturbation	p	p_{real}	Δq_{dir}	Δq_{acyc}	$\Delta \text{Cov.}$	$\Delta \text{Cycle mass}$
Random edge deletion	.10	.068	-.017 [-.031, -.005]	-.012 [-.029, +.003]	-.033	-.006
	.20	.161	-.027 [-.047, -.008]	-.042 [-.061, -.023]	-.091	-.006
	.30	.274	-.067 [-.092, -.044]	-.073 [-.097, -.049]	-.154	-.005
Random false-edge addition	.10	.046	-.013 [-.021, -.004]	-.043 [-.065, -.023]	+.008	+.048
	.20	.124	-.034 [-.050, -.019]	-.099 [-.136, -.064]	+.013	+.111
	.30	.244	-.073 [-.095, -.051]	-.190 [-.244, -.138]	+.031	+.217
Random direction flip	.10	.097	-.052 [-.084, -.023]	-.018 [-.036, -.003]	.000	+.021
	.20	.214	-.104 [-.153, -.057]	-.034 [-.064, -.006]	.000	+.037
	.30	.300	-.180 [-.234, -.129]	-.052 [-.088, -.020]	.000	+.059

Topology is not a standalone correctness verifier. Table 8 rescales the rebuttal-era shared Best-of-8 pool with the canonical gold-first symbolic verifier. Topology-only selection underperforms the dedicated PRM on both subsets, and their simple product ties rather than improves on the PRM. We therefore use topology as an auditable training-time structural source and do not claim that it replaces semantic verification. The candidate-pool and cached-score ledgers have SHA-256 fingerprints 2f4ded08...2664 and 8304d8fc...b7c8.

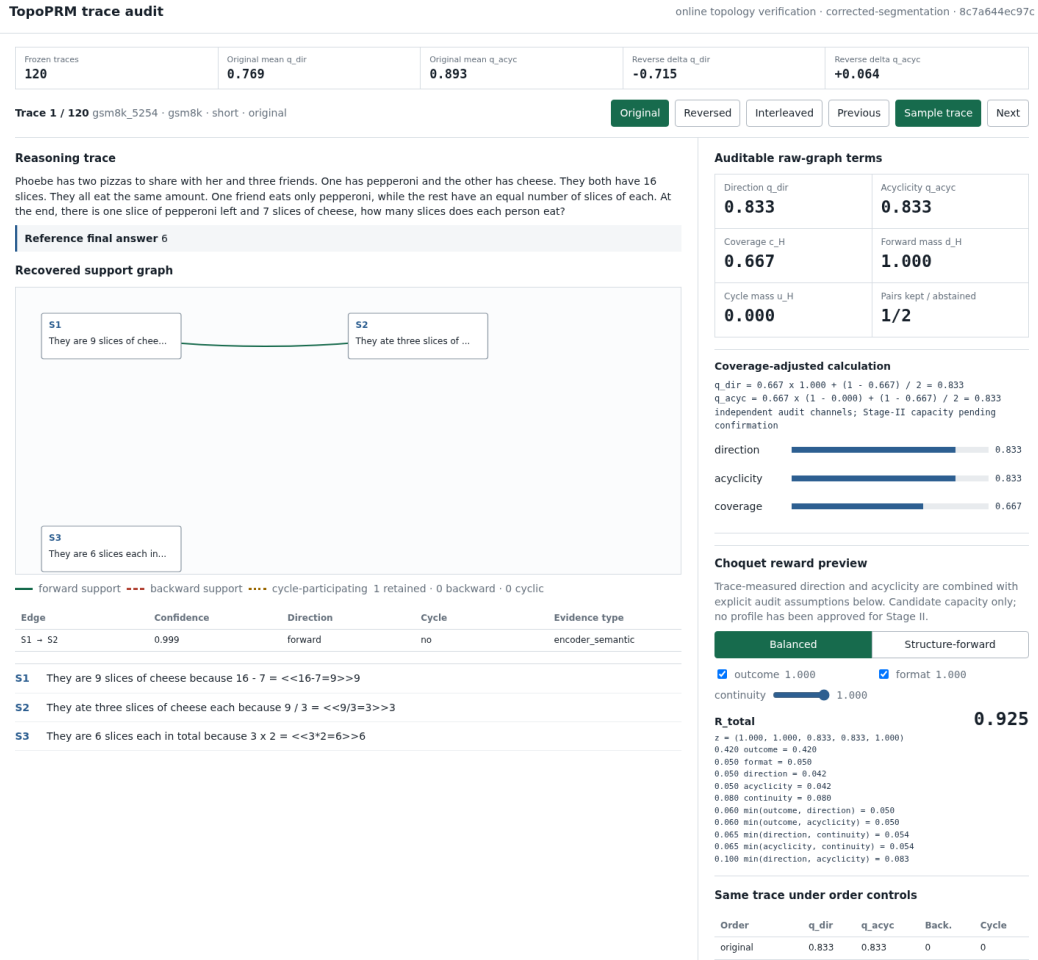


Figure 4: Trace-level topology audit interface. For each sampled reasoning trace, the dashboard exposes the recovered support graph, its edge ledger, and independently recomputed coverage-adjusted direction and acyclicity channels. The interactive reward preview substitutes these trace-measured channels into either preregistered candidate capacity while outcome, format, and continuity remain explicit user-controlled assumptions; it is not a canonical Stage-II result. Original, reversed, and interleaved controls switch the same trace online. Three linked annotation copies hide model predictions and internal record identities, independently shuffle trace presentation, and retain a matched step-order condition for agreement. The illustrated corrected-encoder trace is one auditable sample, not an aggregate evaluation result.

Adaptive sequential weak edges. A weak sequential edge (weight 0.3) is added only when the target step lacks strong incoming support and the source step does not already emit strong outgoing dependencies. This prevents chain-dominant artifacts in long traces.

Node typing. Each node is classified into DEFINITION, DERIVATION, COMPUTATION, CONCLUSION, AUXILIARY, SUBSTITUTION, or CASE_ANALYSIS by keyword rules.

Relationship between DAG quality and reasoning compression. We test whether traces with stronger independent graph diagnostics (acyclicity, low orphan ratio, and directional consistency) admit shorter compressed forms under Topology-Guided Distillation. Intuitively, a well-structured DAG has fewer redundant branches and shorter critical paths, so the transitive reduction yields a more compact chain. The raw-graph scores q_{dir} and q_{acyc} enter Eq. 3 independently, while the projected DAG supplies localized revision targets. This relationship is evaluated rather than assumed.

D REWARD FUNCTION DETAILS

Outcome reward. Extracts the final mathematical answer from the last `\boxed{\}` expression, falling back to an explicit final numeric expression when needed. Symbolic equivalence to the reference answer is checked with `math_verify`; a correct answer receives 1 and an incorrect or unparseable answer receives 0.

Format reward. The released Qwen3.5 template owns the opening `<think>` prefix, so it is part of the prompt rather than the scored continuation. The canonical reward therefore requires the continuation to contain non-empty reasoning, exactly one closing `</think>` tag, and a subsequent boxed answer; it assigns 1 only when all three conditions hold and 0 otherwise. A generated opening tag is rejected, preventing the policy from receiving credit for duplicating a prompt-owned control token. This `prefilled_think` protocol is fixed in the launch environment and recorded in the immutable run manifest.

Length-only control. The explicit-efficiency baseline uses linear decay from 1.0 (at $\leq 2,000$ characters) to 0.0 (at $\geq 8,000$ characters). This signal is absent from the main TopoPRM reward; all evaluation tables independently report tokenizer-token counts.

Raw-graph rewards. The GRPO reward consumes the coverage-adjusted direction and acyclicity scores from Eqs. (1)–(2) as two aligned sources from one frozen encoder pass. It never replaces them with a scalar average. Orphan conclusions on the projected DAG remain an interpretable defect record for distillation but are not a sixth source in the canonical Stage-II capacity.

Continuity reward. Uses regex-based expression matching to verify that each step references expressions or claims from prior steps or the problem statement. The continuity score q_{cont} (Eq. 5) measures the fraction of supported steps.

E SUPPLEMENTARY MATHEMATICAL DEFINITIONS

Independent topology sources. Let c_H be the fraction of nodes incident to a dependency edge in the raw graph, and let d_H and u_H be the confidence-weighted forward-edge and cycle-edge mass defined in Eqs. (1)–(2). The affine shrinkage by c_H maps missing edge evidence to uncertainty at 0.5 rather than structural perfection. The two scores stay separate through aggregation: q_{dir} responds to local reversal without requiring a cycle, whereas q_{acyc} responds to strongly connected support even when backward mass is held fixed. Their interaction in Eq. 3 can reward joint satisfaction without cancelling their distinct counterfactual responses.

Continuity score (full form). The local continuity quality is

$$q_{\text{cont}}(y) = \begin{cases} 1.0, & \eta(y) = 1, \\ 0.8\eta(y), & \eta(y) < 1, \end{cases} \quad \eta(y) = \frac{1}{T} \sum_{i=1}^T \mathbb{1}[\text{supported}(s_i, s_{<i}, q)]. \quad (5)$$

Remark. Eq. (5) is the engineering-tuned continuity term used in the implementation. It is a separate source in the Choquet reward rather than being folded into topology before aggregation.

Topology-conditioned revision dispatch. The revision instruction is generated by

$$P_r = \psi(r_{\text{out}}, q_{\text{dir}}, q_{\text{acyc}}, \rho_{\text{orphan}}, q_{\text{cont}}, \ell), \quad (6)$$

where ℓ identifies the offending edge, cyclic component, or unsupported step when one can be localized.

F TOPOLOGY-GUIDED DISTILLATION DETAILS

Topology-Guided Distillation is an optional compression stage whose goal is to retain the teacher’s structural reasoning profile while reducing generation length in compact students. Table 9 compares

Table 9: Matched 9B-to-4B distillation protocol. All student methods use the same teacher, candidate prompts, accepted-target count, token budget, and optimizer budget. Accuracy is pass@1; Dir. and Acyc. are paired with the independent edge audit. **Cells remain pending until the single canonical run; marginal accuracy uses Wilson intervals and matched differences use paired item-level bootstrap intervals, neither training variance.**

Method	GSM8K	MATH-500	Avg. Tok	Dir.	Acyc.	Accept%
9B teacher	—	—	—	—	—	—
4B initialization	—	—	—	—	—	—
Static teacher traces	—	—	—	—	—	—
Generic revision	—	—	—	—	—	—
Length-only revision	—	—	—	—	—	—
Topology-Guided Distillation	—	—	—	—	—	—

Algorithm 1 Topology-Guided Distillation Training Pipeline

Require: Base policy π_0 ; data $\mathcal{D}_{\text{sft}}, \mathcal{D}_{\text{rl}}$; Choquet coefficients $\{m_i, m_{ij}\}$
Ensure: Teacher π_T , compact student π_S

- 1: $\pi_\theta \leftarrow \text{SFT}(\pi_0, \mathcal{D}_{\text{sft}})$ ▷ Stage I
- 2: **for** each GRPO step on $\mathcal{D}_{\text{rl}}$ **do** ▷ Stage II
- 3: Sample $\{y^{(i)}\}_{i=1}^G \sim \pi_\theta(\cdot | x)$; build $\mathcal{H}^{(i)} = \mathcal{E}_\phi(y^{(i)})$ and $\mathcal{G}^{(i)} = \Pi(\mathcal{H}^{(i)})$
- 4: Compute $r_{\text{out}}^{(i)}, r_{\text{fmt}}^{(i)}, q_{\text{dir}}^{(i)}, q_{\text{acyc}}^{(i)}, q_{\text{cont}}^{(i)}$
- 5: Compute $r_{\text{total}}^{(i)}$ by Eq. 3; normalize within the GRPO group
- 6: Update π_θ with the GRPO objective
- 7: **end for**
- 8: $\pi_T \leftarrow \pi_\theta'$; init. π_S from a smaller pretrained checkpoint
- 9: **for** each prompt $x \in \mathcal{D}_{\text{rl}}$ **do** ▷ Stage III
- 10: $y_{\text{init}} \sim \pi_S(\cdot | x)$; extract defect record $F(y_{\text{init}})$ from $\mathcal{H}$ and $\mathcal{G}$
- 11: $P_r \leftarrow \psi(F(y_{\text{init}}))$; sample $y_{\text{rev}} \sim \pi_T(\cdot | x, y_{\text{init}}, P_r)$
- 12: Retain (x, y_{rev}) only if answer, format, direction, acyclicity, and budget constraints pass
- 13: Update π_S with $\mathcal{L}_{\text{distill}}$ (Eq. 4)
- 14: **end for**
- 15: **return** (π_T, π_S)

static traces, generic revision, length-only revision, and Topology-Guided Distillation under the same teacher, student initialization, accepted-target count, and budgets. Direction and acyclicity retention must survive the independent audit; this extension returns to the main paper only if the single canonical run supports a central claim.

Why static trace distillation is insufficient. Conventional trace distillation (Hinton et al., 2015; Gu et al., 2024) presents every student with the same teacher solutions. It therefore cannot target a student’s particular missing dependency, and a length-only target can become shorter by deleting support that is necessary for the conclusion. Topology-Guided Distillation instead constructs supervision after observing the compact student’s own trace.

Defect localization and revision. For each initial student attempt, we retain the raw edge list and projected DAG. The defect record contains the endpoints of backward edges, nodes in a non-trivial strongly connected component, orphan conclusions in the projected DAG, and unsupported local transitions. The dispatcher selects the highest-priority localized defect and asks the frozen Stage-II teacher to rewrite the solution within a fixed response budget while preserving a boxed final answer.

Target acceptance. A revised trace is retained only if it has a correct boxed answer and valid format, reaches the calibrated topology threshold, does not reduce coverage-adjusted direction consistency or acyclicity relative to the initial trace, and stays within the shared token budget. We log every rejection reason and report acceptance coverage. This makes direction and acyclicity part of the distillation data contract, not merely an evaluation diagnostic.

Student training and controls. The 4B student is trained on accepted revisions with Eq. 4. Matched controls use the same teacher, prompts, number of accepted targets, and optimization budget but replace topology-conditioned dispatch with either a generic revision instruction or an explicit length-only instruction. External answer accuracy and generated tokens are primary; extractor-based structural retention is reported as internal-consistency evidence and paired with the independent human edge audit.

G CLOSEST-METHOD DESIGN AUDIT

Table 10 isolates representation and optimization choices rather than comparing numbers obtained under incompatible models or evaluators. A method may use a graph without exposing violations that a DAG representation has already ruled out. Our distinguishing object is the unconstrained directed support graph before projection: backward and cycle-participating support remain separately measurable and enter the online reward through explicit source interactions.

Table 10: Design audit of the closest topology-based reasoning objectives. “Graph-form output” means that the policy must serialize a graph. “Raw dir.” asks whether support edges that oppose text order remain measurable before a DAG constraint; “cycle signal” asks whether circular semantic support is separately scored. This is a mechanism comparison, not a performance claim.

Method	Relation represented	Graph-form output	Raw dir.	Cycle signal	Online RL	Source interaction
GoT-R1 (Li et al., 2026b)	Declared predecessor in serialized graph	Yes	No	No	Yes	Additive
SARL (Wang et al., 2026c)	Undirected latent-function transition	No	No	No	Yes	Additive
TUR-DPO (Abdullah et al., 2026)	Elicited claim-support graph	No	No	Yes	No (DPO)	Additive
FlowTracer (Dong et al., 2026b)	Attention-induced answer-flow DAG	No	No	No	Yes	No
TOPOPRM	Unconstrained directed semantic support	No	Yes	Yes	Yes	2-additive

H RETAINED EMNLP SUBMISSION RESULTS

Archival evidence, not current ICLR claims. The nine result tables rendered in the EMNLP submission PDF are retained below with their numeric values unchanged. Retired method names are expanded or replaced only to avoid reintroducing abandoned terminology. These tables mix earlier checkpoints, decoders, response budgets, and in some cases explicitly marked approximate entries. The ICLR main tables use only the canonical single-run protocol in Section 4.1; archival rows will be promoted only after row-level provenance and protocol revalidation.

Table 11: Archived EMNLP Table 1: results on nine public benchmarks across four task categories (Elementary: GSM8K, MATH-500; Olympiad: OlympiadBench, Omni-Math; Competition: AIME’24, AIME’25, CNMO’24; General: MMLU, GPQA-Diamond). All numbers are the originally submitted pass@5 accuracies (%).

Model	Params	Elementary		Olympiad		Competition			General		Avg.
		GSM8K	MATH	Olympiad	Omni	AIME'24	AIME'25	CNMO'24	MMLU	GPQA-D	
Reference open-source models											
Qwen3.5-2B	2B	77.2	60.8	48.5	13.3	6.7	13.3	36.1	32.8	37.9	36.3
Qwen3.5-4B	4B	85.1	55.8	32.8	43.4	0.0	13.3	10.0	67.4	29.8	37.5
Qwen2.5-Math-7B-Instruct (Yang et al., 2024b)	7B	95.2	73.8	38.0	22.0	13.0	10.0	13.3	54.8	28.3	38.7
Llama-3.1-8B-Instruct (Grattafiori et al., 2024)	8B	84.5	52.2	15.0	10.0	3.0	2.0	3.0	68.0	30.0	29.7
DeepSeek-R1-Distill-Llama-8B (DeepSeek-AI, 2025)	8B	83.0	89.1	48.0	25.0	50.4	35.0	25.0	52.0	49.0	50.7
Qwen2.5-7B (Base) (Yang et al., 2024a)	7B	85.4	49.8	25.0	18.0	7.0	5.0	8.0	74.2	36.4	34.3
+ SFT	7B	91.3	65.4	42.5	54.2	16.7	12.2	16.7	70.1	28.3	44.2
+ SimpleRL-Zoo (Zeng et al., 2023)	7B	91.7	78.2	38.6	40.4	20.0	15.6	15.0	62.5	55.2	46.4
+ Open-Reasoner-Zero (Hu et al., 2026)	7B	89.0	81.4	47.9	38.4	17.9	15.6	13.3	62.5	36.6	44.6
+ TopoPRM (Ours)	7B	95.8	72.6	47.8	57.2	16.7	26.7	47.0	82.7	51.0	55.3
DeepSeek-R1-Distill-Qwen-7B (Base) (DeepSeek-AI, 2025)	7B	73.8	36.8	25.8	28.0	46.7	30.0	22.9	23.7	12.6	33.4
+ SFT	7B	86.0	92.8	52.0	48.0	55.5	40.0	30.0	50.0	49.1	55.9
+ GRPO (Shao et al., 2024)	7B	85.1	67.4	56.2	72.0	50.0	40.0	57.8	50.1	17.7	55.1
+ TopoPRM (Ours)	7B	84.3	66.6	56.4	72.2	55.5	40.0	55.4	60.3	34.3	58.3
Qwen3.5-9B (Base)	9B	69.5	66.6	55.2	61.5	16.7	21.1	16.7	28.6	37.4	41.5
+ SFT	9B	94.1	68.2	44.0	53.8	30.0	15.6	30.0	78.4	47.0	51.2
+ GRPO (Shao et al., 2024)	9B	93.3	63.0	56.2	72.0	46.7	30.0	57.8	63.0	17.7	55.5
+ TopoPRM (Ours)	9B	97.9	81.9	56.2	76.2	43.3	18.9	47.0	86.4	40.4	60.9

Table 14: Archived EMNLP Table 9: reward aggregation comparison on the in-domain critique set. All numeric values are retained from the submission.

Strategy	Mid	High	Fmt%	Len	Collapse%↓
Linear (weighted sum)	37.2	22.3	94.6	364	75.1
Clipped linear	29.0	21.5	93.0	382	67.7
Legacy hierarchical aggregation	37.2	22.3	94.6	364	37.9

Table 12: Archived EMNLP Table 3: reward component ablation on Qwen3.5-9B. Avg is the originally submitted unweighted mean pass@1 across eight benchmarks.

Reward Config.	Avg (%)	Δ	Tok	Acc/kTok $\uparrow$	Collapse% $\downarrow$
TopoPRM (Full)	45.6	—	1692	26.9	37.9
w/o ACE	41.7	−3.8	1958	21.3	42.1
w/o Topology	41.8	−3.8	1943	21.5	44.6
w/o Continuity	16.3	−29.3	2047	8.0	68.8
Outcome Only	40.8	−4.7	1823	22.4	57.8

Table 13: Archived EMNLP Table 4: cross-scale distillation from the 9B TopoPRM teacher. Token Ratio is mean generation length relative to the teacher; Structural Retention aggregates the originally submitted DAG-profile metrics on MATH-500.

Scale	GSM8K	MATH	Tok	Ratio	Struct. Ret.
9B (Teacher)	84.3	66.6	679	1.00 $\times$	1.00 $\times$
4B	82.8	61.5	382	0.56 $\times$	0.93 $\times$
2B	78.9	57.1	305	0.45 $\times$	0.88 $\times$
0.8B	73.6	52.4	244	0.36 $\times$	0.81 $\times$
4B (SFT-distill)	79.4	58.2	438	0.64 $\times$	0.86 $\times$
4B (Off-policy KL)	76.8	49.3	4096	>1.0 $\times$	0.41 $\times$

Table 15: Per-benchmark results and component ablations of TOPOPRM. We report **p@1**, **p@5**, **maj@5**, mean generated tokens / total wall-clock seconds (**Tok/Time**, in ks), and token-efficiency **Acc/kTok** = $p@1 / (\text{tokens}/1000)$. We use the same protocol in A. **Bold** / underline: best / second-best per column.

Dataset/Method	GSM8K (1,319)					MATH-500 (500)					OlympiadBench (134)					Omni-MATH (262)				
	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok
Base (Qwen3.5-9B)	55.2	87.8	55.7	1536/45.0	35.9	19.8	39.2	21.4	1536/18.7	12.9	11.2	25.4	14.2	2560/26.5	4.4	16.4	33.2	17.9	2559/52.0	6.4
Outcome Only	93.3	97.7	93.6	1047/34.2	89.1	<u>50.8</u>	59.0	55.2	1511/18.9	33.6	31.3	46.3	37.3	<u>2043/19.7</u>	<u>15.3</u>	41.6	<u>55.3</u>	48.9	<u>1845/35.2</u>	<u>22.5</u>
w/o Topology	93.0	97.7	93.9	<u>946/25.0</u>	<u>98.3</u>	<u>51.0</u>	<u>59.4</u>	55.6	<u>1506/16.2</u>	<u>33.9</u>	29.8	<u>50.7</u>	<u>38.1</u>	2468/11.8	12.1	<u>42.4</u>	54.2	45.8	2430/17.8	17.4
w/o Continuity	51.0	86.1	52.6	1528/48.3	33.4	21.6	39.6	23.2	1536/19.1	14.1	8.2	25.4	9.0	2560/ 8.7	3.2	14.9	30.5	17.6	2560/18.5	5.8
w/o ACE	<u>93.8</u>	98.0	<u>94.6</u>	1006/28.3	93.2	<u>50.8</u>	<u>59.8</u>	<u>55.2</u>	1515/18.2	33.5	32.8	47.0	35.8	2491/ 9.6	13.2	42.0	53.8	47.3	2463/19.4	17.1
★ TopoPRM (Full)	94.1	<u>97.9</u>	94.8	796/19.7	118.2	51.0	81.9	54.6	1454/16.6	35.1	32.8	50.7	38.8	1768/18.1	18.6	42.8	55.7	<u>47.7</u>	1592/19.0	26.9

Dataset/Method	AIME 2024 (30)					AIME 2025 (90)					CNMO 2024 (30)					MMLU (500)				
	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok	p@1	p@5	maj@5	Tok/Time	Acc/kTok
Base (Qwen3.5-9B)	16.7	33.3	33.3	4083/ 8.5	4.1	2.2	4.4	1.1	2560/ 5.9	0.9	10.0	16.7	3.3	2560/ 1.9	3.9	28.6	41.4	26.2	512/55.1	55.9
Outcome Only	16.7	33.3	20.0	<u>2546/ 5.6</u>	6.6	<u>13.3</u>	15.6	8.9	<u>2547/16.7</u>	<u>5.2</u>	16.7	<u>33.3</u>	<u>26.7</u>	<u>2530/ 5.5</u>	6.6	63.0	84.0	61.0	512/146.2	123.0
w/o Topology	<u>20.0</u>	<u>40.0</u>	<u>30.0</u>	2560/ 2.1	<u>7.8</u>	12.2	<u>16.7</u>	10.0	2560/ 6.2	4.8	<u>20.0</u>	40.0	<u>26.7</u>	2560/ 2.0	<u>7.8</u>	<u>65.8</u>	<u>85.0</u>	<u>63.0</u>	512/32.3	<u>128.5</u>
w/o Continuity	6.7	13.3	3.3	2560/ 2.3	2.6	1.1	4.4	0.0	2560/ 7.0	0.4	6.7	20.0	3.3	2560/ 2.4	2.6	20.2	34.8	18.6	512/80.4	39.5
w/o ACE	<u>20.0</u>	30.0	23.3	2560/ 2.6	<u>7.8</u>	<u>13.3</u>	<u>17.8</u>	12.2	2560/ 7.7	<u>5.2</u>	<u>20.0</u>	30.0	<u>26.7</u>	2560/ 2.6	<u>7.8</u>	61.1	83.1	61.2	508/58.1	120.3
★ TopoPRM (Full)	30.0	43.3	33.3	2431/ 5.4	12.3	15.6	18.9	<u>11.1</u>	2457/ 7.7	6.4	30.0	47.0	30.0	2530/ 2.5	11.9	68.2	86.4	68.8	508/48.0	134.2

Table 16: Main results on the in-domain mathematical critique benchmark. **Middle** and **High** denote middle-school and high-school subsets. Acc measures critique scoring accuracy; Fmt% measures structural compliance; Len reports mean prediction length in tokens. Best results per column are in **bold**, second-best are underlined.

Model	Middle School			High School			Overall
	Acc	Fmt%	Len	Acc	Fmt%	Len	
Zero-shot Baseline							
Qwen3-32B	16.0	64.5	1 200	7.0	43.5	1 240	11.5
Qwen3-32B + SFT	29.0	78.5	1 164	19.5	65.0	1 233	24.3
RFT Variants							
+ Outcome Only	19.5	89.0	410	13.0	85.5	475	16.3
+ w/o Topology	23.0	91.0	397	16.5	86.0	460	19.8
+ w/o Continuity	27.5	90.5	379	20.0	85.0	436	23.8
+ Clipped Linear	29.0	93.0	382	21.5	88.5	458	25.3
+ TopoPRM (full)	37.2	94.6	364	22.3	89.3	424	29.7
Topology-Guided Distillation Fine-tuned							
Qwen3.5-9B + full	40.5	98.0	1 602	30.5	94.0	2 191	35.5
Qwen3-8B + full	60.0	94.0	532	58.0	97.0	613	59.0

Table 17: Reward component ablation on **Qwen3-32B** over the in-domain critique validation set. Each configuration toggles the topology (r_{topo}) and continuity (r_{cont}) signals on top of the outcome reward r_{out} . **Avg** is the unweighted mean of Middle- and High-school accuracy; Δ denotes the gap relative to the full TopoPRM. Rows are sorted by Avg in ascending order to highlight the monotonic contribution of each component.

Configuration	Reward Components			Accuracy (%)			
	r_{out}	r_{topo}	r_{cont}	Mid	High	Avg	Δ
Outcome Only	✓	✗	✗	19.5	13.0	16.3	−13.0
w/o Topology	✓	✗	✓	23.0	16.5	19.8	−9.5
w/o Continuity	✓	✓	✗	27.5	20.0	23.8	−5.5
TopoPRM (Full)	✓	✓	✓	33.0	25.5	29.3	—

Table 18: Accuracy-length trade-off on the in-domain critique validation set (Qwen3-32B). Acc/k-Tok = Overall-Acc per 1,000 tokens. Len $\uparrow$ is the length increase relative to TopoPRM (lower is better).

Config.	Middle		High		Efficiency	
	Acc	Len	Acc	Len	Acc/kTok $\uparrow$	Len $\uparrow$
TopoPRM (full)	37.2	364	22.3	424	75.5	—
w/o Continuity	27.5	379	20.0	436	58.4	+4.1%
w/o Topology	23.0	397	16.5	460	46.2	+9.0%
Outcome Only	19.5	410	13.0	475	36.8	+12.6%
Qwen3-32B (base)	16.0	1200	7.0	1240	9.4	+230%

Table 19: Structural quality of extracted reasoning DAGs on the in-domain middle-school evaluation set. Acyclic% is the fraction of valid DAGs; No-Orphan% is the fraction of traces without unsupported conclusion nodes; Dir-Cons. is mean directional consistency; DAG Depth is mean dependency depth; Edge-Keep% is the fraction of dependency edges retained after compression.

Model	Acyclic% $\uparrow$	No-Orphan% $\uparrow$	Dir-Cons. $\uparrow$	DAG Depth	Edge-Keep% $\uparrow$
Qwen3-32B + SFT	100	61	0.49	1.05	43
+ GRPO (outcome only)	100	82	0.89	1.12	56
+ GRPO (w/o topology)	100	84	0.91	1.14	58
+ GRPO (w/o continuity)	100	83	0.91	1.13	57
+ GRPO (TopoPRM full)	100	89	1.00	1.17	60

H.1 ADDITIONAL EMNLP-WORKSPACE RESULTS

The following six supplementary tables were present in the accompanying EMNLP source workspace but were not among the nine result tables rendered in the submission PDF. Their values are retained for reference under the same archival restrictions.

Table 20: Accuracy evaluation across 9 public math reasoning benchmarks. All our models follow one transformers-backend protocol with the model’s native chat template, shared answer extraction, and shared matching. Open-reported numbers are taken from official papers and model cards. **Bold** marks the best result per block and underline marks the second best. **This is the retained EMNLP-submission table: entries prefixed by ~ are archival approximations, not ICLR claims, and all rows are being revalidated under the canonical protocol.**

Model	Params	Elem. Math		Olympiad Math		Competition		CNMO	General	
		GSM8K	MATH-500	Olympiad	Omni-MATH	AIME'24	AIME'25	CNMO'24	MMLU	GPQA-D
Reference open-source models (open-reported)										
Qwen2.5-7B-Instruct (Yang et al., 2024a)	7B	91.6	75.5	~25.0	~18.0	~7.0	~5.0	~8.0	75.4	36.4
Qwen2.5-Math-7B-Instruct (Yang et al., 2024b)	7B	95.2	83.6	~38.0	~22.0	~13.0	~10.0	~15.0	~60.0	~35.0
Llama-3.1-8B-Instruct (Grattafiori et al., 2024)	8B	84.5	52.2	~15.0	~10.0	~3.0	~2.0	~3.0	~68.0	~30.0
DeepSeek-R1-Distill-Llama-8B (DeepSeek-AI, 2025)	8B	~83.0	89.1	~48.0	~25.0	50.4	~35.0	~25.0	~52.0	49.0
DeepSeek-R1-Distill-Qwen-7B (DeepSeek-AI, 2025)	7B	~86.0	92.8	~52.0	~28.0	55.5	~40.0	~30.0	~55.0	49.1
SimpleRL-Zoo (Qwen2.5-7B) (Zeng et al., 2023)	7B	91.7	78.2	~38.6	~40.4	20.0	~15.6	~15.0	~62.5	~55.2
Ours: 9B family (Qwen3.5-9B base, unified transformers protocol)										
Qwen3.5-9B (base)	9B	67.1	43.8	29.8	51.4	6.7	3.3	12.1	52.9	27.3
+ SFT	9B	94.1	50.8	32.8	42.4	30.0	15.6	30.0	68.2	20.7
+ GRPO (outcome-only)	9B	93.3	50.8	31.3	41.6	16.7	13.3	16.7	63.0	—
+ GRPO (w/o topology)	9B	93.0	51.0	29.8	42.4	20.0	12.2	20.0	65.8	16.2
+ GRPO (w/o continuity)	9B	—	—	—	—	—	—	—	—	—
+ TopoPRM (hierarchical)	9B	93.5	49.8	32.8	42.8	26.7	12.2	26.7	61.8	16.7
+ TopoPRM (gated)	9B	<u>93.8</u>	50.8	32.8	42.0	20.0	13.3	20.0	61.1	21.2
Ours: DR1-7B family (DeepSeek-R1-Distill-Qwen-7B base, unified transformers protocol)										
DR1-7B (base, measured)	7B	60.8	<u>68.4</u>	<u>57.8</u>	72.8	46.7	<u>33.3</u>	23.3	42.4	<u>16.2</u>
+ SFT	7B	73.8	36.8	25.8	48.0	0.0	0.0	22.9	23.7	12.6
+ GRPO (outcome-only)	7B	85.1	67.4	56.2	72.0	46.7	30.0	57.8	50.1	17.7
+ GRPO (w/o topology)	7B	84.5	68.8	—	—	36.7	<u>33.3</u>	53.0	—	11.1
+ GRPO (w/o continuity)	7B	85.1	66.4	—	—	36.7	26.7	50.6	—	14.6
+ TopoPRM (hierarchical)	7B	84.3	66.6	56.4	<u>72.2</u>	50.0	40.0	55.4	49.6	15.2
Ours: Distilled students (Topology-Guided Distillation, target $\leq 4B$)										
Student (4B, SFT distill)	4B	~89.4	~49.8	~24.0	~30.2	~12.1	~8.5	~6.8	~82.6	~40.7
Student (4B, Topology-Guided Distillation)	4B	~91.3	~52.1	~27.6	~34.8	~15.4	~10.2	~8.1	~83.5	~41.9
Student (2B, Topology-Guided Distillation)	2B	~86.7	~46.5	~21.3	~27.1	~9.2	~6.4	~5.7	~80.8	~38.4
Student (0.8B, Topology-Guided Distillation)	0.8B	~79.8	~39.4	~14.0	~18.6	~6.1	~4.5	~3.9	~76.3	~34.2

Table 21: Detailed unified metrics across our measured configurations (num_samples=5, native chat template, same extractor). **Cor** = correct count; **Err** = error count; **Tok** = mean generation tokens; **Acc/kTok** = pass@1 points per 1 000 tokens (higher is more token-efficient). Best within each block in **bold**.

GSM8K									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	60.8	80.8	64.2	60.8	60.8	800/519	617	98.5
+ SFT	7B	73.8	91.5	80.2	73.7	73.8	970/349	634	116.4
+ GRPO (outcome-only)	7B	85.1	94.0	88.3	85.1	85.1	1120/199	674	126.3
+ GRPO (w/o topology)	7B	84.5	93.9	87.5	84.5	84.5	1113/206	686	123.2
+ GRPO (w/o continuity)	7B	85.1	94.2	87.9	85.1	85.1	1121/198	640	133.0
+ TopoPRM (hierarchical)	7B	84.3	93.6	88.3	84.3	84.3	1110/209	679	124.2
<i>9B family (our unified protocol)</i>									
Qwen3.5-9B (base)	9B	67.1	—	—	—	67.1	850/469	1536	43.7
MATH-500									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	68.4	70.6	67.4	68.4	68.4	322/178	4957	13.8
+ SFT	7B	36.8	54.8	43.4	36.8	36.8	185/315	3731	9.9
+ GRPO (outcome-only)	7B	67.4	72.6	67.4	67.4	67.4	318/182	6001	11.2
+ GRPO (w/o topology)	7B	68.8	73.2	68.8	68.8	68.8	324/176	5973	11.5
+ GRPO (w/o continuity)	7B	66.4	71.8	67.8	66.4	66.4	312/188	6088	10.9
+ TopoPRM (hierarchical)	7B	66.6	72.8	68.2	66.6	66.6	314/186	6053	11.0
<i>9B family (our unified protocol)</i>									
Qwen3.5-9B (base)	9B	43.8	—	—	—	43.8	203/297	3072	14.3
AIME 2024									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	46.7	46.7	36.7	46.7	46.7	10/20	8138	5.7
+ SFT	7B	0.0	3.3	0.0	0.0	0.0	0/30	4301	0.0
+ GRPO (outcome-only)	7B	46.7	36.7	33.3	46.7	46.7	9/21	8192	5.7
+ GRPO (w/o topology)	7B	36.7	46.7	33.3	36.7	36.7	8/22	8192	4.5
+ GRPO (w/o continuity)	7B	36.7	40.0	33.3	36.7	36.7	8/22	8192	4.5
+ TopoPRM (hierarchical)	7B	50.0	40.0	33.3	50.0	50.0	10/20	8192	6.1
<i>9B family</i>									
Qwen3.5-9B (base)	9B	6.7	—	—	—	6.7	1/29	4096	1.6
Omni-MATH (500 subset)									
Model	Params	pass@1	pass@5	maj@5	prm@5	F1	Cor/Err	Tok	Acc/kTok
<i>DRI-7B family</i>									
DR1-7B (base, measured)	7B	72.8	—	—	72.8	72.8	348/152	4720	15.4
+ SFT	7B	48.0	—	—	48.0	48.0	235/265	7114	6.7
+ GRPO (outcome-only)	7B	72.0	—	—	72.0	72.0	347/153	3562	<u>20.2</u>
+ TopoPRM (hierarchical)	7B	72.2	—	—	72.2	72.2	348/152	3529	20.5

Table 22: Ablation of TopoPRM reward components on the Qwen3.5-9B backbone. All variants start from the same SFT checkpoint and run identical GRPO steps. We report pass@1 / pass@5 (%), mean generation tokens (Tok), and Acc/kTok = pass@1 per 1 000 tokens. Δ shows the change in pass@5 relative to the full TopoPRM (hierarchical). Best in **bold**.

Reward Config	GSM8K				MATH-500				AIME'24				Omni-MATH				MMLU			
	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5	p@1	p@5	Tok	Δ p@5
TopoPRM (full, hier.)	93.5	97.9	1026	—	49.8	59.0	1506	—	26.7	43.3	2560	—	42.8	55.7	2465	—	61.8	81.9	508	—
TopoPRM (gated)	93.8	<u>97.9</u>	1006	+0.1	50.8	59.8	1515	+0.8	20.0	30.0	2560	−13.3	42.0	53.8	2463	−1.9	61.1	83.1	508	+1.3
w/o Topology	93.0	97.7	946	−0.2	51.0	59.4	1506	+0.4	20.0	40.0	2560	−3.3	42.4	54.2	2430	−1.5	65.8	<u>85.0</u>	512	+3.1
w/o Continuity	51.0	86.1	1528	−11.8	21.6	39.6	1536	−19.4	6.7	13.3	2560	−30.0	14.9	30.5	2560	−25.2	20.2	34.8	512	−47.1
Outcome Only GRPO	93.3	97.7	1047	−0.2	50.8	59.0	1511	+0.0	16.7	33.3	2546	−10.0	41.6	55.3	1845	−0.4	63.0	84.0	512	+2.1
SFT baseline	94.1	97.7	796	−0.2	50.8	<u>60.4</u>	1454	+1.4	<u>30.0</u>	43.3	2431	+0.0	<u>42.4</u>	53.1	1592	−2.7	<u>68.2</u>	86.4	512	+4.5

Table 23: Generation efficiency of Qwen3.5-9B variants. All runs use 5 samples per item on the same hardware. **Tok** = mean generation tokens per sample; **Time** = total wall-clock seconds for the benchmark; **Tok/s** = throughput (total tokens / time); **Acc/kTok** = pass@1 per 1 000 tokens (higher = more token-efficient). Best in **bold**.

Model	GSM8K (1,319 items)				MATH-500				AIME'24 (30 items)				Omni-MATH (262 items)			
	Tok	Time	Tok/s	Acc/kTok	Tok	Time	Tok/s	Acc/kTok	Tok	Time	Tok/s	Acc/kTok	Tok	Time	Tok/s	Acc/kTok
TopoPRM (Full)	796	19.7k	266	118.2	1454	16.6k	220	34.9	2431	5.4k	67	<u>12.3</u>	1592	31.7k	66	26.6
Outcome Only GRPO	1047	34.2k	202	89.1	1511	18.9k	200	33.6	2546	5.6k	68	6.5	1845	35.2k	69	22.5
w/o Topology	946	25.0k	250	98.3	1506	16.2k	232	33.9	2560	2.1k	186	7.8	2430	17.8k	179	17.4
w/o Continuity	1528	48.3k	209	33.4	1536	19.1k	201	14.1	2560	2.3k	165	2.6	2560	18.5k	181	5.8
Legacy gated agg.	1006	28.3k	234	93.2	1515	18.2k	208	33.5	2560	2.6k	147	7.8	2463	19.4k	167	17.0
TopoPRM (Hier.)	1026	33.4k	203	91.1	1506	18.0k	209	33.1	2560	2.6k	149	10.4	2465	19.0k	170	<u>17.3</u>

Table 24: First-Attempt vs Revised accuracy on 1K sampled AIME-2024 / MATH problems (Qwen3.5-9B family). First Attempt: one sample from the model. Revised Attempt: a second sample conditioned on the first attempt and a topology-aware revision prompt P_r (Eq. 6). Baseline = plain distillation reviser without DAG signal. Avg-Len refers to the revised response only.

Model	First Att.	Revised Att.	Gain	Revised Len	Revised / First
Qwen3.5-9B (base, no SRT)	~31.8	~34.6	~+2.8	~742	~1.07
+ Plain distillation reviser (no topology)	~33.1	~37.0	~+3.9	~688	~1.02
+ Topology-Guided Distillation reviser (Eq. 6)	~33.0	~39.8	~+6.8	~624	~0.94

Table 25: Case study comparing teacher and topology-guided distilled-student traces on the same prompt.

	Teacher (TopoPRM)	Distilled student
Problem	High-school algebra problem requiring multi-step substitution and verification (same prompt hash for both rows).	
# Steps	~14	~9
# Tokens	~612	~428
Score-Acc	~0.93	~0.91
Key reasoning	Core conclusion supported by an explicit minimal chain of derivations and substitutions.	Preserves key support edges while compressing redundant branches and intermediate restatements.